

Charles University in Prague
Faculty of Mathematics and Physics

BACHELOR THESIS



Stanislav Kučera

Generating random pattern-avoiding matrices

Computer Science Institute of Charles University

Supervisor of the bachelor thesis: RNDr. Vít Jelínek, Ph.D.

Study programme: Computer Science

Study branch: General Computer Science

Prague 2016

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University in Prague has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

signature of the author

Title: Generating random pattern-avoiding matrices

Author: Stanislav Kučera

Department: Computer Science Institute of Charles University

Supervisor: RNDr. Vít Jelínek, Ph.D., department

Abstract: Abstract.

Keywords: binary matrix pattern-avoiding Markov chain Monte Carlo

Dedication.

Contents

Preface	3
Introduction	5
1 Markov chain Monte Carlo	6
1.1 Markov chain for pattern-avoiding binary matrices	6
2 An algorithm for testing pattern-avoidance of a general pattern	7
2.1 Sketch of a brute force algorithm	7
2.2 Equivalent mappings	7
3 An algorithm for testing pattern-avoidance of a special pattern	9
3.1 Walking pattern	9
3.2 Dynamic program	9
3.2.1 The algorithm	10
3.2.2 Inner structure	11
3.2.3 Correctness	11
3.2.4 Generalization	12
4 Improvements to basic algorithms	13
4.1 General pattern	13
4.1.1 Improving memory consumption	13
4.1.2 Not mapping empty lines	13
4.1.3 Using the last changed position	13
4.1.4 Line order	14
4.1.5 Mapping approaches	15
4.1.6 Using the whole structure in the next iteration	16
4.2 Walking pattern	17
4.2.1 Using the last changed position	17
4.2.2 Lazy avoid	18
4.3 Parallel computing	19
4.3.1 MCMC parallelism	19
5 Technical documentation	22
5.1 General pattern	22
5.1.1 Construction	22
5.2 Parallel computing	23
5.2.1 MCMC parallelism	23
5.3 Library interface	24
6 User documentation	25
6.1 Configuration file	25
6.1.1 Input	25
6.1.2 Pattern	26
6.1.3 Output	27
6.1.4 Statistics	28

6.2	File input	29
6.2.1	Matrix file	29
6.2.2	Order file	29
6.3	File output	30
6.3.1	Matrix text file	30
6.3.2	Matrix bmp file	30
Conclusion		31
List of Figures		32
List of Tables		33
List of Abbreviations		34
Attachments		35

1 Preface

2 Theses at the faculty of mathematics and physics usually fit into one of three
3 categories:

- 4 1. Theoretical thesis
- 5 2. Experimental thesis
- 6 3. Implementation thesis

7 My thesis does not fit entirely into only one category and it does not try to. The
8 project consists of several similarly important parts which are:

- 9 • Design of the algorithm for generating a special binary matrix
- 10 • Making it run fast on inputs which are usual for researchers
- 11 • Implementing the algorithm to provide practical tool

12 None of these points would make sense alone but together the thesis may become
13 very useful for scientists as it is a common practice to test hypothesis on random
14 data.

Introduction

The area of pattern avoidance is heavily studied for permutations and it also becomes studied more for its generalization - binary matrices.

We denote by $M \in \{0,1\}^{n \times m}$ a binary matrix of size n by m , calling n the number of rows of M - the height of the matrix M and m the number of columns - its width. A line of a matrix is one of its rows or columns and for matrix M , we denote $L(M)$ the ordered set of all lines of M . Its order is given by the natural indexing of rows and columns.

We say a binary matrix M contains a binary matrix P , which we call a “pattern”, as a submatrix, if there is a mapping $f : L(P) \rightarrow L(M)$, such that

- $l \in L(P)$ is a row of P iff $f(l) \in L(M)$ is a row of M
- $\forall l, l' \in L(P) : l < l' \Rightarrow f(l) < f(l')$ (preserves the order)
- $\forall l, l' \in L(P) : \text{if lines } l \text{ and } l' \text{ intersect and there is a one-entry at the intersection, then there is a one-entry at the intersection of } f(l) \text{ and } f(l').$

otherwise, it avoids the pattern P .

$$\begin{array}{ccc}
 \begin{array}{cc} & \begin{array}{cc} 2 & 3 \end{array} \\ \begin{array}{c} 0 \\ 1 \end{array} & \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \end{array} &
 \begin{array}{ccc} & \begin{array}{ccc} 3 & 4 & 5 \end{array} \\ \begin{array}{c} 0 \\ 1 \\ 2 \end{array} & \begin{pmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} &
 \begin{array}{ccc} & \begin{array}{ccc} 3 & 4 & 5 \end{array} \\ \begin{array}{c} 0 \\ 1 \\ 2 \end{array} & \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{array}
 \end{array}$$

Let the pattern P be the left matrix from the picture. The middle matrix contains the pattern P , because a mapping $\{(0,0), (1,2), (2,3), (3,4)\}$ satisfies all the conditions. On the other hand the right matrix avoids the pattern as there is no such mapping.

There are other possible definitions of containing but from now on we will always consider a pattern to be contained as a submatrix. We will always denote M the binary matrix for which we test the containing and by P the pattern that is being tested. Moreover, we denote h the height (the number of rows) of P and w its width.

Generating random matrix

Our main goal is to show an iterated algorithm, which, with high enough number of iterations, for a given forbidden pattern and number n generates a uniformly random binary matrix $M \in \{0,1\}^{n \times n}$ avoiding the pattern.

To create such an algorithm we generalize an algorithm generating fixed size permutations without a forbidden pattern introduced by [Madras-Liu]. To achieve randomness it uses Markov chain Monte Carlo method, which is explained in the next section.

The process starts with a binary matrix of the correct size avoiding the pattern (zero matrix is always sufficient) and in each step it changes uniformly at random one bit of the matrix and tests whether the matrix still avoids the pattern after the change. If it does the step is over, if it does not, we revert the change (flip the bit back) and end the step.

52 It can be proven that if the Markov chain is chosen carefully, this leads to a
53 uniformly random matrix avoiding the pattern, after a sufficient number of steps.
54 Sadly no-one knows what is “sufficient” in this manner of speaking.

55 **Testing avoidance**

56 The hardest task to do in a step is to test whether the new matrix avoids the
57 pattern. We show a very fast algorithm which only works for a special class
58 of binary matrices as well as a slightly less performing algorithm for a general
59 pattern, which again comes as a generalization of an algorithm for permutations
60 from the article by Madras and Liu.

1. Markov chain Monte Carlo

We shall consider discrete-time Markov chains X_0, X_1, X_2, \dots , where each X_i takes values in a finite or countably infinite state space S . For $k = 0, 1, 2, \dots$, the k -step transition probabilities are:

$$p_{i,j}^{(k)} = \Pr(X_{t+k} = j | X_t = i) \quad (i, j \in S)$$

A Markov chain is **irreducible** if the chain can eventually get from each state to every other state, that is, for every $i, j \in S$ there exists a $k \geq 0$ (depending on i and j) such that $p_{i,j}^{(k)} > 0$.

An irreducible chain has period D if D is the greatest common divisor of $\{k \geq 1 | p_{i,i}^{(k)} > 0\}$ for some $i \in S$ (equivalently, for all $i \in S$). A chain is called **aperiodic** if its period is 1. In particular, if an irreducible chain has $P_{i,i} > 0$ for some i , then it is aperiodic.

A Markov chain is said to be **symmetric** if $P_{i,j} = P_{j,i}$ for every pair of states i and j .

[Madras] [Madras-Liu]

1.1 Markov chain for pattern-avoiding binary matrices

74 2. An algorithm for testing 75 pattern-avoidance of a general 76 pattern

77 We begin with a very basic algorithm that we improve a lot to get a fast algorithm
78 for testing avoidance of a general pattern.

79 2.1 Sketch of a brute force algorithm

80 Let $L = \{l_1, l_2, \dots, l_{w+h}\}$ be the lines (rows and columns) of the pattern P .
81 Partial mapping of lines of P is a function f from $L' \subseteq L$ to lines of the big
82 matrix M satisfying two conditions:

- 83 • Both $l' \in L'$ and $f(l')$ are rows or they are both columns.
- 84 • If $l' \in L'$ and $l'' \in L'$ are both rows or columns and $l' < l''$, then $f(l') < f(l'')$.
85 This means partial mapping keeps the order of the lines.
- 86 • If $l' \in L'$ is a row of P and $l'' \in L'$ is a column of P and there is a one-entry
87 at the intersection of l' and l'' , then there is a one-entry at the intersection
88 of $f(l')$ and $f(l'')$.

89 The basic algorithm we use goes as follows. First it maps l_1 to all possible lines
90 of M , creating partial mappings of $\{l_1\} \subseteq L$. For $k = 2, \dots, w + h$ it takes each
91 partial mapping from previous iteration and extends it by adding line l_k to the
92 partial mapping in all possible ways. If we manage to map all the lines of P , then
93 M does not avoid it and if at some point there are no partial mappings to extend
94 it means M avoids P .

95 Note that we do not use even the slightest heuristics like not to map a line
96 having k one-entries to a line with fewer one-entries. The fact we made a useless
97 mapping will be discovered at the time we try to map the k -th orthogonal line
98 (or maybe earlier). This leads to possibly very large number of partial mappings
99 and our next goal will be to reduce this number significantly.

100 2.2 Equivalent mappings

101 There are a lot of possible heuristics we can check before adding a line to a partial
102 mapping, but the most fundamental operation we do to decrease the number of
103 found mapping is saying that two mapping are equivalent.

104 The idea behind it is very basic. There is no need to remember two different
105 mappings if they can be both extended exactly the same way as our function is
106 only supposed to check whether a pattern can be mapped to a big matrix not to
107 find all such mappings.

108 A **level** of a partial mapping is the number of lines mapped by the mapping.

109 We call a line l of a pattern **important** in a partial mapping if one of the
110 conditions is met:

111 • An adjacent line of the pattern has not been mapped yet.
 112 • There is a one-entry on the line l at the intersection with line l' that has
 113 not been mapped yet.
 114 . Otherwise the line is **unimportant** in the mapping.
 115 At the beginning, when no line is mapped, all lines are important. After some
 116 lines get mapped a line can become unimportant in the partial mapping as all
 117 lines that bound in are in the mapping as well. If a line is unimportant in a
 118 partial mapping of some level, it will stay unimportant in all extensions of the
 119 mapping we can find.
 120 We say two partial mappings of the same level are **equivalent** if all important
 121 lines in the mapping of that level are mapped to the same lines of the big matrix
 122 in both mappings.
 123 Picture and use...


3. An algorithm for testing pattern-avoidance of a special pattern

In the previous chapter, we have seen an algorithm for a general forbidden pattern which, using some heuristics, runs pretty fast. In this chapter, we introduce a special kind of a pattern, satisfying additional conditions, for which we can produce much faster algorithm.

3.1 Walking pattern

We call the specific pattern a walking pattern. The additional condition we want the pattern to satisfy is that there is a walk from one corner to the opposite one and all the one-entries of the pattern are contained on the walk.

The **walk** in a matrix is a sequence of some of its entries beginning in the top left corner and ending in the bottom right one. If an entry at the position $[i, j]$ is in the sequence, the next one is either $[i + 1, j]$ or $[i, j + 1]$. Therefore the length of an arbitrary walk is equal to $w + h - 1$.



In the picture you can see a matrix that is a walking pattern as all the one-entries are included in a walk. Not all entries of a walk need to be one-entries though.

It can be shown a walking pattern with a walk is exactly a matrix avoiding a forbidden pattern

$$\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

3.2 Dynamic program

Now that we know what the walking pattern is, we show an algorithm deciding whether a pattern P is contained in a big matrix M or not.

The pattern P is a walking pattern so there is a walk containing all the one-entries of the pattern. We choose one such walk arbitrarily and index its entries $w_1, w_2, \dots, w_{m+n-1}$ starting from the beginning of the walk. For each entry of the walk we remember whether its value is one or zero and whether the walk continues

149 from the entry vertically, which we call a **vertical entry** or horizontally, calling
150 it a **horizontal entry**.

151 For an element e of M at the position $[i, j]$, the matrix $M_{\leq e}$ is a $(i+1) \times (j+1)$
152 submatrix of M consisting of rows with the index smaller than or equal to i and
153 columns with the index smaller than or equal to j . The element e lies in its bottom
154 right corner. Similarly $M_{\geq e}$ is a $(n-i) \times (n-j)$ submatrix of M consisting of
155 rows with the index greater than or equal to i and columns with index greater
156 then or equal to j . The element e is its first element.

157 To determine whether P is contained in M we find out for each element e of
158 M what is the longest part of the pattern that can be found in $M_{\leq e}$. If there is
159 an element for which we manage to find the last entry of the pattern, the pattern
160 is contained in the matrix, otherwise it is avoided.

161 For each element e of M at the position $[i, j]$ we remember two numbers. The
162 number $c_v(e)$ says what is the longest part of the walk in $M_{\leq e}$ with the last entry
163 in j -th column and being a vertical entry. The number $c_h(e)$, symmetrically, says
164 what is the longest part of the walk in $M_{\leq e}$ with the last entry in i -th row and
165 and being a horizontal entry.

166 An observation we make is that if we have a fixed element e of M and any
167 other element e' above e in the same column then if $c_v(e')$ is equal to some k ,
168 then $c_v(e)$ is at least k . This means that for e we can find the maximum part of
169 the pattern ending in the column of e and continuing vertically by looking only
170 to elements in that column above e and since this is true for all of them, it is
171 sufficient to only check the value of the element right above e (at the position
172 $[i-1, j]$). Similarly the argument goes for the value of c_h in horizontal way.


173 The algorithm iterates through diagonals. A diagonal in this matter of speak-
174 ing is a subset of elements of M , such that all elements have the same sum of
175 their coordinates. For example, zero diagonal only consists of an element $[0, 0]$,
176 the first diagonal contains elements $[0, 1]$ and $[1, 0]$, and so on.

177 For simplicity, in the pseudo-code below we do not deal with elements outside
178 M (like $-1, 0$) explicitly. Instead for those elements we just assume the values of
179 c_v and c_h are always equal to zero.

180 3.2.1 The algorithm

- 181 1. For $d = 0, \dots, w + h - 1$
- 182 2. For e element of d -th diagonal at the position $[i, j]$
- 183 3. $e_v := [i - 1, j]$
- 184 4. $e_h := [i, j - 1]$
- 185 5. $c_v(e) := c_v(e_v)$
- 186 6. $c_h(e) := c_h(e_h)$
- 187 7. If $w_{c_v(e)+1}$ can be mapped to e
- 188 8. If $c_v(e) + 1 = w + h + 1$
- 189 9. Terminate - M contains P as a submatrix

190 10. If $w_{c_v(e)+1}$ is a vertical entry
 191 11. $c_v(e) := c_v(e) + 1$
 192 12. Else
 193 13. $c_h(e) := \max\{c_h(e), c_v(e) + 1\}$
 194 14. If $w_{c_h(e)+1}$ can be mapped to e
 195 15. If $c_h(e) + 1 = w + h + 1$
 196 16. Terminate - M contains P as a submatrix
 197 17. If $w_{c_h(e)+1}$ is a vertical entry
 198 18. $c_v(e) := \max\{c_v(e), c_h(e) + 1\}$
 199 19. Else
 200 20. $c_h(e) := \max\{c_h(e), c_h(e) + 1\}$



201 3.2.2 Inner structure

202 To run properly the algorithm needs two structures. The first one is a description
 203 of the walk, which is just an array of the values of its entries as well as the
 204 information whether the entry is vertical or horizontal. The second structure is
 205 a matrix of the values c_v and c_h as described above.

206 3.2.3 Correctness

207 We need to show that the values of c_v and c_h are always correct for the recomputed
 208 elements after at the end of the function. We proceed by induction.

209 For the first element it is definitely true since there can be only the first entry
 210 of the pattern mapped and we check just that.

211 When we compute an element e of a computed diagonal d , by induction as-
 212 sumption all the diagonals $d' < d$ are correctly computed. In particular, the
 213 values are correct in the diagonal $d - 1$. To compute the correct values of e , we
 214 use the values of two element on the diagonal $d - 1$: e_v , which is right above e and

215 e_h , which is the first element to the left of e . If e_v or e_h are outside the matrix
 216 then from that direction we cannot expect to find anything more than just the
 217 first entry of the pattern and that is what we check for.

218 Let v be the true length of the longest part of P in $M_{\leq e}$ continuing vertically
 219 in the same column as e . Now if e itself is not an entry of that part of the pattern,
 220 it is a different element e' in the same column. But then the value of $c_v(e')$ is
 221 correctly computed by the inductive hypothesis and it is copied to all element
 222 underneath. Especially e_v gets the value and the algorithm copies the value from
 223 it to e . On the other hand if e is an entry of the part of the pattern we work with,
 224 it is the last entry. The entry right before the last one needs to be mapped to
 225 the same row or column; therefore, either e_v or e_h contain the part of the pattern
 226 shorter by one and the algorithm extends it to a correct value.

227 3.2.4 Generalization

228 The same algorithm, just rotated by 90 degrees, can be also used for a pattern
 229 where all one-entries are contained on a walk from top right corner to the bottom
 230 left one. Indeed the program uses it and if given a walking pattern it determines
 231 by itself which walk it is.

232 On the other hand a direct generalization for a general pattern does not work.
 233 While we can index all entries of the pattern, when trying to map a certain w_k
 234 to an element it is not sufficient to just check whether w_l is above and w'_l to the
 235 left from the element.

$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & \boxed{1} \\ 1 & \textcircled{1} & \diamond 1 \end{pmatrix} \quad \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & \boxed{1} \\ 0 & 1 & \textcircled{1} & \diamond 1 \end{pmatrix}$$

236 In the picture, let the matrix on the left side be the pattern P and let M
 237 be the other matrix. The entry in the square can be mapped to the element in
 238 the square and the same holds for entries in the circle but it is not a sufficient
 239 condition for the entry in the triangle to be mapped to the element in the triangle.

240 4. Improvements to basic 241 algorithms

242 4.1 General pattern

243 4.1.1 Improving memory consumption

244 As the algorithm for testing avoidance of a general pattern was described in
245 [chapter2], it creates all possible partial mappings and checks whether at least
246 one can be extended to a full mapping. Note that to compute all the partial
247 mappings of some level l , it only uses mappings of level $l - 1$; therefore, it is
248 enough to only store partial mappings of two levels in memory.

249 In [chapter2] we also introduced the idea of (un)important lines for a partial
250 mapping of level l and equivalence based on not using unimportant lines at all -
251 as they are fully bounded by other already mapped lines. When a line becomes
252 unimportant it stays unimportant till the end of the run; as a result, we can forget
253 where we mapped those lines to save memory. This is not as big of a deal as the
254 previous observation was but note there are cases of patterns in which each line
255 becomes unimportant just two levels after it gets added.

256 4.1.2 Not mapping empty lines

257 An empty line is a row or a column that does not contain any one-entries. Such
258 a line can be mapped to any line and if the algorithm leaves space for it (which
259 it does), we do not need to map it at all.

260 4.1.3 Using the last changed position

261 As the MCMC process works, it always changes one element of the big matrix
262 and asks whether it still avoids the pattern. If it does not and we know that
263 before the change it did, we are sure the changed element $[r, c]$ is a part of the
264 pattern. It is hard to use this fact in the algorithm. It just maps one line after
265 another and we do not know at the beginning to which line the changed position
266 lines should be mapped.

267 What we can do is to enforce that neither the r -th line nor the c -th one
268 get skipped. We will only look at the restriction for rows. The restrictions for
269 columns are symmetrical. There are three situations we want to avoid:

- 270 • The first row of P is mapped under the r -th row. This prevents any other
271 row to be mapped to r -th one and we don't want that.
- 272 • The last row of P is mapped above the r -th row. This again prevents any
273 other row to be mapped to r -th one.
- 274 • Two adjacent rows $l, l + 1$ of P are mapped to $L < L'$ respectively and
275 $L < r < L'$ which leaves no other row to be mapped to r .

276 4.1.4 Line order

277 An important thing, if we want the algorithm to run fast, is to choose a good
278 line order. A line which is unimportant in level l in a line order may easily be
279 important till the nearly last level in a different order.

280 We choose line order to hopefully enforce two things:

- 281 • Make as many unimportant lines as possible. This really allows the equiv-
282 alence based improvements to kick in. The more lines are unimportant the
283 more mappings become equivalent and the faster it is to iterate through all
284 of them.
- 285 • Recognize hopeless partial mappings as soon as possible. A partial mapping
286 gets extended if the line does not break the rule that there is a one-entry
287 where it needs to be. If we map all the rows first, the rule will get broken
288 only after we start to map columns and we probably want to find out sooner.

289
290 In the program a user can either choose their own custom order or one of four
291 algorithms with different main purposes:

- 292 • AUTO - this one tries the other three line orders and chooses the one which
293 shows the best performance over some iterations on a matrix. While this
294 may sound like a good thing to use, it is only so if an initial matrix is chosen
295 and it takes a lot of time since a lot of iterations need to be made in order
296 to make a good sample. I would recommend not to use AUTO order at all
297 and instead to try all the line orders by hand with a number of iterations
298 depending on the pattern and a good initial matrix; for instance, generated
299 with a smaller number of iterations on the same pattern and with any line
300 order.
- 301 • DESC - the lines are ordered in descending order depending on the number
302 of one-entries. This follows the idea to start with the lines that are the
303 hardest to map. Note that this algorithm does poorly if there are a lot of
304 lines with the same number of one-entries (for example an identity matrix).
- 305 • MAX - it orders the lines so that the maximum number of important lines
306 throughout the levels is as small as possible. This focuses straightforwardly
307 to having many unimportant lines, which the program does not remember.
- 308 • SUM - it orders the lines so that the sum of the numbers of the important
309 lines is the smallest possible throughout all levels. The purpose is the same
310 as in the MAX order and quite often it is the case both approaches produce
311 the same order.
- 312 • TWO - it orders the lines so that the maximum number of important lines
313 in two consecutive levels throughout all the levels is as small as possible.
314 This again focuses to having many unimportant lines, which the program
315 does not remember. The constant two is chosen due to the fact general
316 pattern always stores two levels of partial mapping at a time.

317 4.1.5 Mapping approaches

318 The one thing the approaches we will introduce have in common is that they try
 319 to recognize those partial mappings that have no chance to be extended to a full
 320 mapping as early as possible.

321 While the algorithm introduced in [chapter2] finds out the partial mapping
 322 is invalid only at the time it maps two lines having a one-entry at their inter-
 323 section to two lines having a zero-entry at the intersection, different approaches
 324 try to reveal the fact we would end up in the situation earlier by checking more
 325 conditions.


326 Imaging this as the pattern, where $0, \dots, 8$ are indices of the lines:

	5	6	7	8
0	1	1	1	1
1	0	1	1	0
2	1	1	0	0
3	0	1	0	1
4	1	1	1	0

327 We are in a situation when only lines 0, 3 and 7 are mapped and line 6 is
 328 currently being mapped. Here are some mapping approaches:

329 Enough one-entries


330 In the situation above, we not only want to check there are one-entries at the
 331 intersections of line 6 with lines 0 and 3, but we also check if there are enough
 332 one-entries in lines between where lines 0 and 3 are mapped so that there is a
 333 hope we can map lines 1 and 2 there and if there is a one-entry below the line
 334 where line 3 is mapped so we can map line 4 there later.



335 Recursive mapping


336 While we were only testing whether there are enough one-entries in between
 337 already mapped lines in the previous approach, this time, we also check whether
 338 those one-entries can be used for the lines that are intended to be mapped there.
 339 For example, when we check there is a one-entry to be used for line 1 later, we
 340 also check the line 1 can be mapped to that row, which in this situation means

341 to also check there is a one-entry at the intersection with the line to which the
 342 line 7 is mapped.



343 **Orthogonal bounds**

344 When adding line 6, we check whether there are enough one-entries on the already
 345 mapped lines orthogonal to line 6, in between line 6 and the closest mapped lines
 346 next to line 6. The same idea as in “Enough one-entries”, but checking different
 347 lines.



348 **Usage**

349 These restrictions on the added lines are not a fixed part of the program. A user
 350 can decide which approaches they want to use in the configuration file. This is
 351 due to the fact that there is no right path to choose.

352 In the testing that was done for a fixed pattern, we found out it is useful to
 353 use all the mentioned restrictions when generating a matrix of size 100×100 , as
 354 it turned out to be much faster than without the restrictions. On the other hand,
 355 in the same test for a generated matrix of size 500×500 , it was much better not
 356 to use any of those restrictions.

357 **4.1.6 Using the whole structure in the next iteration**

358 It may seem like a good idea to remember all the partial mappings, to propagate
 359 them to the next iteration of the MCMC process and alter them depending upon

360 the change.

361 This really can be done. If the change is from zero-entry to one-entry, for each
362 partial mapping we already have we want to try to extend it by the line that just
363 changed and if we manage to do that we then try to extend it to a full mapping
364 in all possible ways if it is a new mapping or do nothing if it is equivalent with a
365 partial mapping of higher level. This can be easily done by means already used
366 in the standard algorithm and may lead to a better performing one.

367 However, if the element gets changed from one-entry to zero-entry we need
368 to go through the partial mappings and delete those that used the currently
369 changed one-entry. This gets a bit messy as we can no longer forget unimportant
370 lines and moreover for each partial mapping we need to remember how many
371 partial mappings of the previous level can be extended to that one, to delete that
372 mapping from the list if there are no longer any mappings extensible to that one.

373 This can all be done, but the whole thing comes with three huge inconven-
374 niences:

- 375 • Memory consumption - there can be a LOT of partial mappings and we
376 need to remember them all. Of course we can still use the equivalence but
377 we need to remember mappings of all levels.
- 378 • The change from one-entry to zero-entry is no longer for free. If this change
379 is done, we already know the pattern is not contained in M , but we still
380 need to do a lot of work to change the structure in order to use it in the
381 next iteration.
- 382 • Reverting - if the change is unsuccessful (the pattern is contained) we need
383 to revert the change which means to completely revert all changes we did to
384 the list of partial mappings. This can be either done by making a backup
385 copy of the whole structure and override the structure if needed, which again
386 is very costly as the structure is huge, or we can remember what partial
387 mappings are new and we go through all partial mappings and remove those
388 new ones. This again means to iterate through the big structure one more
389 time.

390 After realizing these issues it no longer looks useful to me to implement this
391 version of the algorithm.

392 4.2 Walking pattern

393 While the brute force implementation of an avoid algorithm for a general pattern
394 was improved heavily, the algorithm for a walking pattern is very fast in its nature
395 and cannot be much better. Or can it be?

396 4.2.1 Using the last changed position

397 As the MCMC process works it always changes one element of the big matrix and
398 asks whether it still avoids the pattern. If it does not and we know that before the
399 change it did, we are sure the changed element is a part of the pattern. Knowing
400 that and using the same inductive proof as we did in the proof of correctness of

401 the avoid algorithm (see [chapter2]) it is sufficient to only recompute the part of
402 the inner structure under the changed element and check if the last entry of the
403 pattern can be found there.

404 Not only that. We also know, using the fact the structure was completely
405 correct before the change, that if the values of both c_v and c_h of an element did
406 not change, the element won't cause the element underneath it to change and we
407 no longer have to recompute the other parts of the structure.

408 To use both these facts we replace the cycle through the diagonals by a simple
409 queue, starting at the position of the last changed element and putting more
410 positions in if the values of c_v or c_h are different than they were before. The
411 function ends either when the pattern was discovered or when the queue becomes
412 empty.

413 4.2.2 Lazy avoid

414 Lazy avoid is a variant of avoid function used when the MCMC parallelism is
415 chosen. While all the other types of patterns have a trivial implementation of
416 revert function, when using the walking pattern the inner structure needs to be
417 modified even when reverting. The MCMC parallelism turned out to work much
418 better if the revert calls are handled by the main thread (more in [chapter4]) and
419 it requires the function to run as fast as possible so the other threads are not
420 blocked by the call for too long. That is a reason why functions lazy revert and
421 lazy avoid were created.

422 The avoid function expects the inner structure of the walking pattern (see
423 [chapter3]) to be in a valid state and that requires some effort. To make lazy
424 revert the fastest possible, we postpone the work until the next call of lazy avoid,
425 meaning that lazy avoid then needs to do more things at once. It is no longer
426 sufficient to only compute the submatrix under the position changed last as we
427 did above, but it needs to also compute changes in the positions changed in those
428 lazy revert calls that are postponed.

429 We discuss several approaches, starting with the easiest one and ending with
430 the one that is fast and used in the final implementation.

431 Recompute the whole structure every time

432 The easiest way how to implement lazy avoid would be to always recompute the
433 whole inner structure. In that case we do not worry which positions are correct
434 and which are not, because every time we find the pattern, we recomputed all
435 the entries that form it, so we know it really is there. On the other hand, if we
436 manage to recompute the whole structure without finding the last entry of the
437 pattern, it just is not there.

438 The issue is efficiency. If the whole structure was correct and there was a
439 change of the last entry of the matrix it is sufficient to only recompute that one
440 entry. Instead we recompute a possibly very big structure. This results in a very
441 bad performance negating the advantage of parallel computation.

442 **Recompute only a part of the structure diagonal by diagonal**

443 A simple improvement would be to remember the changes done in previous calls
444 of lazy revert and together with the change done in lazy avoid call only recompute
445 the part of the structure that has possibly altered.

446 This gets a bit tricky when lazy avoid call actually discovers the pattern
447 because we cannot be sure the rest of the structure is in a correct order. It is
448 still possible to remember some horizontal, vertical and diagonal bounds and use
449 them to restrict the recomputed part of the matrix. The improvement is not that
450 significant though and we can do better.

451 **Queue of positions to recompute**

452 A different approach is closer to the one used in a standard avoid function. In-
453 stead of going through diagonal one after another, we have a queue of entries-to-
454 recompute. It is no longer sufficient to have a standard queue since in different
455 calls of lazy revert/avoid we can possibly change an entry of different priority (the
456 higher the more important) so we need to have some kind of a priority queue.
457 That is exactly what I tried.

458 Using `std::priority_queue` the function had no more problems with recompute-
459 ing the entries that were not influenced by the changes and used all the benefits
460 mentioned in the previous section. But the container does not come for free and
461 in the end I found out the price I paid for the operations on the priority queue
462 made the whole implementation comparably slow as in the previous attempt.

463 **Two leveled queue of positions to recompute**

464 The final solution comes with the same idea, but a different storage type. As
465 the priority depends upon a diagonal (two entries on the same diagonal can be
466 recomputed in any order) we only remember a priority queue of diagonals and an
467 array of diagonals saying whether a diagonal is already a member of the priority
468 queue. As far as the entries are concerned for every diagonal we have a `std::vector`
469 of entries-to-recompute as well as an array saying whether an entry is already a
470 member of the vector. So finally it is the case that the storage used is not only
471 good theoretically but as the numbers say, also practically. [reference to a table
472 of measurements or something]

473 **4.3 Parallel computing**

474 **4.3.1 MCMC parallelism**

475 The first approach makes the MCMC generator parallel and it can be used with
476 all types of patterns. While the serial MCMC generator changes one element in
477 the generated matrix and checks whether it still avoids forbidden patterns, the
478 parallel one counts on the fact at some point it is unlikely a change does not
479 create a mapping of the pattern. So as it expects almost every change of one
480 element to form a pattern, it tries to deal with those changes as fast as possible.

481 When the parallel version of MCMC generator is chosen and it is assigned n
482 threads, it creates $n - 1$ copies of the generated matrix and assigns one thread,

483 called worker, to each of them. The last thread, which we call the main thread,
484 makes one change of a bit in each matrix and makes the corresponding worker
485 check the avoidance.

486 The job of a worker is, when woken up, only to check if its copy of the
487 matrix still avoids the pattern when one bit is changed. On the other hand all
488 synchronization is left to the main thread. We still want the generator to satisfy
489 the conditions we have for the Markov chain (more in [chapter1]) in order to
490 generate a random matrix. To achieve that we have an ID for each task it assigns
491 and it always waits for the task with the lowest ID to end and depending on its
492 result it continues:

- 493 • success: If the task with the lowest ID is successful, we can propagate
494 the changed bit into the generated matrix. Also, since all other threads ex-
495 pected this one to fail, they compute a change that will never be propagated
496 to the generated matrix and therefore is useless. What we do with that is we
497 stop all the workers that compute a task with higher ID, revert the changes
498 they made on their matrices and we synchronize the change computed by
499 the successful thread. After all of this is done, all the matrices are the same.
- 500 • fail: If the task with the lowest ID is not successful, there is no change to
501 propagate to the generated matrix. This still makes a valid iteration of the
502 MCMC process though and since all the other threads expected this one to
503 fail, they compute correct changes.

504 In both cases a new task is then assigned to the worker.

505 **Speculative computing**

506 It may easily happen that a task not having the lowest ID ends first. In that
507 case we could just wait until it has the lowest ID and process it later. This
508 would be pretty costly though. Instead we process the task immediately, just as
509 described for the task with the lowest ID, but we don't propagate the changes
510 to the generated matrix and we do not stop the workers processing a task with
511 lower ID. This gets a bit tricky when the task succeeds. Not only we revert
512 all the changes computed by tasks with higher ID and synchronize the changes
513 computed by the worker (we would do that even if waiting for the task to have
514 the lowest ID), but it might happen a task with even lower ID succeeds as well.
515 This leads to a need of reverting the synchronizations we made. Luckily this is
516 the only precarious situation we may encounter and it is not that hard to deal
517 with it.

518 The way we deal with these inconveniences is described in [chapter5] and
519 should be clear from the code itself.

520 **Reverting and synchronizing in the main thread**

521 The speculative computing discussed above is not the only improvement we can
522 make. It turns out that it is costly to wake a thread so it computes a trivial
523 function, sets a few atomic variables and falls asleep again. This happens a lot.
524 Every time a call of avoid succeeds it makes other workers revert their changes
525 and synchronize the successful change, which are both trivial functions.

526 To workaround this problem we make a theoretically bad decision which has
527 very nice practical results. All the reverts and synchronizations are computed by
528 the main thread instead of by an appropriate worker. There is no problem with
529 concurrency because the worker is always asleep when a revert or sync calls are to
530 be assigned and using the fact those calls are really trivial, it does not make the
531 rest of threads wait for the main thread for too long while it computes changes.

532 5. Technical documentation

533 In this chapter, we cover those parts of the algorithm that may be hard to un-
534 derstand just from the code. This only means functions that are technically hard
535 - functions with unexpected dependencies, side effects and so on. Algorithmic
536 difficult tasks are explained in [chapter4].

537 5.1 General pattern

538 The general pattern class contains a lot of function. Most of them are easy to
539 follow and they all should be commented enough in the code. The only part
540 which deserves more attention is the constructor.

541 5.1.1 Construction

542 In the constructor of a general pattern, there are a few function that are easy in
543 nature but as they somehow use each other it is hard not to lose track of their
544 dependencies and results. In order to make this part of the code, which is a very
545 important part indeed, more understandable, we go through the constructor and
546 explain all that is happening in the order it is happening in.

547 Storing the pattern

548 The first thing, which is done right after initialization of variable, is storing the
549 pattern. Instead of storing the pattern in a `Matrix<bool>`, I decided to use to
550 store lines into a number, where in the binary coding a one-entry in the position
551 i means there is a one-entry in the line at the intersection with i -th orthogonal
552 line. This comes handy when computing lines orders. At the same time we also
553 find those lines that are empty (more in [chapter4]) and remember them, because
554 we do not have to map them at all.

555 Choosing the line order

556 After that we need to choose the right line order (again more in [chapter4]).
557 To compute MAX or SUM order we just use a brute force algorithm that checks
558 sequences of line adding and for each it computes how many lines are unimportant.
559 Then it just chooses the order which is the best in chosen metric.

560 What to remember

561 In the next step, we find what do we need to remember in each level of partial
562 mappings with respect to chosen order. As mention earlier, for MAX or SUM
563 order it is already computed when finding the order, but for other variants of
564 orders it is not, so we just compute it every time. What to remember is based
565 on the equivalence introduced in [chapter2] and the decision not to remember
566 unimportant lines (which we explained in [chapter4]).

567 **Parallel bound indices**

568 Now comes the hardest to follow part - precomputing the indices for searching for
569 parallel bounds. The idea behind is simple. When we are adding a new line and
570 we already have a partial mapping, it restricts to where we can add the line. For
571 example, if there are three rows in the pattern and the rows 1 and 3 are mapped,
572 the second one need to be mapped in between those two. The question is, where
573 are those two lines mapped to? First we add in a chosen order and second we do
574 not remember all lines, as some are unimportant. What do we want is to have
575 a instant access to the index of the line, which bounds added line, in the partial
576 mapping so we do not need to compute the index over and over again. That is
577 exactly what gets computed when the function “find_parallel_bound_indices” is
578 called. The series of other function calls follows just because we compute the
579 indices for all added lines in the order in which they are going to be added.

580 **Extending order**

581 The last function, “find_extending_order” just specifies how the next partial map-
582 ping will look like a from where from the previous mapping the values will be
583 copied. Again, unimportant lines play their role here and it may easily be the
584 case from a partial mapping storing k lines, after mapping one more line, we end
585 up with a partial mapping only storing $k - 1$ lines, because two lines become
586 unimportant by adding the line.

587 **5.2 Parallel computing**

588 **5.2.1 MCMC parallelism**

589 While the idea behind MCMC parallelism is described in [chapter4.3] and the
590 code is heavily commented, the work done by the main thread may still be hard
591 to understand.

592 Let I be the ID the process is currently waiting for, that is, the lowest ID of
593 a task that is being tested by a worker. In a structure called “queue” (which is
594 `std::vector<std::deque>`) each worker has a queue of tasks related to it. In the
595 queue, there are tasks that are either being computed or have been computed.
596 The history of tasks is needed to allow reverting changes that should have not
597 happen when the main thread encounters a different successful task with lower
598 ID. There is no need to have a complete history of all tasks computed. There
599 are only those tasks, that have higher ID than I or have lower ID, but those are
600 going to be removed from the “queue” as soon as possible. The name “queue” is
601 not random, it describes the order in which the tasks are being stored - the tasks
602 with lower ID have been inserted earlier and therefore they are at the bottom.

603 Now that we know the most important structure let's see how the main thread
604 works with that and what are the situations.

- 605 • **pop_front:** The main thread deletes the first tasks (the one with the lowest
606 ID) if one of two things happen:
 - 607 – The ID of the task being deleted is equal to I . That means the change
608 computed by the task is being propagated to the generated matrix and

609 there is no need to remember the task anymore. This also increases I ,
610 not necessarily by one.

611 – The ID of the task being deleted is less than I . This situation is due
612 to synchronization. The worker was supposed to synchronize a task
613 computed by a different worker that did not have the lowest ID at the
614 time. Therefore the task needs to be in the list of tasks so we can
615 revert it if needed. If there is no need to revert it and the lowest ID
616 gets greater or equal to the ID of the task, we can just delete it from
617 the “queue”.

618 • `pop_back`: There is only one reason to delete tasks from the end of the
619 “queue” and that is reverting. Imagine there is a task with id J at the
620 end of the “queue”. Now a different worker computes a task with lower
621 ID and finds out the change is successful. This means the task J won’t
622 propagate to the generated matrix and there is no use for it. If it is still
623 being computed, we cannot do much about it, so we just tell the worker to
624 stop computing and deal with it later. If the task is finished, we need to
625 revert it, but only in case the task was successful, because if it was not, it
626 had already been reverted by the worker. So we revert the task if needed
627 and we can just delete it from “queue” as it will never be used.

628 • `emplace_back`: The main thread only inserts new tasks to the end of the
629 “queue” and there are two reasons to insert:

630 – Worker is assigned a completely new task to check the avoidance. In
631 this situation the task is given a new, globally highest ID and we add
632 the task at the end of the list.

633 – The second reason to insert into “queue” are, again, synchronizations.
634 The situation is the same as it was in the case, when we `pop_back` -
635 after we revert all the tasks in the list, we need to synchronize changes
636 that forced reverting and if their ID is not lower or equal to I , we need
637 to add them to the list so they can be reverted if needed.

638 5.3 Library interface

6. User documentation

In the last chapter of the thesis we show how to make the program generate random matrices or to test whether a certain matrix avoids a given forbidden pattern. In the first section we find out how to set a configuration file. After that we go through necessary input and output files and see how they are formatted.

6.1 Configuration file

In order to modify what the program computes, we use a configuration file. There is only command line options to be used and that is a path to the configuration file. If no path is inserted the configuration file is expected to be located in the same directory as the executable file is and its name is “config.txt”.

The file is a standard text file which can be modified by any text editor and is structured into four sections:

- input
- pattern
- output
- statistics

The order of the sections is not fixed and there can be additional empty lines for better readability. In each section there is a list of values that can be set either to arbitrary value or to a specific one. There is at most one command of format “option=value” per line and there might be additional white spaces surrounding the “=” sign.

If an option is set more than once, the latter value is always used. If, on the other hand, an option is not set at all, the default value is used. If there is a line encountered that sets a wrong option, for instance when the user mistypes a valid option, the line is skipped and the user gets a warning in the standard error output.

Let us provide a list of all options for each section together with their default values.

6.1.1 Input

In the first section of the configuration file we set the generating process.

- size: The size of the generated matrix. Results in $M \in \{0, 1\}^{size \times size}$.
Possible value: $s \in \mathbb{N}$
Default value: 100
- iterations: The number of iterations of the MCMC process.
Possible value: $i \in \mathbb{N}$
Default value: 10,000

- 673 • `random_seed`: The random seed for the MCMC process.
 Possible value: $s \in \mathbb{N}$
 “random” - chooses a random seed
 Default value: “random”
- 675 • `init_matrix`: A $size \times size$ matrix the MCMC process starts with.
 Possible value: *input file path*
 “zero” - a matrix containing no one-entries
 Default value: “zero”
- 677 • `parallel_mode`: choose whether to compute in parallel or serial.
 Possible value: “serial”
 “mcmc” - more iterations of the MCMC generator are tested in parallel
 “map” - more partial mappings are being extended in parallel
 Default value: “serial”
- 679 • `threads_count`: The number of threads that are going to be used if a parallel
 mode is chosen.
 Possible value: $t \in \mathbb{N}$
 -1 - chosen according to the number of cores
 Default value: 1

682 6.1.2 Pattern

683 It is time to set the options that matter the most - matrix patterns. As you
 684 generate a matrix which avoids more than just one pattern, the section [pattern]
 685 can be used multiple times, specifying one pattern for each occurrence.

- 686 • `pattern_file`: A path to a input matrix file - the pattern.
 Possible value: matrix file path
 Default value: “input.txt”
- 688 • `pattern_type`: The type of the pattern. Determines the method used for
 testing avoidance.
 Possible value: “general”
 “walking” - see [chapter3]
 “slow” - brute force algorithm for a general pattern
 Default value: “general”

691 The next options are only useful if the general pattern type is chosen. It
 692 specifies how the mappings are stored as well as what the map function tests.

- 693 • `map_one_entries`: If set to “yes”, the map function not only checks that
 already mapped lines do cross the currently added one in a one-entry if
 needed but also tests whether there is enough one-entries in between those
 already mapped lines. The result is that the map function takes a little
 longer but it recognizes hopeless mappings earlier.
 Possible value: “yes”
 “no”
 Default value: “yes”

- 699 • `map_recursion`: If set to “yes” and the `map_one_entries` is also set to “yes”,
700 the `map` function not only tests there are enough one-entries in between
701 already mapped lines on the currently added one, but for those non-mapped
702 lines crossing the current one in a one-entry, it also tests if the partial
703 mapping has one-entries everywhere they are supposed to be. Although
704 this comes with a big time consumption it can filter a big portion of partial
705 mapping that would later become unfeasible.
 Possible value: “yes”
706 “no”
 Default value: “yes”
- 707 • `map_orthogonal_bounds`: If set to “yes”, the `map` function also tests the
708 orthogonal bounds of added line - see [chapter2].
 Possible value: “yes”
709 “no”
 Default value: “no”
- 710 • `map_container`: A container in which the partial mappings are stored.
 Possible value: “set” - `std::set` (red-black tree)
 “hash” - `std::unordered_set` (hash table)
711 “vector” - `std::vector` (dynamic array)
 Default value: “hash”
- 712 • `line_order`: Choose the order in which the lines are being added to the
713 partial mapping. See [chapter2]
 Possible value: “max”
 “two”
 “sum”
714 “desc”
 “auto”
 “order file path”
 Default value: “max”

715 6.1.3 Output

716 There is no reason to generate a matrix noone will ever see. In this section we
717 make the matrix show. As the matrix can be output to console, a text file or a
718 bmp file an option in the section can be set more than once and all the settings
719 make difference.

- 720 • `matrix_output`: The generated matrix can be output as a bmp file in which
721 one-entries are black pixels and zero-entries white. To do that, the file path
722 has to have a pattern “path.bmp”. If a different path is given the file is
723 stored as a matrix text file. It can also be output into a console if “console”
724 is set. In that case it has the text format.
 Possible value: “console”
 matrix bmp file path
 matrix text file path
725 “no”
 Default value: “no”

- 726 • `performance_stats`: If the serial computation or the map parallelism is cho-
727 sen, the program can output a statistics like the percentage of avoid call
728 success, how long did one call take on average and what was the average
729 size of structures. Note that this is not very useful as a matrix statistics,
730 see the next section for that. If more patterns are set, the statistics may get
731 misleading as it also counts the cases when the first pattern is contained in
732 the matrix and the other patterns are not tested at all.
Possible value: "console"
performance file path
733 "no"
Default value: "no"
- 734 • `performance_csv_stats`: The same information as above but formatted to a
735 csv file so the data can be more easily worked with.
Possible value: "console"
csv file path
736 "no"
Default value: "no"
- 737 • `time_to_console`: Prints how long the computation took into a console.
Possible value: "yes"
738 "no"
Default value: "no"
- 739 • `patterns_to_console`: Prints all the used patterns into the console.
Possible value: "yes"
740 "no"
Default value: "no"

741 6.1.4 Statistics

742 The last section handles the options important for scientists. While generating
743 a random matrix is a great result, on its way the program can also create some
744 statistics, namely make a histogram of occurrences of one-entries in a generated
745 matrix as the MCMC iterates as well as store the matrix with the highest amount
746 of one-entries. As the process usually does not start with a random matrix, the
747 user can decide to only compute the statistics after a certain number of iterations
748 has been done and to only check a small portion of iterations, every 10th for
749 instance, as a single iteration may not make any difference and counting the
750 histogram takes time.

- 751 • `histogram_frequency`: Sets how often the histogram gets refreshed.
Possible value: $f \in \mathbb{N}$
752 0 - the histogram is not computed at all
Default value: 0
- 753 • `histogram_initial`: Sets the initial iteration of the MCMC process when the
754 histogram gets refreshed.
Possible value: $i \in \mathbb{N}$
755 Default value: 1,000

- 756 • histogram_final: Sets the last iteration of the MCMC process when the
757 histogram gets refreshed.
Possible value: $f \in \mathbb{N}$
758 -1 - the histogram is computed till the end
Default value: -1
- 759 • histogram_file: Sets where to output the histogram computed during the
760 MCMC process.
Possible value: matrix bmp file path
matrix text file path
761 "console"
"no"
Default value: "no"
- 762 • max_ones_matrix_file: Sets where to output the matrix that had the most
763 one-entries among all matrices iterated through during the MCMC process.
Possible value: matrix bmp file path
matrix text file path
764 "console"
"no"
Default value: "no"

765 6.2 File input

766 There are only two types of input files expected by the program. Either you want
767 to read a matrix file, which can be a pattern or an initial matrix, or an order file
768 that determines an order in which the lines are going to be mapped if the general
769 pattern is chosen.

770 6.2.1 Matrix file

771 A matrix file is a standard text file having the format as follows:

- 772 • 2 natural numbers specifying the number of rows and columns in this order.
- 773 • a sequence of zeros and ones of length rows×columns specifying the matrix
774 from the top left corner one row after another.

775 **Example:** 2 3
 1 0 1
 1 1 0

776 6.2.2 Order file

777 If you want to choose the order in which the lines are going to be mapped when
778 a general pattern is chosen, it is your responsibility to check that all lines that
779 need to be mapped are mapped. It is for example possible to only map three lines
780 even if the pattern consists of six lines just because there is for example no need

781 to map empty lines at all. Therefore the program does not check the validity of
782 the order and just uses it.

783 Now that the user has been warned, the format of the custom order file is
784 simple. It consist of the indices of the lines of the pattern numbered starting
785 with 0 and starting from the top row and ending with the right column.

786 One possible order for the matrix given as an example in [6.2.1] is this file:

787
$$2\ 1\ 0\ 3\ 4$$

788 First mapping the left column, the second and first row after that and finishing
789 the mapping with the middle column and the right one.

790 **6.3 File output**

791 Let us now find out what the output files look like.

792 **6.3.1 Matrix text file**

793 The matrix text file has the same format as the input one. It consists of:

- 794 • 2 natural numbers specifying the number of rows and columns in this order.
- 795 • a sequence of zeros and ones of length rows×columns specifying the matrix
796 from the top left corner one row after another.

797 The matrix is binary except for the one produced as a histogram, which can have
798 higher natural numbers and contains the number of samples as the last number.
799 If you then divide all the entries by the last number, you get a percentage of the
800 entry being a one-entry.

801 **6.3.2 Matrix bmp file**

802 For an $n \times n$ matrix the standard bmp file contains $n \times n$ pixel of black color
803 meaning a one-entry and a white color for a zero-entry. If the histogram is output
804 as a bmp file, the pixels are greyscaled and the darker a pixel is the more often
805 the entry was a one-entry during the MCMC process.

806 Conclusion

807 List of Figures

808 List of Tables

809 List of Abbreviations

810 Attachments