

Charles University in Prague
Faculty of Mathematics and Physics

BACHELOR THESIS



Stanislav Kučera

Generating random pattern-avoiding matrices

Computer Science Institute of Charles University

Supervisor of the bachelor thesis: RNDr. Vít Jelínek, Ph.D.

Study programme: Computer Science

Study branch: General Computer Science

Prague 2016

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University in Prague has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

signature of the author

Title: Generating random pattern-avoiding matrices

Author: Stanislav Kučera

Department: Computer Science Institute of Charles University

Supervisor: RNDr. Vít Jelínek, Ph.D., department

Abstract: Abstract.

Keywords: binary matrix pattern-avoiding Markov chain Monte Carlo

Dedication.

Contents

Preface	3
Introduction	5
1 Markov chain Monte Carlo	6
1.1 Markov chains	6
1.2 Markov chain for pattern-avoiding binary matrices	6
2 An algorithm for testing pattern-avoidance of a general pattern	8
2.1 Sketch of a brute force algorithm	8
2.2 Equivalent mappings	8
3 An algorithm for testing pattern-avoidance of a special pattern	10
3.1 Walking pattern	10
3.2 Dynamic program	10
3.2.1 The algorithm	11
3.2.2 Inner structure	12
3.2.3 Correctness	13
3.2.4 Generalization	13
4 Improvements to basic algorithms	14
4.1 General pattern	14
4.1.1 Improving memory consumption	14
4.1.2 Not mapping empty lines	14
4.1.3 Using the last changed position	14
4.1.4 Line order	15
4.1.5 Mapping approaches	16
4.1.6 Using the whole structure in the next iteration	18
4.2 MCMC parallelism	19
4.2.1 Example of the MCMC process for n threads	20
4.2.2 Speculative computing	20
4.2.3 Reverting and synchronizing in the main thread	21
4.3 Walking pattern	21
4.3.1 Using the last changed position	21
4.3.2 Lazy avoid	21
5 Technical documentation	24
5.1 General pattern	24
5.1.1 Construction	24
5.2 Parallel computing	25
5.2.1 MCMC parallelism	25
5.3 Library interface	26

6	User documentation	27
6.1	Configuration file	27
6.1.1	Input	27
6.1.2	Pattern	28
6.1.3	Output	29
6.1.4	Statistics	30
6.2	File input	31
6.2.1	Matrix file	31
6.2.2	Order file	31
6.3	File output	32
6.3.1	Matrix text file	32
6.3.2	Matrix bmp file	32
	Conclusion	33
	Bibliography	34
	List of Figures	35
	List of Tables	36
	List of Abbreviations	37
	Attachments	38

1 Preface

2 Theses at the faculty of mathematics and physics usually fit into one of three
3 categories:

- 4 1. Theoretical thesis
- 5 2. Experimental thesis
- 6 3. Implementation thesis

7 My thesis does not fit entirely into only one category and it does not try to. The
8 project consists of several similarly important parts which are:

- 9 • Design of the algorithm for generating a special binary matrix
- 10 • Making it run fast on inputs which are usual for researchers
- 11 • Implementing the algorithm to provide practical tool

12 None of these points would make sense alone but together the thesis may be-
13 come very useful for scientists as it provides with a process generating random
14 matrices and it is a common practice to test hypothesis on random data.

Introduction

We denote by $M \in \{0, 1\}^{n \times m}$ a *binary matrix* of size n by m , calling n the number of rows of M - the *height* of the matrix M and m the number of columns - its *width*. A *line* of a matrix is one of its rows or columns and we denote by $L(M)$ the ordered set of all lines of M . Its order is given by the natural indexing of rows and columns.

Definition 1. We say a binary matrix M contains a binary matrix P , which we call a “*pattern*”, as a submatrix, if there is a mapping $f : L(P) \rightarrow L(M)$, such that

- $l \in L(P)$ is a row of P iff $f(l) \in L(M)$ is a row of M
- $\forall l, l' \in L(P) : l < l' \Rightarrow f(l) < f(l')$ (preserves the order)
- $\forall l, l' \in L(P) : \text{if lines } l \text{ and } l' \text{ intersect and there is a one-entry at the intersection, then there is a one-entry at the intersection of } f(l) \text{ and } f(l').$

otherwise, it avoids the pattern P .

$$P = \begin{matrix} & \begin{matrix} 2 & 3 \end{matrix} \\ \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \end{matrix} \quad M_1 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix} \quad M_2 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix}$$

Figure 1: Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$ satisfies all the conditions. On the other hand, matrix M_2 avoids P as there is no such mapping.

The interesting cases are square matrices of size n by n , where n is big (going to infinity) and the size of a pattern (not necessarily square matrix) is small (constant). Even for a constant size forbidden pattern it is hard to determine the number of matrices of size n that avoid it or to characterize, what properties they have. Sometimes we consider matrices avoiding more than just one forbidden pattern, in which case we denote the set of all forbidden matrices by \mathcal{P} . When a matrix avoids \mathcal{P} , it avoids every $P \in \mathcal{P}$.

Definition 2. We denote by $\mathcal{M}_n(\mathcal{P})$ a set of all binary matrices of size n by n avoiding \mathcal{P} as submatrices. We always denote M the square binary matrix for which we test the containing and by P the pattern (if there is one) that is being tested. Moreover, we denote h the height (the number of rows) of P and w its width.

The area of pattern avoidance has been heavily studied for permutations and it also becomes more popular for their generalization - binary matrices. In most of the areas in combinatorics it is useful to explore properties of random objects and a lot of attention is directed towards random matrices when considering pattern avoidance. The goal of the work is, for given $n \in \mathbb{N}$ and set of forbidden patterns \mathcal{P} , to generate a uniformly random $M \in_R \mathcal{M}_n(\mathcal{P})$.

47 **Generating random matrix**

48 One way to get $M \in_R \mathcal{M}_n(\mathcal{P})$ is to choose a matrix of required size completely
49 at random, for such, test whether it avoids the pattern and simply repeat the
50 process until we find one, which does. However, in the most interesting cases,
51 only a small fraction of all matrices avoid the pattern and the process takes too
52 long, to be practically useful.

53 For generating random permutations avoiding forbidden pattern, a different
54 technique was introduced in Madras and Liu [2010]. It uses a randomized pro-
55 cess called Markov chain Monte Carlo, which we will abbreviate by MCMC. It
56 is an iterative process, which for a well chosen Markov chain (more in chapter 1)
57 approximates a random object. The algorithm by Madras and Liu was devel-
58 oped for permutations (permutation matrices) and it cannot be used for general
59 matrices. In section 1.2 we show how to adapt the algorithm, which will lead
60 us to a MCMC algorithm that approximates $M \in_R \mathcal{M}_n(\mathcal{P})$. To produce a good
61 approximation the process needs to do a lot of iterations and despite the fact
62 it is unknown what is the mixing time (the number of iterations required) of a
63 MCMC process, in practice, the method does better than the trivial algorithm.

64 **Testing avoidance**

65 In each step of our MCMC process we need to test whether a matrix avoids a
66 pattern. We will show a very fast algorithm that only works for a special class of
67 binary matrices (explained in chapter 3) together with a slightly less performing
68 algorithm for a general pattern, which, again, comes as a generalization of an
69 algorithm for permutations from the article by Madras and Liu and is described
70 in chapter 2.

71 In chapter 4 we improve both our algorithms and introduce a parallel version
72 of MCMC process, which further increases the performance of matrix generating.

73 In chapter 5 some technical details are explained to make reading the code
74 easier for reader and to describe user interface. The last chapter (chapter 6)
75 contains user documentation.

1. Markov chain Monte Carlo

Our goal to generate $M \in_R \mathcal{M}(\mathcal{P})$ heavily depends on the theory of Markov chains. In this work we only define useful terms and state two important theorems. If you are interested in more details, see Madras [2002].

1.1 Markov chains

Definition 3. We shall consider discrete-time Markov chain X_0, X_1, \dots , where $X_i \in \mathcal{S}$, for a finite state space \mathcal{S} and every i (number of steps). The k -step transition probabilities are:

$$p_{i,j}^{(k)} = \Pr(X_{t+k} = j | X_t = i) \quad (i, j \in \mathcal{S})$$

Definition 4. A Markov chain is said to be symmetric if $p_{i,j}^{(1)} = p_{j,i}^{(1)}$ for every pair of states i and j .

Definition 5. A Markov chain is irreducible if the chain can eventually get from each state to every other state, that is, for every $i, j \in \mathcal{S}$ there exists a $k \geq 0$ (depending on i and j) such that $p_{i,j}^{(k)} > 0$.

Definition 6. An irreducible chain has period D if D is the greatest common divisor of $\{k \geq 1 | p_{i,i}^{(k)} > 0\}$ for some $i \in \mathcal{S}$ (equivalently, for all $i \in \mathcal{S}$). A chain is called aperiodic if its period is 1. In particular, if an irreducible chain has $p_{i,i}^{(1)} > 0$ for some i , then it is aperiodic.

Theorem 1. Consider an aperiodic irreducible Markov chain with state space \mathcal{S} . For every $i, j \in \mathcal{S}$, the limit $\lim_{k \rightarrow \infty} p_{i,j}^{(k)}$ exists and is independent of i ; call it π_j . Furthermore, if \mathcal{S} is finite, then

$$\sum_{j \in \mathcal{S}} \pi_j = 1 \quad \wedge \quad \sum_{i \in \mathcal{S}} \pi_i p_{i,j}^{(1)} = \pi_j$$

for every $j \in \mathcal{S}$. That is, if we write π to denote the row vector whose entries are π_i , then $\pi P = \pi$.

Theorem 2. Suppose that an irreducible Markov chain on the finite state space \mathcal{S} is symmetric. Then the equilibrium distribution is uniform on \mathcal{S} .

In other words, the theorems together give us a guarantee that if we choose an irreducible, symmetric and aperiodic Markov chain with state space \mathcal{S} then the probability distribution of X_i converges to uniform distribution on \mathcal{S} independently of the initial state.

1.2 Markov chain for pattern-avoiding binary matrices

To generate a binary matrix $M \in \{0, 1\}^{n \times n}$ avoiding patterns in \mathcal{P} , we create a Markov chain, whose states space is $\mathcal{M}_n(\mathcal{P})$. After sufficiently many iterations (m) of MCMC process we set $M := X_m \in \mathcal{M}_n(\mathcal{P})$. We always begin with an initial matrix X_0 and the process looks like this:

- 104 1. For $i := 1, 2, \dots, m$:
- 105 2. Set $X_i := X_{i-1}$.
- 106 3. Choose $r \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 107 4. Choose $c \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 108 5. Flip the bit at $X_i[r, c]$.
- 109 6. If X_i contains \mathcal{P} , flip the bit back.

110 If the process starts with a matrix X_0 that avoids \mathcal{P} , then after every step it
 111 still avoids \mathcal{P} . Note that an iteration does not change the matrix if the condition
 112 6 is satisfied. We need to show the Markov chain we presented meets all the
 113 conditions of both theorems:

114 **Symmetry**

115 Imagine a sequence of bits flipping which changes the i -th matrix to j -th one.
 116 The reversed order of the same sequence changes the j -th matrix to the i -th one.

117 **Irreducibility**

118 As the steps go, it is easy to see we can with non-zero probability create any
 119 matrix $M \in \mathcal{M}_n(\mathcal{P})$ from the zero matrix $0_n = 0^{n \times n}$ by choosing the one-entries
 120 of M . When we can get from 0_n to M by a sequence of flip changes, the reversed
 121 sequence is a sequence of steps from any matrix $\in \mathcal{M}_n(\mathcal{P})$ to 0_n . Thus the Markov
 122 chain is irreducible.

123 **Aperiodicity**

124 The Markov chain is irreducible so it suffices to show that there is an i for which
 125 $p_{i,i}^{(1)} > 0$. Clearly there is a matrix for which there is at least one bit that cannot
 126 be flipped without creating a pattern and this forces $p_{i,i}^{(1)} > 0$.

127 2. An algorithm for testing 128 pattern-avoidance of a general 129 pattern

130 In this chapter and chapter 3 we show algorithms for testing whether a pattern
131 P is contained in a square binary matrix M .

132 We begin with a very basic algorithm, which we then improve a lot to get a
133 fast algorithm for testing avoidance of a general pattern.

134 2.1 Sketch of a brute force algorithm

135 Let $L = (l_1, l_2, \dots, l_{w+h})$ be a permutation of lines (rows and columns) of the
136 pattern P and $k \in [w+h]$. *Partial mapping of level k* of lines of P is a function
137 f from $L' := \{l_1, l_2, \dots, l_k\} \subseteq L$ to lines of the big matrix M satisfying two
138 conditions:

- 139 • Both $l' \in L'$ and $f(l')$ are rows or they are both columns.
- 140 • If $l' \in L'$ and $l'' \in L'$ are both rows or columns and $l' < l''$, then $f(l') < f(l'')$.
141 This means partial mapping keeps the order of the lines.
- 142 • If $l' \in L'$ is a row of P and $l'' \in L'$ is a column of P and there is a one-entry
143 at the intersection of l' and l'' , then there is a one-entry at the intersection
144 of $f(l')$ and $f(l'')$.

145 The basic algorithm we use goes as follows. First it maps l_1 to all possible lines
146 of M , creating partial mappings of $\{l_1\} \subseteq L$. For $k = 2, \dots, w+h$ it takes each
147 partial mapping from the previous iteration and extends it by adding line l_k to
148 the partial mapping in all possible ways. If we manage to map all the lines of P ,
149 then M does not avoid it and if at some point there are no partial mappings to
150 extend it means M avoids P .

151 The algorithm can be improved in two ways. Firstly, we can try to recognize
152 unextendable partial mappings earlier than at the moment a line can no longer be
153 mapped, for example by counting whether there is enough one-entries in between
154 already mapped lines (more in subsection 4.1.5). Secondly, which is going to be
155 fundamental for us, we can try not to remember more copies of different mapping,
156 which can be extended in the same way.

157 2.2 Equivalent mappings

158 There is no need to remember two different mappings if they can be both extended
159 exactly the same way as our function is only supposed to check whether a pattern
160 can be mapped to a big matrix not to find all such mappings.

161 **Definition 7.** We call a line l of a pattern important in a partial mapping if one
162 of the conditions is met:

- 163 • An adjacent line of the pattern has not been mapped yet.
- 164 • There is a one-entry on the line l at the intersection with line l' that has
- 165 not been mapped yet.
- 166 . Otherwise the line is unimportant in the mapping.

167 Whether a line is important or not only depends on the permutation, we have
 168 chosen for partial functions, because that is what forces the order in which we
 169 map lines.

170 At the beginning, when no line is mapped, all lines are important. After some
 171 lines get mapped, a line can become unimportant in the partial mapping as all
 172 lines that bound it are in the mapping as well. If a line is unimportant in a partial
 173 mapping of some level, it will stay unimportant in all extensions of the mapping
 174 we can find.

175 **Definition 8.** We say two partial mappings of the same level are equivalent if
 176 all important lines in the mapping of that level are mapped to the same lines of
 177 the big matrix in both mappings.

$$\begin{array}{c}
 \begin{array}{c}
 0 \\
 1 \\
 2 \\
 3
 \end{array}
 \begin{array}{c}
 4 \quad 5 \quad 6 \quad 7 \\
 \left(\begin{array}{cccc}
 1 & 1 & 0 & 0 \\
 0 & 1 & 1 & 1 \\
 0 & 1 & 0 & 0 \\
 0 & 1 & 0 & 1
 \end{array} \right)
 \end{array}
 \end{array}
 \quad
 \begin{array}{c}
 \begin{array}{c}
 0 \\
 1 \\
 2 \\
 3 \\
 4
 \end{array}
 \begin{array}{c}
 5 \quad 6 \quad 7 \quad 8 \quad 9 \\
 \left(\begin{array}{ccccc}
 1 & 1 & 1 & 1 & 1 \\
 0 & 1 & 1 & 0 & 1 \\
 1 & 1 & 0 & 0 & 0 \\
 0 & 1 & 1 & 0 & 1 \\
 1 & 1 & 1 & 0 & 1
 \end{array} \right)
 \end{array}
 \end{array}
 \quad
 \begin{array}{c}
 P= \\
 M=
 \end{array}$$

Figure 2.1: An example showing unimportant line and equivalent mappings.

178 For P and M , matrices in Figure 2.1, in partial mapping $f = \{(1, 1), (2, 2), (3, 4), (5, 6)\}$
 179 line 2 is unimportant because both lines 1 and 3 are mapped and so is line 5 -
 180 the only line to intersect line 2 in a one-entry. Line 3 is important, because there
 181 is line 7 intersecting it in one-entry, which is not mapped.

182 In the same situation as above, consider a different partial mapping $f' =$
 183 $\{(1, 1), (2, 3), (3, 4), (5, 6)\}$, which is a mapping of the same level as f and only
 184 differs from f in mapping line 2. The line 2 is unimportant and by the definition
 185 of equivalent partial mappings, f and f' are equivalent. The idea behind this
 186 notion is simple. It is not important where we map line 2, because it does not
 187 restrict where we can map any other line that has not been mapped yet. This
 188 means that if a partial mapping f can be somehow extended, the equivalent
 189 partial mapping f' can be extended in the same way; therefore, it is sufficient to
 190 only extend one of them in order to find one full mapping. Note that it would
 191 be also sufficient to only extend one of the partial mappings if we were looking
 192 for all full mappings, but, in that case, we would need to keep the information
 193 about where the unimportant lines were mapped to.

3. An algorithm for testing pattern-avoidance of a special pattern

In the previous chapter, we have seen an algorithm for a general forbidden pattern which, using some heuristics, runs pretty fast. In this chapter, we introduce a special kind of a pattern, satisfying additional conditions, for which we can produce much faster algorithm.

3.1 Walking pattern

We call the specific pattern a walking pattern. The additional condition we want the pattern to satisfy is that there is a walk from one corner to the opposite one and all the one-entries of the pattern are contained on the walk.

Definition A **walk** in a matrix is a sequence of some of its entries beginning in the top left corner and ending in the bottom right one. If an entry at the position $[i, j]$ is in the sequence, the next one is either $[i + 1, j]$ or $[i, j + 1]$. Therefore, the length of an arbitrary walk is equal to $w + h - 1$.

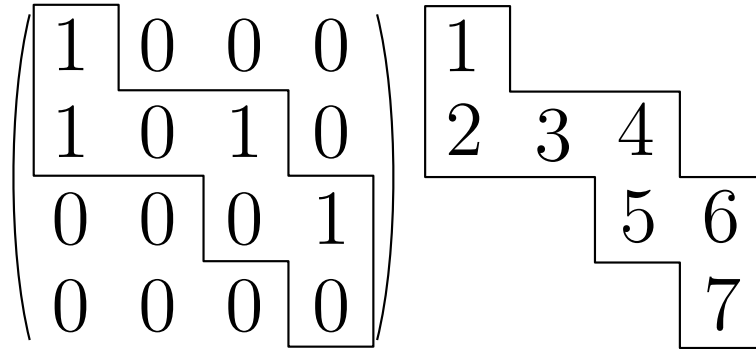


Figure 3.1: An example of a walk and the order of its entries.

In Figure 3.1 you can see a matrix that is a walking pattern as all the one-entries are included in a walk. Not all entries of a walk need to be one-entries though.

It can be shown a walking pattern with a walk is exactly a matrix avoiding a forbidden pattern

$$\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

3.2 Dynamic program

Now that we know what walking pattern is, we show an algorithm deciding whether such a pattern P is contained in a big matrix M or not.

215 The pattern P is a walking pattern so there is a walk containing all the one-
 216 entries of the pattern. We choose one such walk arbitrarily and index its entries
 217 $e_1, e_2, \dots, e_{h+w-1}$ starting from the beginning of the walk. For each entry of the
 218 walk we remember whether its value is one or zero and whether the walk continues
 219 from the entry vertically, in which case we call it a **vertical entry** or horizontally,
 220 calling it a **horizontal entry**.

221 For an element e of M at the position $[i, j]$, the matrix $M_{\leq e}$ is a $(i+1) \times (j+1)$
 222 submatrix of M consisting of rows with the index smaller than or equal to i and
 223 columns with the index smaller than or equal to j . The element e then lies in
 224 the bottom right corner. Similarly, $M_{\geq e}$ is a $(n-i) \times (n-j)$ submatrix of M
 225 consisting of rows with the index greater than or equal to i and columns with
 226 index greater than or equal to j . The element e is its first element.

227 To determine whether P is contained in M we find out for each element e of
 228 M what is the longest part of the pattern that can be found in $M_{\leq e}$. If there is
 229 an element for which we manage to find the last entry of the pattern, the pattern
 230 is contained in the matrix; otherwise, it is avoided.

231 For each element e of M at the position $[i, j]$ we remember two numbers. The
 232 number $c_v(e)$ says what is the longest part of the walk in $M_{\leq e}$ with the last entry
 233 in j -th column and being a vertical entry. The number $c_h(e)$, symmetrically, says
 234 what is the longest part of the walk in $M_{\leq e}$ with the last entry in i -th row and
 235 and being a horizontal entry.

236 An observation we make is that if we have a fixed element e of M and any
 237 other element e' above e in the same column then $c_v(e')$ is less than or equal to
 238 $c_v(e)$. This means that for e we can find the maximum part of the pattern ending
 239 in the column of e and continuing vertically by looking only to elements in that
 240 column above e and since this is true for all of them, it is sufficient to only check
 241 the value of the element right above e (at the position $[i-1, j]$). Similarly the
 242 argument goes for the value of c_h in horizontal way.

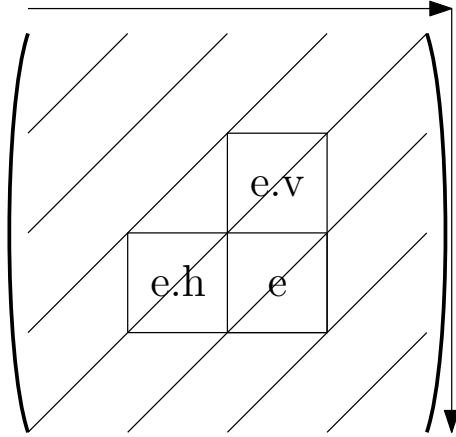
243 The algorithm iterates through diagonals. A diagonal in this matter of speak-
 244 ing is a subset of elements of M , such that all elements have the same sum of
 245 their coordinates. For example, zero diagonal only consists of an element $[0, 0]$,
 246 the first diagonal contains elements $[0, 1]$ and $[1, 0]$, and so on.

247 For simplicity, in the pseudo-code below we do not deal with elements outside
 248 M (like $-1, 0$) explicitly. Instead for those elements we just assume the values of
 249 c_v and c_h are always equal to zero.

250 3.2.1 The algorithm

- 251 1. For $d = 0, \dots, w + h - 1$
- 252 2. For e element of d -th diagonal at the position $[i, j]$
- 253 3. $e_v := [i - 1, j]$
- 254 4. $e_h := [i, j - 1]$
- 255 5. $c_v(e) := c_v(e_v)$
- 256 6. $c_h(e) := c_h(e_h)$

257 7. If $w_{c_v(e)+1}$ can be mapped to e
 258 8. If $c_v(e) + 1 = w + h + 1$
 259 9. Terminate - M contains P as a submatrix
 260 10. If $w_{c_v(e)+1}$ is a vertical entry
 261 11. $c_v(e) := c_v(e) + 1$
 262 12. Else
 263 13. $c_h(e) := \max\{c_h(e), c_v(e) + 1\}$
 264 14. If $w_{c_h(e)+1}$ can be mapped to e
 265 15. If $c_h(e) + 1 = w + h + 1$
 266 16. Terminate - M contains P as a submatrix
 267 17. If $w_{c_h(e)+1}$ is a vertical entry
 268 18. $c_v(e) := \max\{c_v(e), c_h(e) + 1\}$
 269 19. Else
 270 20. $c_h(e) := \max\{c_h(e), c_h(e) + 1\}$



271 3.2.2 Inner structure

272 To run properly the algorithm needs two structures. The first one is a description
 273 of the walk, which is just an array of the values of its entries as well as the
 274 information whether the entry is vertical or horizontal. The second structure is
 275 a matrix of the values c_v and c_h as described above.

276 3.2.3 Correctness

277 We need to show that the values of c_v and c_h are always correct for the recomputed
278 elements after at the end of the function. We proceed by induction.

279 For the first element it is definitely true since there can be only the first entry
280 of the pattern mapped and we check just that.

281 When we compute an element e of a computed diagonal d , by induction as-
282 sumption all the diagonals $d' < d$ are correctly computed. In particular, the
283 values are correct in the diagonal $d - 1$. To compute the correct values of e , we
284 use the values of two element on the diagonal $d - 1$: e_v , which is right above e and
285 e_h , which is the first element to the left of e . If e_v or e_h are outside the matrix
286 then from that direction we cannot expect to find anything more than just the
287 first entry of the pattern and that is what we check for.

288 Let v be the true length of the longest part of P in $M_{\leq e}$ continuing vertically
289 in the same column as e . Now if e itself is not an entry of that part of the pattern,
290 it is a different element e' in the same column. But then the value of $c_v(e')$ is
291 correctly computed by the inductive hypothesis and it is copied to all element
292 underneath. Especially e_v gets the value and the algorithm copies the value from
293 it to e . On the other hand if e is an entry of the part of the pattern we work with,
294 it is the last entry. The entry right before the last one needs to be mapped to
295 the same row or column; therefore, either e_v or e_h contain the part of the pattern
296 shorter by one and the algorithm extends it to a correct value.

297 3.2.4 Generalization

298 The same algorithm, just rotated by 90 degrees, can be also used for a pattern
299 where all one-entries are contained on a walk from top right corner to the bottom
300 left one. Indeed the program uses it and if given a walking pattern it determines
301 by itself which walk it is.

302 On the other hand a direct generalization for a general pattern does not work.
303 While we can index all entries of the pattern, when trying to map a certain w_k
304 to an element it is not sufficient to just check whether w_l is above and w'_l to the
305 left from the element.

$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & \boxed{1} \\ 1 & \textcircled{1} & \blacklozenge 1 \end{pmatrix} \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & \boxed{1} \\ 0 & 1 & \textcircled{1} & \blacklozenge 1 \end{pmatrix}$$

306 In the picture, let the matrix on the left side be the pattern P and let M
307 be the other matrix. The entry in the square can be mapped to the element in
308 the square and the same holds for entries in the circle but it is not a sufficient
309 condition for the entry in the triangle to be mapped to the element in the triangle.

310 4. Improvements to basic 311 algorithms

312 4.1 General pattern

313 4.1.1 Improving memory consumption

314 As the algorithm for testing avoidance of a general pattern was described in
315 [chapter2], it creates all possible partial mappings and checks whether at least
316 one can be extended to a full mapping. Note that to compute all the partial
317 mappings of some level l , it only uses mappings of level $l - 1$; therefore, it is
318 enough to only store partial mappings of two levels in memory.

319 In [chapter2] we also introduced the idea of (un)important lines for a partial
320 mapping of level l and equivalence based on not using unimportant lines at all -
321 as they are fully bounded by other already mapped lines. When a line becomes
322 unimportant it stays unimportant till the end of the run; as a result, we can forget
323 where we mapped those lines to save memory. This is not as big of a deal as the
324 previous observation was but note there are cases of patterns in which each line
325 becomes unimportant just two levels after it gets added.

326 4.1.2 Not mapping empty lines

327 An empty line is a row or a column that does not contain any one-entries. Such
328 a line can be mapped to any line and if the algorithm leaves space for it (which
329 it does), we do not need to map it at all.

330 4.1.3 Using the last changed position

331 As the MCMC process works, it always changes one element of the big matrix
332 and asks whether it still avoids the pattern. If it does not and we know that
333 before the change it did, we are sure the changed element $[r, c]$ is a part of the
334 pattern. It is hard to use this fact in the algorithm. It just maps one line after
335 another and we do not know at the beginning to which line the changed position
336 lines should be mapped.

337 What we can do is to enforce that neither the r -th line nor the c -th one
338 get skipped. We will only look at the restriction for rows. The restrictions for
339 columns are symmetrical. There are three situations we want to avoid:

- 340 • The first row of P is mapped under the r -th row. This prevents any other
341 row to be mapped to r -th one and we don't want that.
- 342 • The last row of P is mapped above the r -th row. This again prevents any
343 other row to be mapped to r -th one.
- 344 • Two adjacent rows $l, l + 1$ of P are mapped to $L < L'$ respectively and
345 $L < r < L'$ which leaves no other row to be mapped to r .

346 4.1.4 Line order

347 An important thing, if we want the algorithm to run fast, is to choose a good
348 line order. A line which is unimportant in level l in a line order may easily be
349 important till the nearly last level in a different order.

350 We choose line order to hopefully enforce two things:

- 351 • Make as many unimportant lines as possible. This really allows the equiv-
352 alence based improvements to kick in. The more lines are unimportant the
353 more mappings become equivalent and the faster it is to iterate through all
354 of them.
- 355 • Recognize hopeless partial mappings as soon as possible. A partial mapping
356 gets extended if the line does not break the rule that there is a one-entry
357 where it needs to be. If we map all the rows first, the rule will get broken
358 only after we start to map columns and we probably want to find out sooner.

359
360 In the program a user can either choose their own custom order or one of four
361 algorithms with different main purposes:

- 362 • AUTO - this one tries the other three line orders and chooses the one which
363 shows the best performance over some iterations on a matrix. While this
364 may sound like a good thing to use, it is only so if an initial matrix is chosen
365 and it takes a lot of time since a lot of iterations need to be made in order
366 to make a good sample. I would recommend not to use AUTO order at all
367 and instead to try all the line orders by hand with a number of iterations
368 depending on the pattern and a good initial matrix; for instance, generated
369 with a smaller number of iterations on the same pattern and with any line
370 order.
- 371 • DESC - the lines are ordered in descending order depending on the number
372 of one-entries. This follows the idea to start with the lines that are the
373 hardest to map. Note that this algorithm does poorly if there are a lot of
374 lines with the same number of one-entries (for example an identity matrix).
- 375 • MAX - it orders the lines so that the maximum number of important lines
376 throughout the levels is as small as possible. This focuses straightforwardly
377 to having many unimportant lines, which the program does not remember.
- 378 • SUM - it orders the lines so that the sum of the numbers of the important
379 lines is the smallest possible throughout all levels. The purpose is the same
380 as in the MAX order and quite often it is the case both approaches produce
381 the same order.
- 382 • TWO - it orders the lines so that the maximum number of important lines
383 in two consecutive levels throughout all the levels is as small as possible.
384 This again focuses to having many unimportant lines, which the program
385 does not remember. The constant two is chosen due to the fact general
386 pattern always stores two levels of partial mapping at a time.

387 4.1.5 Mapping approaches

388 The one thing the approaches we will introduce have in common is that they try
 389 to recognize those partial mappings that have no chance to be extended to a full
 390 mapping as early as possible.

391 While the algorithm introduced in [chapter2] finds out the partial mapping
 392 is invalid only at the time it maps two lines having a one-entry at their inter-
 393 section to two lines having a zero-entry at the intersection, different approaches
 394 try to reveal the fact we would end up in the situation earlier by checking more
 395 conditions.

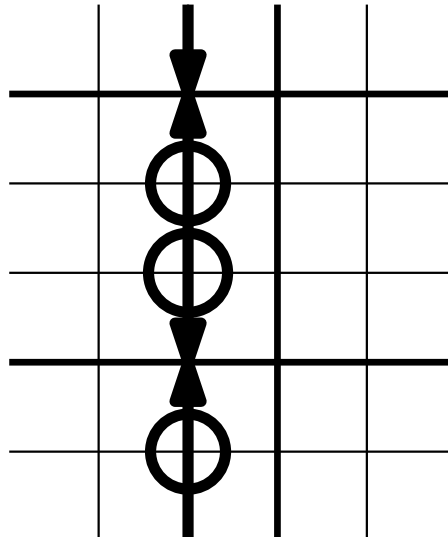
396 Imaging this as the pattern, where $0, \dots, 8$ are indices of the lines:

$$\begin{array}{c} \\ \\ \\ \\ \\ \end{array} \begin{array}{c} 5 \quad 6 \quad 7 \quad 8 \\ \left(\begin{array}{cccc} 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{array} \right) \end{array}$$

397 We are in a situation when only lines 0, 3 and 7 are mapped and line 6 is
 398 currently being mapped. Here are some mapping approaches:

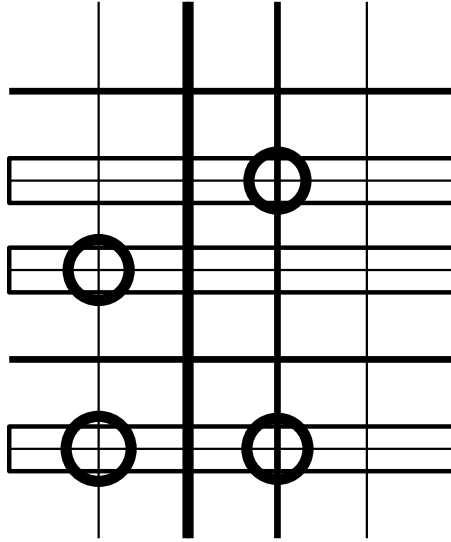
399 Enough one-entries

400 In the situation above, we not only want to check there are one-entries at the
 401 intersections of line 6 with lines 0 and 3, but we also check if there are enough
 402 one-entries in lines between where lines 0 and 3 are mapped so that there is a
 403 hope we can map lines 1 and 2 there and if there is a one-entry below the line
 404 where line 3 is mapped so we can map line 4 there later.



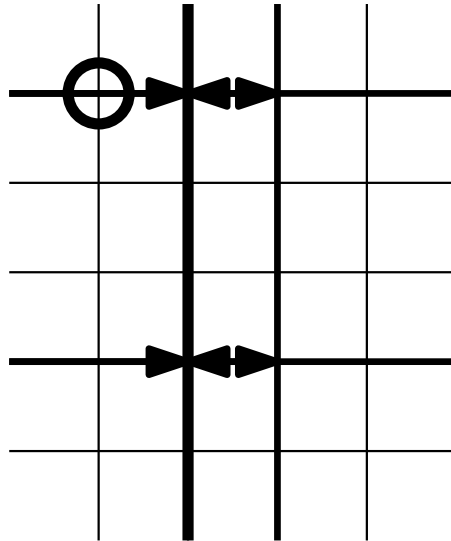
405 Recursive mapping

406 While we were only testing whether there are enough one-entries in between
407 already mapped lines in the previous approach, this time, we also check whether
408 those one-entries can be used for the lines that are intended to be mapped there.
409 For example, when we check there is a one-entry to be used for line 1 later, we
410 also check the line 1 can be mapped to that row, which in this situation means
411 to also check there is a one-entry at the intersection with the line to which the
412 line 7 is mapped.



413 Orthogonal bounds

414 When adding line 6, we check whether there are enough one-entries on the already
415 mapped lines orthogonal to line 6, in between line 6 and the closest mapped lines
416 next to line 6. The same idea as in “Enough one-entries”, but checking different
417 lines.



418 Usage

419 These restrictions on the added lines are not a fixed part of the program. A user
420 can decide which approaches they want to use in the configuration file. This is

421 due to the fact that there is no right path to choose.

422 In the testing that was done for a fixed pattern, we found out it is useful to
423 use all the mentioned restrictions when generating a matrix of size 100×100 , as
424 it turned out to be much faster than without the restrictions. On the other hand,
425 in the same test for a generated matrix of size 500×500 , it was much better not
426 to use any of those restrictions.

427 4.1.6 Using the whole structure in the next iteration

428 It may seem like a good idea to remember all the partial mappings, to propagate
429 them to the next iteration of the MCMC process and alter them depending upon
430 the change.

431 This really can be done. If the change is from zero-entry to one-entry, for each
432 partial mapping we already have we want to try to extend it by the line that just
433 changed and if we manage to do that we then try to extend it to a full mapping
434 in all possible ways if it is a new mapping or do nothing if it is equivalent with a
435 partial mapping of higher level. This can be easily done by means already used
436 in the standard algorithm and may lead to a better performing one.

437 However, if the element gets changed from one-entry to zero-entry we need
438 to go through the partial mappings and delete those that used the currently
439 changed one-entry. This gets a bit messy as we can no longer forget unimportant
440 lines and moreover for each partial mapping we need to remember how many
441 partial mappings of the previous level can be extended to that one, to delete that
442 mapping from the list if there are no longer any mappings extensible to that one.

443 This can all be done, but the whole thing comes with three huge inconven-
444 niences:

- 445 • Memory consumption - there can be a LOT of partial mappings and we
446 need to remember them all. Of course we can still use the equivalence but
447 we need to remember mappings of all levels.
- 448 • The change from one-entry to zero-entry is no longer for free. If this change
449 is done, we already know the pattern is not contained in M , but we still
450 need to do a lot of work to change the structure in order to use it in the
451 next iteration.
- 452 • Reverting - if the change is unsuccessful (the pattern is contained) we need
453 to revert the change which means to completely revert all changes we did to
454 the list of partial mappings. This can be either done by making a backup
455 copy of the whole structure and override the structure if needed, which again
456 is very costly as the structure is huge, or we can remember what partial
457 mappings are new and we go through all partial mappings and remove those
458 new ones. This again means to iterate through the big structure one more
459 time.

460 After realizing these issues it no longer looks useful to me to implement this
461 version of the algorithm.

4.2 MCMC parallelism

To speed up the computations, it is often possible to use parallelism. In this section, we show how to make the MCMC generator parallel, while still allowing both types of the pattern.

While the serial MCMC generator in each iteration changes one element in the generated matrix and checks whether it still avoids forbidden patterns, the parallel one makes several iterations at once, one on each copy of the generated matrix. This means that while iteration x is being computed by a thread, iteration $x + 1$ can at the same time be computed by a different thread. The only issue is that iteration $x + 1$ does not know what is going to be the state of the generated matrix at the time it should start. It expects iteration x to fail - not change the generated matrix at all, counting on the fact at some point it is unlikely a change does not create a mapping of the pattern, and starts with the same matrix as iteration x . If iteration x succeeds though, then the computed iteration $x + 1$ is invalid and the iteration is going to be recomputed again, starting with the altered matrix.

When the parallel version of MCMC generator is chosen and it is assigned n threads, it creates $n - 1$ private copies of the generated matrix and assigns one thread, called worker, to each of them. The last thread, which we call the main thread and which has exclusive access to the master copy of the generated matrix, makes one change of a bit in each private copy of the matrix and makes the corresponding worker check the avoidance.

The job of a worker is only to check if its copy of the matrix still avoids the pattern when one bit is changed. On the other hand all synchronization is left to the main thread. As mentioned before, one iteration of the MCMC process can be recomputed several times. We still want the generator to satisfy the conditions we have for the Markov chain (more in [chapter1]) in order to generate a random matrix. To achieve that, if a computed iteration x succeeds (and changes the generated matrix), all the other computed iterations that would follow after the iteration x become invalid and they all have to be recomputed. The process ends when all iterations get computed.

From now on, we won't be talking about iterations but about tasks. A task is basically one iteration of the MCMC process. The usefulness of this notation comes with an ID - a number, unique per task, assigned to each task, starting with 1 and always increasing. For a pair of consecutive iterations x and $x + 1$ it will always be the case that if task a is the last task to compute iteration x (which means the iteration does not get recomputed ever again after) and task b is the last task to compute iteration $x + 1$, then the ID of a is lower than the ID of b . Also there is no point in which two different tasks would be computing the same iteration at the same time. If tasks with IDs $a < b$ computed the same iteration, it must have been the case an earlier iteration succeeded when task with ID a was computed and after it got removed, task with ID b was assigned to recompute.

At any point in time, we only consider those tasks, that are being computed or those that wait to be processed (not those that have been processed), which means the lowest ID of tasks we consider increases in time.

When a task ends and it has the lowest ID (we can always wait for the task

509 with the lowest ID) we do:

510 • if it fails:

511 – Do nothing - there is no change to propagate to the master copy of the
512 generated matrix and all the tasks with higher ID expected this task
513 to fail, which it did.

514 – This increases the lowest ID by exactly one, as the task we speak of
515 got processed.

516 • if it succeeds:

517 – The main thread propagates the change tested by the task to the
518 master copy of the generated matrix.

519 – All the rest of the task get removed as they all had a higher ID -
520 computed iterations that follow after the one just computed and they
521 expected the task to fail, which it did not.

522 – This increases the lowest ID by more then one, because there are tasks
523 that got removed and one that got processed.

524 4.2.1 Example of the MCMC process for n threads

525 At first, iterations 1 to $n - 1$ are assigned one to each worker as tasks with ID 1 to
526 $n - 1$ with the same order as the order of iterations. If iteration 1 is not successful
527 (which all the other iterations count on), everything is alright. However, if the
528 iteration (its task) is successful, all the results of other tasks (and some of them
529 might have been already finished) are cleared and those iterations get recomputed
530 in tasks n to $2n - 3$ and the worker that computed task with ID 1 is assigned a
531 new task with ID $2n - 2$ - to compute iteration n . The result of the task gets
532 propagated to the master copy of the generated matrix only if all the tasks n to
533 $2n - 3$ fail, else is gets recomputed. This is what happens till the end.

534 4.2.2 Speculative computing

535 It may easily happen that a task not having the lowest ID ends first. In that
536 case we could just wait until it has the lowest ID and process it later. This is
537 not a very efficient approach. Instead we process the task immediately, but we
538 don't propagate the changes to the master copy of the generated matrix until all
539 tasks with lower ID fail and we do not stop the workers processing a task with
540 lower ID. If the task succeeds we remove all the changes computed by tasks with
541 higher ID and override their private copy of the generated matrix. Also it might
542 happen a task with even lower ID succeeds as well. This leads to more and more
543 overriding. Luckily this is the only precarious situation we may encounter and it
544 is not that hard to deal with it, even without copying the possibly huge generated
545 matrix all over the place.

546 The way we deal with these inconveniences is described in [chapter5] and
547 should be clear from the code itself.

548 4.2.3 Reverting and synchronizing in the main thread

549 The speculative computing discussed above is not the only improvement we can
550 make. It turns out to be costly to wake a thread so it computes a trivial function,
551 sets a few atomic variables and falls asleep again. This happens a lot in the
552 MCMC process. Every time a task succeeds it makes other workers revert the
553 changes they computed and synchronize the successful change, which are both
554 trivial functions.

555 To workaround this problem we make a theoretically bad decision which has
556 very nice practical results. All the reverts and synchronizations are computed by
557 the main thread instead of by an appropriate worker. There is no problem with
558 concurrency because the worker is always asleep when a task is to be assigned and
559 using the fact those tasks are really trivial, it does not make the rest of threads
560 wait for the main thread for too long while it computes changes.

561 4.3 Walking pattern

562 While the brute force implementation of an avoid algorithm for a general pattern
563 was improved heavily, the algorithm for a walking pattern is very fast in its nature
564 and cannot be much better. Or can it be?

565 4.3.1 Using the last changed position

566 As the MCMC process works it always changes one element of the big matrix and
567 asks whether it still avoids the pattern. If it does not and we know that before the
568 change it did, we are sure the changed element is a part of the pattern. Knowing
569 that and using the same inductive proof as we did in the proof of correctness of
570 the avoid algorithm (see [chapter2]) it is sufficient to only recompute the part of
571 the inner structure under the changed element and check if the last entry of the
572 pattern can be found there.

573 Not only that. We also know, using the fact the structure was completely
574 correct before the change, that if the values of both c_v and c_h of an element did
575 not change, the element won't cause the element underneath it to change and we
576 no longer have to recompute the other parts of the structure.

577 To use both these facts we replace the cycle through the diagonals by a simple
578 queue, starting at the position of the last changed element and putting more
579 positions in if the values of c_v or c_h are different than they were before. The
580 function ends either when the pattern was discovered or when the queue becomes
581 empty.

582 4.3.2 Lazy avoid

583 Lazy avoid is a variant of avoid function used when the MCMC parallelism is
584 chosen. While all the other types of patterns have a trivial implementation of
585 revert function, when using the walking pattern the inner structure needs to be
586 modified even when reverting. The MCMC parallelism turned out to work much
587 better if the revert calls are handled by the main thread (more in [chapter4]) and
588 it requires the function to run as fast as possible so the other threads are not

589 blocked by the call for too long. That is a reason why functions lazy revert and
590 lazy avoid were created.

591 The avoid function expects the inner structure of the walking pattern (see
592 [chapter3]) to be in a valid state and that requires some effort. To make lazy
593 revert the fastest possible, we postpone the work until the next call of lazy avoid,
594 meaning that lazy avoid then needs to do more things at once. It is no longer
595 sufficient to only compute the submatrix under the position changed last as we
596 did above, but it needs to also compute changes in the positions changed in those
597 lazy revert calls that are postponed.

598 We discuss several approaches, starting with the easiest one and ending with
599 the one that is fast and used in the final implementation.

600 **Recompute the whole structure every time**

601 The easiest way how to implement lazy avoid would be to always recompute the
602 whole inner structure. In that case we do not worry which positions are correct
603 and which are not, because every time we find the pattern, we recomputed all
604 the entries that form it, so we know it really is there. On the other hand, if we
605 manage to recompute the whole structure without finding the last entry of the
606 pattern, it just is not there.

607 The issue is efficiency. If the whole structure was correct and there was a
608 change of the last entry of the matrix it is sufficient to only recompute that one
609 entry. Instead we recompute a possibly very big structure. This results in a very
610 bad performance negating the advantage of parallel computation.

611 **Recompute only a part of the structure diagonal by diagonal**

612 A simple improvement would be to remember the changes done in previous calls
613 of lazy revert and together with the change done in lazy avoid call only recompute
614 the part of the structure that has possibly altered.

615 This gets a bit tricky when lazy avoid call actually discovers the pattern
616 because we cannot be sure the rest of the structure is in a correct order. It is
617 still possible to remember some horizontal, vertical and diagonal bounds and use
618 them to restrict the recomputed part of the matrix. The improvement is not that
619 significant though and we can do better.

620 **Queue of positions to recompute**

621 A different approach is closer to the one used in a standard avoid function. In-
622 stead of going through diagonal one after another, we have a queue of entries-to-
623 recompute. It is no longer sufficient to have a standard queue since in different
624 calls of lazy revert/avoid we can possibly change an entry of different priority (the
625 higher the more important) so we need to have some kind of a priority queue.
626 That is exactly what I tried.

627 Using `std::priority_queue` the function had no more problems with recomput-
628 ing the entries that were not influenced by the changes and used all the benefits
629 mentioned in the previous section. But the container does not come for free and
630 in the end I found out the price I paid for the operations on the priority queue
631 made the whole implementation comparably slow as in the previous attempt.

632 **Two leveled queue of positions to recompute**

633 The final solution comes with the same idea, but a different storage type. As
634 the priority depends upon a diagonal (two entries on the same diagonal can be
635 recomputed in any order) we only remember a priority queue of diagonals and an
636 array of diagonals saying whether a diagonal is already a member of the priority
637 queue. As far as the entries are concerned for every diagonal we have a `std::vector`
638 of entries-to-recompute as well as an array saying whether an entry is already a
639 member of the vector. So finally it is the case that the storage used is not only
640 good theoretically but as the numbers say, also practically. [reference to a table
641 of measurements or something]

642 5. Technical documentation

643 In this chapter, we cover those parts of the algorithm that may be hard to un-
644 derstand just from the code. This only means functions that are technically hard
645 - functions with unexpected dependencies, side effects and so on. Algorithmic
646 difficult tasks are explained in [chapter4].

647 5.1 General pattern

648 The general pattern class contains a lot of function. Most of them are easy to
649 follow and they all should be commented enough in the code. The only part
650 which deserves more attention is the constructor.

651 5.1.1 Construction

652 In the constructor of a general pattern, there are a few function that are easy in
653 nature but as they somehow use each other it is hard not to lose track of their
654 dependencies and results. In order to make this part of the code, which is a very
655 important part indeed, more understandable, we go through the constructor and
656 explain all that is happening in the order it is happening in.

657 Storing the pattern

658 The first thing, which is done right after initialization of variable, is storing the
659 pattern. Instead of storing the pattern in a `Matrix<bool>`, I decided to use to
660 store lines into a number, where in the binary coding a one-entry in the position
661 i means there is a one-entry in the line at the intersection with i -th orthogonal
662 line. This comes handy when computing lines orders. At the same time we also
663 find those lines that are empty (more in [chapter4]) and remember them, because
664 we do not have to map them at all.

665 Choosing the line order

666 After that we need to choose the right line order (again more in [chapter4]).
667 To compute MAX or SUM order we just use a brute force algorithm that checks
668 sequences of line adding and for each it computes how many lines are unimportant.
669 Then it just chooses the order which is the best in chosen metric.

670 What to remember

671 In the next step, we find what do we need to remember in each level of partial
672 mappings with respect to chosen order. As mention earlier, for MAX or SUM
673 order it is already computed when finding the order, but for other variants of
674 orders it is not, so we just compute it every time. What to remember is based
675 on the equivalence introduced in [chapter2] and the decision not to remember
676 unimportant lines (which we explained in [chapter4]).

677 **Parallel bound indices**

678 Now comes the hardest to follow part - precomputing the indices for searching for
679 parallel bounds. The idea behind is simple. When we are adding a new line and
680 we already have a partial mapping, it restricts to where we can add the line. For
681 example, if there are three rows in the pattern and the rows 1 and 3 are mapped,
682 the second one need to be mapped in between those two. The question is, where
683 are those two lines mapped to? First we add in a chosen order and second we do
684 not remember all lines, as some are unimportant. What do we want is to have
685 a instant access to the index of the line, which bounds added line, in the partial
686 mapping so we do not need to compute the index over and over again. That is
687 exactly what gets computed when the function “find_parallel_bound_indices” is
688 called. The series of other function calls follows just because we compute the
689 indices for all added lines in the order in which they are going to be added.

690 **Extending order**

691 The last function, “find_extending_order” just specifies how the next partial map-
692 ping will look like a from where from the previous mapping the values will be
693 copied. Again, unimportant lines play their role here and it may easily be the
694 case from a partial mapping storing k lines, after mapping one more line, we end
695 up with a partial mapping only storing $k - 1$ lines, because two lines become
696 unimportant by adding the line.

697 **5.2 Parallel computing**

698 **5.2.1 MCMC parallelism**

699 While the idea behind MCMC parallelism is described in [chapter4.3] and the
700 code is heavily commented, the work done by the main thread may still be hard
701 to understand.

702 Let I be the ID the process is currently waiting for, that is, the lowest ID of
703 a task that is being tested by a worker. In a structure called “queue” (which is
704 `std::vector<std::deque>`) each worker has a queue of tasks related to it. In the
705 queue, there are tasks that are either being computed or have been computed.
706 The history of tasks is needed to allow reverting changes that should have not
707 happen when the main thread encounters a different successful task with lower
708 ID. There is no need to have a complete history of all tasks computed. There
709 are only those tasks, that have higher ID than I or have lower ID, but those are
710 going to be removed from the “queue” as soon as possible. The name “queue” is
711 not random, it describes the order in which the tasks are being stored - the tasks
712 with lower ID have been inserted earlier and therefore they are at the bottom.

713 Now that we know the most important structure let's see how the main thread
714 works with that and what are the situations.

- 715 • `pop_front`: The main thread deletes the first tasks (the one with the lowest
716 ID) if one of two things happen:
 - 717 – The ID of the task being deleted is equal to I . That means the change
718 computed by the task is being propagated to the generated matrix and

719 there is no need to remember the task anymore. This also increases I ,
720 not necessarily by one.

721 – The ID of the task being deleted is less than I . This situation is due
722 to synchronization. The worker was supposed to synchronize a task
723 computed by a different worker that did not have the lowest ID at the
724 time. Therefore the task needs to be in the list of tasks so we can
725 revert it if needed. If there is no need to revert it and the lowest ID
726 gets greater or equal to the ID of the task, we can just delete it from
727 the “queue”.

728 • `pop_back`: There is only one reason to delete tasks from the end of the
729 “queue” and that is reverting. Imagine there is a task with id J at the
730 end of the “queue”. Now a different worker computes a task with lower
731 ID and finds out the change is successful. This means the task J won’t
732 propagate to the generated matrix and there is no use for it. If it is still
733 being computed, we cannot do much about it, so we just tell the worker to
734 stop computing and deal with it later. If the task is finished, we need to
735 revert it, but only in case the task was successful, because if it was not, it
736 had already been reverted by the worker. So we revert the task if needed
737 and we can just delete it from “queue” as it will never be used.

738 • `emplace_back`: The main thread only inserts new tasks to the end of the
739 “queue” and there are two reasons to insert:

740 – Worker is assigned a completely new task to check the avoidance. In
741 this situation the task is given a new, globally highest ID and we add
742 the task at the end of the list.

743 – The second reason to insert into “queue” are, again, synchronizations.
744 The situation is the same as it was in the case, when we `pop_back` -
745 after we revert all the tasks in the list, we need to synchronize changes
746 that forced reverting and if their ID is not lower or equal to I , we need
747 to add them to the list so they can be reverted if needed.

748 5.3 Library interface

749 6. User documentation

750 In the last chapter of the thesis we show how to make the program generate
751 random matrices or to test whether a certain matrix avoids a given forbidden
752 pattern. In the first section we find out how to set a configuration file. After that
753 we go through necessary input and output files and see how they are formatted.

754 6.1 Configuration file

755 In order to modify what the program computes, we use a configuration file. There
756 is only command line options to be used and that is a path to the configuration
757 file. If no path is inserted the configuration file is expected to be located in the
758 same directory as the executable file is and its name is “config.txt”.

759 The file is a standard text file which can be modified by any text editor and
760 is structured into four sections:

- 761 • input
- 762 • pattern
- 763 • output
- 764 • statistics

765 The order of the sections is not fixed and there can be additional empty lines for
766 better readability. In each section there is a list of values that can be set either
767 to arbitrary value or to a specific one. There is at most one command of format
768 “option=value” per line and there might be additional white spaces surrounding
769 the “=” sign.

770 If an option is set more than once, the latter value is always used. If, on the
771 other hand, an option is not set at all, the default value is used. If there is a
772 line encountered that sets a wrong option, for instance when the user mistypes a
773 valid option, the line is skipped and the user gets a warning in the standard error
774 output.

775 Let us provide a list of all options for each section together with their default
776 values.

777 6.1.1 Input

778 In the first section of the configuration file we set the generating process.

- 779 • size: The size of the generated matrix. Results in $M \in \{0, 1\}^{size \times size}$.
Possible value: $s \in \mathbb{N}$
780 Default value: 100
- 781 • iterations: The number of iterations of the MCMC process.
Possible value: $i \in \mathbb{N}$
782 Default value: 10,000

- 783 • `random_seed`: The random seed for the MCMC process.
 Possible value: $s \in \mathbb{N}$
 “random” - chooses a random seed
 784 Default value: “random”
- 785 • `init_matrix`: A $size \times size$ matrix the MCMC process starts with.
 Possible value: *input file path*
 “zero” - a matrix containing no one-entries
 786 Default value: “zero”
- 787 • `parallel_mode`: choose whether to compute in parallel or serial.
 Possible value: “serial”
 “mcmc” - more iterations of the MCMC generator are tested in parallel
 788 “map” - more partial mappings are being extended in parallel
 Default value: “serial”
- 789 • `threads_count`: The number of threads that are going to be used if a parallel
 790 mode is chosen.
 Possible value: $t \in \mathbb{N}$
 -1 - chosen according to the number of cores
 791 Default value: 1

792 6.1.2 Pattern

793 It is time to set the options that matter the most - matrix patterns. As you
 794 generate a matrix which avoids more than just one pattern, the section [pattern]
 795 can be used multiple times, specifying one pattern for each occurrence.

- 796 • `pattern_file`: A path to a input matrix file - the pattern.
 Possible value: matrix file path
 797 Default value: “input.txt”
- 798 • `pattern_type`: The type of the pattern. Determines the method used for
 799 testing avoidance.
 Possible value: “general”
 “walking” - see [chapter3]
 800 “slow” - brute force algorithm for a general pattern
 Default value: “general”

801 The next options are only useful if the general pattern type is chosen. It
 802 specifies how the mappings are stored as well as what the map function tests.

- 803 • `map_one_entries`: If set to “yes”, the map function not only checks that
 804 already mapped lines do cross the currently added one in a one-entry if
 805 needed but also tests whether there is enough one-entries in between those
 806 already mapped lines. The result is that the map function takes a little
 807 longer but it recognizes hopeless mappings earlier.
 Possible value: “yes”
 “no”
 808 Default value: “yes”

- 809 • `map_recursion`: If set to “yes” and the `map_one_entries` is also set to “yes”,
810 the `map` function not only tests there are enough one-entries in between
811 already mapped lines on the currently added one, but for those non-mapped
812 lines crossing the current one in a one-entry, it also tests if the partial
813 mapping has one-entries everywhere they are supposed to be. Although
814 this comes with a big time consumption it can filter a big portion of partial
815 mapping that would later become unfeasible.
816 Possible value: “yes”
817 “no”
818 Default value: “yes”
- 817 • `map_orthogonal_bounds`: If set to “yes”, the `map` function also tests the
818 orthogonal bounds of added line - see [chapter2].
819 Possible value: “yes”
820 “no”
821 Default value: “no”
- 820 • `map_container`: A container in which the partial mappings are stored.
821 Possible value: “set” - `std::set` (red-black tree)
822 “hash” - `std::unordered_set` (hash table)
823 “vector” - `std::vector` (dynamic array)
824 Default value: “hash”
- 822 • `line_order`: Choose the order in which the lines are being added to the
823 partial mapping. See [chapter2]
824 Possible value: “max”
825 “two”
826 “sum”
827 “desc”
828 “auto”
829 “order file path”
830 Default value: “max”

825 6.1.3 Output

826 There is no reason to generate a matrix noone will ever see. In this section we
827 make the matrix show. As the matrix can be output to console, a text file or a
828 bmp file an option in the section can be set more than once and all the settings
829 make difference.

- 830 • `matrix_output`: The generated matrix can be output as a bmp file in which
831 one-entries are black pixels and zero-entries white. To do that, the file path
832 has to have a pattern “path.bmp”. If a different path is given the file is
833 stored as a matrix text file. It can also be output into a console if “console”
834 is set. In that case it has the text format.
835 Possible value: “console”
836 matrix bmp file path
837 matrix text file path
838 “no”
839 Default value: “no”

- 836 • `performance_stats`: If the serial computation or the map parallelism is cho-
837 sen, the program can output a statistics like the percentage of avoid call
838 success, how long did one call take on average and what was the average
839 size of structures. Note that this is not very useful as a matrix statistics,
840 see the next section for that. If more patterns are set, the statistics may get
841 misleading as it also counts the cases when the first pattern is contained in
842 the matrix and the other patterns are not tested at all.
Possible value: "console"
performance file path
843 "no"
Default value: "no"
- 844 • `performance_csv_stats`: The same information as above but formatted to a
845 csv file so the data can be more easily worked with.
Possible value: "console"
csv file path
846 "no"
Default value: "no"
- 847 • `time_to_console`: Prints how long the computation took into a console.
Possible value: "yes"
848 "no"
Default value: "no"
- 849 • `patterns_to_console`: Prints all the used patterns into the console.
Possible value: "yes"
850 "no"
Default value: "no"

851 6.1.4 Statistics

852 The last section handles the options important for scientists. While generating
853 a random matrix is a great result, on its way the program can also create some
854 statistics, namely make a histogram of occurrences of one-entries in a generated
855 matrix as the MCMC iterates as well as store the matrix with the highest amount
856 of one-entries. As the process usually does not start with a random matrix, the
857 user can decide to only compute the statistics after a certain number of iterations
858 has been done and to only check a small portion of iterations, every 10th for
859 instance, as a single iteration may not make any difference and counting the
860 histogram takes time.

- 861 • `histogram_frequency`: Sets how often the histogram gets refreshed.
Possible value: $f \in \mathbb{N}$
862 0 - the histogram is not computed at all
Default value: 0
- 863 • `histogram_initial`: Sets the initial iteration of the MCMC process when the
864 histogram gets refreshed.
Possible value: $i \in \mathbb{N}$
865 Default value: 1,000

866 • histogram_final: Sets the last iteration of the MCMC process when the
867 histogram gets refreshed.
Possible value: $f \in \mathbb{N}$
868 -1 - the histogram is computed till the end
Default value: -1

869 • histogram_file: Sets where to output the histogram computed during the
870 MCMC process.
Possible value: matrix bmp file path
matrix text file path
871 “console”
“no”
Default value: “no”

872 • max_ones_matrix_file: Sets where to output the matrix that had the most
873 one-entries among all matrices iterated through during the MCMC process.
Possible value: matrix bmp file path
matrix text file path
874 “console”
“no”
Default value: “no”

875 6.2 File input

876 There are only two types of input files expected by the program. Either you want
877 to read a matrix file, which can be a pattern or an initial matrix, or an order file
878 that determines an order in which the lines are going to be mapped if the general
879 pattern is chosen.

880 6.2.1 Matrix file

881 A matrix file is a standard text file having the format as follows:

- 882 • 2 natural numbers specifying the number of rows and columns in this order.
- 883 • a sequence of zeros and ones of length rows×columns specifying the matrix
884 from the top left corner one row after another.

885 **Example:** 2 3
 1 0 1
 1 1 0

886 6.2.2 Order file

887 If you want to choose the order in which the lines are going to be mapped when
888 a general pattern is chosen, it is your responsibility to check that all lines that
889 need to be mapped are mapped. It is for example possible to only map three lines
890 even if the pattern consists of six lines just because there is for example no need

891 to map empty lines at all. Therefore the program does not check the validity of
892 the order and just uses it.

893 Now that the user has been warned, the format of the custom order file is
894 simple. It consist of the indices of the lines of the pattern numbered starting
895 with 0 and starting from the top row and ending with the right column.

896 One possible order for the matrix given as an example in [6.2.1] is this file:

897 2 1 0 3 4

898 First mapping the left column, the second and first row after that and finishing
899 the mapping with the middle column and the right one.

900 6.3 File output

901 Let us now find out what the output files look like.

902 6.3.1 Matrix text file

903 The matrix text file has the same format as the input one. It consists of:

- 904 • 2 natural numbers specifying the number of rows and columns in this order.
- 905 • a sequence of zeros and ones of length rows×columns specifying the matrix
906 from the top left corner one row after another.

907 The matrix is binary except for the one produced as a histogram, which can have
908 higher natural numbers and contains the number of samples as the last number.
909 If you then divide all the entries by the last number, you get a percentage of the
910 entry being a one-entry.

911 6.3.2 Matrix bmp file

912 For an $n \times n$ matrix the standard bmp file contains $n \times n$ pixel of black color
913 meaning a one-entry and a white color for a zero-entry. If the histogram is output
914 as a bmp file, the pixels are greyscaled and the darker a pixel is the more often
915 the entry was a one-entry during the MCMC process.

⁹¹⁶ Conclusion

917 Bibliography

918 Neal Madras. *Lectures on monte carlo methods*, volume 16. Springer Science &
919 Business, 2002.

920 Neal Madras and Hailong Liu. Random pattern-avoiding permutations. *Algo-*
921 *rithmic Probability and Combinatorics, AMS, Providence, RI*, pages 173–194,
922 2010.

923 List of Figures

924	1	Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$	
925		satisfies all the conditions. On the other hand, matrix M_2 avoids	
926		P as there is no such mapping.	4
927	2.1	An example showing unimportant line and equivalent mappings. .	9
928	3.1	An example of a walk and the order of its entries.	10

929 List of Tables

⁹³⁰ List of Abbreviations

931 Attachments