

Charles University in Prague
Faculty of Mathematics and Physics

BACHELOR THESIS



Stanislav Kučera

Generating random pattern-avoiding matrices

Computer Science Institute of Charles University

Supervisor of the bachelor thesis: RNDr. Vít Jelínek, Ph.D.

Study programme: Computer Science

Study branch: General Computer Science

Prague 2016

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University in Prague has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

signature of the author

Title: Generating random pattern-avoiding matrices

Author: Stanislav Kučera

Department: Computer Science Institute of Charles University

Supervisor: RNDr. Vít Jelínek, Ph.D., department

Abstract: Abstract.

Keywords: binary matrix pattern-avoiding Markov chain Monte Carlo

Dedication.

Contents

Preface	3
Introduction	5
1 Markov chain Monte Carlo	6
1.1 Markov chains	6
1.2 Markov chain for pattern-avoiding binary matrices	6
2 An algorithm for testing pattern-avoidance of a general pattern	8
2.1 Sketch of a brute force algorithm	8
2.2 Equivalent mappings	8
3 An algorithm for testing pattern-avoidance of a special pattern	10
3.1 Walking pattern	10
3.2 Dynamic program	10
3.2.1 The algorithm	11
3.2.2 Inner structure	12
3.2.3 Correctness	13
3.2.4 Generalization	13
4 Improvements to basic algorithms	14
4.1 General pattern	14
4.1.1 Improving memory consumption	14
4.1.2 Not mapping empty lines	14
4.1.3 Using the last changed position	14
4.1.4 Line order	15
4.1.5 Mapping approaches	16
4.1.6 Using the whole structure in the next iteration	18
4.2 MCMC parallelism	19
4.2.1 Example of the MCMC process for n threads	20
4.2.2 Speculative computing	20
4.2.3 Reverting and synchronizing in the main thread	21
4.3 Walking pattern	21
4.3.1 Using the last changed position	21
4.3.2 Lazy avoid	21
5 Technical documentation	24
5.1 General pattern	24
5.1.1 Construction	24
5.2 Parallel computing	25
5.2.1 MCMC parallelism	25
5.3 Library interface	26

6	User documentation	27
6.1	Configuration file	27
6.1.1	Input	27
6.1.2	Pattern	28
6.1.3	Output	29
6.1.4	Statistics	30
6.2	File input	31
6.2.1	Matrix file	31
6.2.2	Order file	31
6.3	File output	32
6.3.1	Matrix text file	32
6.3.2	Matrix bmp file	32
	Conclusion	33
	Bibliography	34
	List of Figures	35
	List of Tables	36
	List of Abbreviations	37
	Attachments	38

1 Preface

2 Theses at the faculty of mathematics and physics usually fit into one of three
3 categories:

- 4 1. Theoretical thesis
- 5 2. Experimental thesis
- 6 3. Implementation thesis

7 My thesis does not fit entirely into only one category and it does not try to. The
8 project consists of several similarly important parts which are:

- 9 • Design of the algorithm for generating a special binary matrix
- 10 • Making it run fast on inputs which are usual for researchers
- 11 • Implementing the algorithm to provide practical tool

12 None of these points would make sense alone but together the thesis may be-
13 come very useful for scientists as it provides with a process generating random
14 matrices and it is a common practice to test hypothesis on random data.

Introduction

We let $M \in \{0, 1\}^{n \times m}$ denote a *binary matrix* of size n by m . The *height* of M , denoted by n , is the number of rows of M and m is its width (the number of columns). A *line* of a matrix is one of its rows or columns and we denote by $L(M)$ the ordered set of all lines of M . Its order is given by the standard indexing of rows and columns.

Definition 1. We say a binary matrix M contains a binary matrix P , which we call a “pattern”, as a submatrix, if there is a mapping $f : L(P) \rightarrow L(M)$, such that

- $l \in L(P)$ is a row of P iff $f(l) \in L(M)$ is a row of M
 - $\forall l, l' \in L(P) : l < l' \Rightarrow f(l) < f(l')$ (preserves the order)
 - $\forall l, l' \in L(P) : \text{if lines } l \text{ and } l' \text{ intersect and there is a one-entry at the intersection, then there is a one-entry at the intersection of } f(l) \text{ and } f(l').$
- otherwise, it avoids the pattern P .

$$P = \begin{matrix} & \begin{matrix} 2 & 3 \end{matrix} \\ \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \end{matrix} \quad M_1 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix} \quad M_2 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix}$$

Figure 1: Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$ satisfies all the conditions. On the other hand, matrix M_2 avoids P as there is no such mapping.

The interesting cases are square matrices of size n by n , where n is big (going to infinity) and the size of a pattern (not necessarily square matrix) is small (constant). Even for a constant size forbidden pattern it is hard to determine the number of matrices of size n that avoid it or to characterize, what properties they have. Sometimes we consider matrices avoiding more than just one forbidden pattern, in which case we denote the set of all forbidden matrices by \mathcal{P} . When a matrix avoids \mathcal{P} , it avoids every $P \in \mathcal{P}$.

Definition 2. We denote by $\mathcal{M}_n(\mathcal{P})$ a set of all binary matrices of size n by n avoiding \mathcal{P} as submatrices. We always call M the square binary matrix for which we test the containing and P the pattern (if there is one) that is being tested. Moreover, we denote by h the height (the number of rows) of P and w its width.

The area of pattern avoidance has been heavily studied for permutations and it also becomes more popular for their generalization - binary matrices. In most of the areas in combinatorics it is useful to explore properties of random objects and a lot of attention is directed towards random matrices when considering pattern avoidance. The goal of the work is, for given $n \in \mathbb{N}$ and set of forbidden patterns \mathcal{P} , to generate a uniformly random $M \in_R \mathcal{M}_n(\mathcal{P})$.

46 **Generating random matrix**

47 One way to get $M \in_R \mathcal{M}_n(\mathcal{P})$ is to choose a matrix of required size completely
48 at random, for such, test whether it avoids the pattern and simply repeat the
49 process until we find one, which does. However, in the most interesting cases,
50 only a small fraction of all matrices avoid the pattern and the process takes too
51 long, to be practically useful.

52 For generating random permutations avoiding forbidden pattern, a different
53 technique was introduced in Madras and Liu [2010]. It uses a randomized pro-
54 cess called Markov chain Monte Carlo, which we will abbreviate by MCMC. It
55 is an iterative process, which for a well chosen Markov chain (more in chapter 1)
56 approximates a random object. The algorithm by Madras and Liu was devel-
57 oped for permutations (permutation matrices) and it cannot be used for general
58 matrices. In section 1.2 we show how to adapt the algorithm, which will lead
59 us to a MCMC algorithm that approximates $M \in_R \mathcal{M}_n(\mathcal{P})$. To produce a good
60 approximation the process needs to do a lot of iterations and despite the fact
61 it is unknown what is the mixing time (the number of iterations required) of a
62 MCMC process, in practice, the method does better than the trivial algorithm.

63 **Testing avoidance**

64 In each step of our MCMC process we need to test whether a matrix avoids a
65 pattern. We will show a very fast algorithm that only works for a special class of
66 binary matrices (explained in chapter 3) together with a slightly less performing
67 algorithm for a general pattern, which, again, comes as a generalization of an
68 algorithm for permutations from the article by Madras and Liu and is described
69 in chapter 2.

70 In chapter 4 we improve both our algorithms and introduce a parallel version
71 of MCMC process, which further increases the performance of matrix generating.

72 In chapter 5 some technical details are explained to make reading the code
73 easier for reader and to describe user interface. The last chapter (chapter 6)
74 contains user documentation.

1. Markov chain Monte Carlo

Our goal to generate $M \in_R \mathcal{M}(\mathcal{P})$ heavily depends on the theory of Markov chains. In this work we only define useful terms and state two important theorems. If you are interested in more details, see Madras [2002].

1.1 Markov chains

Definition 3. We shall consider discrete-time Markov chain X_0, X_1, \dots , where $X_i \in \mathcal{S}$, for a finite state space \mathcal{S} and every i (number of steps). The k -step transition probabilities are:

$$p_{i,j}^{(k)} = \Pr(X_{t+k} = j | X_t = i) \quad (i, j \in \mathcal{S})$$

Definition 4. A Markov chain is said to be symmetric if $p_{i,j}^{(1)} = p_{j,i}^{(1)}$ for every pair of states i and j .

Definition 5. A Markov chain is irreducible if the chain can eventually get from each state to every other state, that is, for every $i, j \in \mathcal{S}$ there exists a $k \geq 0$ (depending on i and j) such that $p_{i,j}^{(k)} > 0$.

Definition 6. An irreducible chain has period D if D is the greatest common divisor of $\{k \geq 1 | p_{i,i}^{(k)} > 0\}$ for some $i \in \mathcal{S}$ (equivalently, for all $i \in \mathcal{S}$). A chain is called aperiodic if its period is 1. In particular, if an irreducible chain has $p_{i,i}^{(1)} > 0$ for some i , then it is aperiodic.

Theorem 1. Consider an aperiodic irreducible Markov chain with state space \mathcal{S} . For every $i, j \in \mathcal{S}$, the limit $\lim_{k \rightarrow \infty} p_{i,j}^{(k)}$ exists and is independent of i ; call it π_j . Furthermore, if \mathcal{S} is finite, then

$$\sum_{j \in \mathcal{S}} \pi_j = 1 \quad \wedge \quad \sum_{i \in \mathcal{S}} \pi_i p_{i,j}^{(1)} = \pi_j$$

for every $j \in \mathcal{S}$. That is, if we write π to denote the row vector whose entries are π_i , then $\pi P = \pi$.

Theorem 2. Suppose that an irreducible Markov chain on the finite state space \mathcal{S} is symmetric. Then the equilibrium distribution is uniform on \mathcal{S} .

In other words, the theorems together give us a guarantee that if we choose an irreducible, symmetric and aperiodic Markov chain with state space \mathcal{S} then the probability distribution of X_i converges to uniform distribution on \mathcal{S} independently of the initial state.

1.2 Markov chain for pattern-avoiding binary matrices

To generate a binary matrix $M \in \{0, 1\}^{n \times n}$ avoiding patterns in \mathcal{P} , we create a Markov chain, whose states space is $\mathcal{M}_n(\mathcal{P})$. After sufficiently many iterations (m) of MCMC process we set $M := X_m \in \mathcal{M}_n(\mathcal{P})$. We always begin with an initial matrix X_0 and the process looks like this:

- 103 1. For $i := 1, 2, \dots, m$:
- 104 2. Set $X_i := X_{i-1}$.
- 105 3. Choose $r \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 106 4. Choose $c \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 107 5. Flip the bit at $X_i[r, c]$.
- 108 6. If X_i contains \mathcal{P} , flip the bit back.

109 If the process starts with a matrix X_0 that avoids \mathcal{P} , then after every step it
 110 still avoids \mathcal{P} . Note that an iteration does not change the matrix if the condition
 111 6 is satisfied. We need to show the Markov chain we presented meets all the
 112 conditions of both theorems:

113 Symmetry

114 Imagine a sequence of bits flipping which changes the i -th matrix to j -th one.
 115 The reversed order of the same sequence changes the j -th matrix to the i -th one.

116 Irreducibility

117 As the steps go, it is easy to see we can with non-zero probability create any
 118 matrix $M \in \mathcal{M}_n(\mathcal{P})$ from the zero matrix $0_n = 0^{n \times n}$ by choosing the one-entries
 119 of M . When we can get from 0_n to M by a sequence of flip changes, the reversed
 120 sequence is a sequence of steps from any matrix $\in \mathcal{M}_n(\mathcal{P})$ to 0_n . Thus the Markov
 121 chain is irreducible.

122 Aperiodicity

123 The Markov chain is irreducible so it suffices to show that there is an i for which
 124 $p_{i,i}^{(1)} > 0$. Clearly there is a matrix for which there is at least one bit that cannot
 125 be flipped without creating a pattern and this forces $p_{i,i}^{(1)} > 0$.

126 2. An algorithm for testing 127 pattern-avoidance of a general 128 pattern

129 In this chapter and chapter 3 we show algorithms for testing whether a pattern
130 P is contained in a square binary matrix M .

131 We begin with a very basic algorithm, which we then improve a lot to get a
132 fast algorithm for testing avoidance of a general pattern.

133 2.1 Sketch of a brute force algorithm

134 Let $L = (l_1, l_2, \dots, l_{w+h})$ be a permutation of lines (rows and columns) of the
135 pattern P and $k \in [w+h]$. *Partial mapping of level k* of lines of P is a function
136 f from $L' := \{l_1, l_2, \dots, l_k\} \subseteq L$ to lines of the big matrix M satisfying two
137 conditions:

- 138 • Both $l' \in L'$ and $f(l')$ are rows or they are both columns.
- 139 • If $l' \in L'$ and $l'' \in L'$ are both rows or columns and $l' < l''$, then $f(l') < f(l'')$.
140 This means partial mapping keeps the order of the lines.
- 141 • If $l' \in L'$ is a row of P and $l'' \in L'$ is a column of P and there is a one-entry
142 at the intersection of l' and l'' , then there is a one-entry at the intersection
143 of $f(l')$ and $f(l'')$.

144 The basic algorithm we use goes as follows. First it maps l_1 to all possible lines
145 of M , creating partial mappings of $\{l_1\} \subseteq L$. For $k = 2, \dots, w+h$ it takes each
146 partial mapping from the previous iteration and extends it by adding line l_k to
147 the partial mapping in all possible ways. If we manage to map all the lines of P ,
148 then M does not avoid it and if at some point there are no partial mappings to
149 extend it means M avoids P .

150 The algorithm can be improved in two ways. Firstly, we can try to recognize
151 unextendable partial mappings earlier than at the moment a line can no longer be
152 mapped, for example by counting whether there is enough one-entries in between
153 already mapped lines (more in subsection 4.1.5). Secondly, which is going to be
154 fundamental for us, we can try not to remember more copies of different mappings
155 that can be extended in the same way.

156 2.2 Equivalent mappings

157 There is no need to remember two different mappings if they can be both extended
158 exactly the same way as our function is only supposed to check whether a pattern
159 can be mapped to a big matrix not to find all such mappings.

160 **Definition 7.** We call a line l of a pattern important in a partial mapping if one
161 of the conditions is met:

- 162 • An adjacent line of the pattern has not been mapped yet.
- 163 • There is a one-entry on the line l at the intersection with line l' that has
- 164 not been mapped yet.
- 165 . Otherwise the line is unimportant in the mapping.

166 Whether a line is important or not only depends on the permutation, we have
 167 chosen for partial functions, because that is what forces the order in which we
 168 map lines.

169 At the beginning, when no line is mapped, all lines are important. After some
 170 lines get mapped, a line can become unimportant in the partial mapping as all
 171 lines that bound it are in the mapping as well. If a line is unimportant in a partial
 172 mapping of some level, it will stay unimportant in all extensions of the mapping
 173 we can find.

174 **Definition 8.** We say two partial mappings of the same level are equivalent if
 175 all important lines in the mapping of that level are mapped to the same lines of
 176 the big matrix in both mappings.

$$\begin{array}{c}
 \begin{array}{ccccc}
 & 4 & 5 & 6 & 7 \\
 0 & \begin{pmatrix} 1 & 1 & 0 & 0 \end{pmatrix} \\
 1 & \begin{pmatrix} 0 & 1 & 1 & 1 \end{pmatrix} \\
 2 & \begin{pmatrix} 0 & 1 & 0 & 0 \end{pmatrix} \\
 3 & \begin{pmatrix} 0 & 1 & 0 & 1 \end{pmatrix}
 \end{array}
 &
 \begin{array}{ccccc}
 & 5 & 6 & 7 & 8 & 9 \\
 0 & \begin{pmatrix} 1 & 1 & 1 & 1 & 1 \end{pmatrix} \\
 1 & \begin{pmatrix} 0 & 1 & 1 & 0 & 1 \end{pmatrix} \\
 2 & \begin{pmatrix} 1 & 1 & 0 & 0 & 0 \end{pmatrix} \\
 3 & \begin{pmatrix} 0 & 1 & 1 & 0 & 1 \end{pmatrix} \\
 4 & \begin{pmatrix} 1 & 1 & 1 & 0 & 1 \end{pmatrix}
 \end{array}
 \end{array}
 \quad
 P=
 \quad
 M=$$

Figure 2.1: An example showing unimportant line and equivalent mappings.

177 For P and M , matrices in Figure 2.1, in partial mapping $f = \{(1, 1), (2, 2), (3, 4), (5, 6)\}$
 178 line 2 is unimportant because both lines 1 and 3 are mapped and so is line 5 -
 179 the only line to intersect line 2 in a one-entry. Line 3 is important, because there
 180 is line 7 intersecting it in one-entry, which is not mapped.

181 In the same situation as above, consider a different partial mapping $f' =$
 182 $\{(1, 1), (2, 3), (3, 4), (5, 6)\}$, which is a mapping of the same level as f and only
 183 differs from f in mapping line 2. The line 2 is unimportant and by the definition
 184 of equivalent partial mappings, f and f' are equivalent. The idea behind this
 185 notion is simple. It is not important where we map line 2, because it does not
 186 restrict where we can map any other line that has not been mapped yet. This
 187 means that if a partial mapping f can be somehow extended, the equivalent
 188 partial mapping f' can be extended in the same way; therefore, it is sufficient to
 189 only extend one of them in order to find one full mapping. Note that it would
 190 be also sufficient to only extend one of the partial mappings if we were looking
 191 for all full mappings, but, in that case, we would need to keep the information
 192 about where the unimportant lines were mapped to.

3. An algorithm for testing pattern-avoidance of a special pattern

In the previous chapter, we have seen an algorithm for a general forbidden pattern which, using some heuristics, runs pretty fast. In this chapter, we introduce a special kind of a pattern, satisfying additional conditions, for which we can produce much faster algorithm.

3.1 Walking pattern

We call the specific pattern a walking pattern. The additional condition we want the pattern to satisfy is that there is a walk from one corner to the opposite one and all the one-entries of the pattern are contained on the walk.

Definition A **walk** in a matrix is a sequence of some of its entries beginning in the top left corner and ending in the bottom right one. If an entry at the position $[i, j]$ is in the sequence, the next one is either $[i + 1, j]$ or $[i, j + 1]$. Therefore, the length of an arbitrary walk is equal to $w + h - 1$.

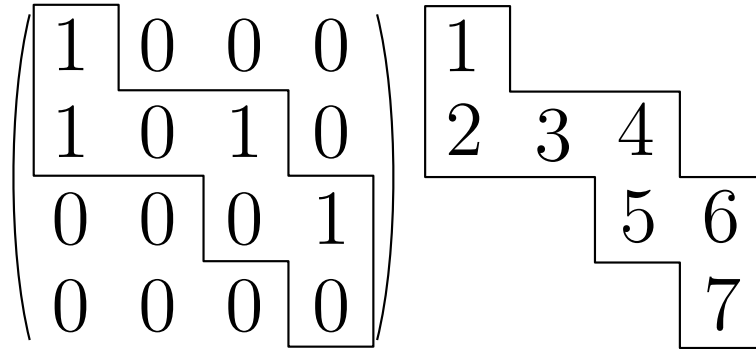


Figure 3.1: An example of a walk and the order of its entries.

In Figure 3.1 you can see a matrix that is a walking pattern as all the one-entries are included in a walk. Not all entries of a walk need to be one-entries though.

It can be shown a walking pattern with a walk is exactly a matrix avoiding a forbidden pattern

$$\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

3.2 Dynamic program

Now that we know what walking pattern is, we show an algorithm deciding whether such a pattern P is contained in a big matrix M or not.

214 The pattern P is a walking pattern so there is a walk containing all the one-
 215 entries of the pattern. We choose one such walk arbitrarily and index its entries
 216 $e_1, e_2, \dots, e_{h+w-1}$ starting from the beginning of the walk. For each entry of the
 217 walk we remember whether its value is one or zero and whether the walk continues
 218 from the entry vertically, in which case we call it a **vertical entry** or horizontally,
 219 calling it a **horizontal entry**.

220 For an element e of M at the position $[i, j]$, the matrix $M_{\leq e}$ is a $(i+1) \times (j+1)$
 221 submatrix of M consisting of rows with the index smaller than or equal to i and
 222 columns with the index smaller than or equal to j . The element e then lies in
 223 the bottom right corner. Similarly, $M_{\geq e}$ is a $(n-i) \times (n-j)$ submatrix of M
 224 consisting of rows with the index greater than or equal to i and columns with
 225 index greater than or equal to j . The element e is its first element.

226 To determine whether P is contained in M we find out for each element e of
 227 M what is the longest part of the pattern that can be found in $M_{\leq e}$. If there is
 228 an element for which we manage to find the last entry of the pattern, the pattern
 229 is contained in the matrix; otherwise, it is avoided.

230 For each element e of M at the position $[i, j]$ we remember two numbers. The
 231 number $c_v(e)$ says what is the longest part of the walk in $M_{\leq e}$ with the last entry
 232 in j -th column and being a vertical entry. The number $c_h(e)$, symmetrically, says
 233 what is the longest part of the walk in $M_{\leq e}$ with the last entry in i -th row and
 234 and being a horizontal entry.

235 An observation we make is that if we have a fixed element e of M and any
 236 other element e' above e in the same column then $c_v(e')$ is less than or equal to
 237 $c_v(e)$. This means that for e we can find the maximum part of the pattern ending
 238 in the column of e and continuing vertically by looking only to elements in that
 239 column above e and since this is true for all of them, it is sufficient to only check
 240 the value of the element right above e (at the position $[i-1, j]$). Similarly the
 241 argument goes for the value of c_h in horizontal way.

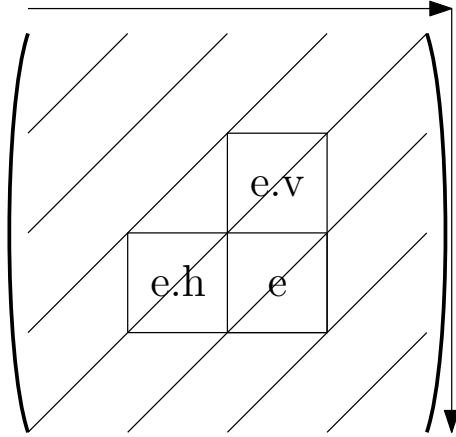
242 The algorithm iterates through diagonals. A diagonal in this matter of speak-
 243 ing is a subset of elements of M , such that all elements have the same sum of
 244 their coordinates. For example, zero diagonal only consists of an element $[0, 0]$,
 245 the first diagonal contains elements $[0, 1]$ and $[1, 0]$, and so on.

246 For simplicity, in the pseudo-code below we do not deal with elements outside
 247 M (like $-1, 0$) explicitly. Instead for those elements we just assume the values of
 248 c_v and c_h are always equal to zero.

249 3.2.1 The algorithm

- 250 1. For $d = 0, \dots, w + h - 1$
- 251 2. For e element of d -th diagonal at the position $[i, j]$
- 252 3. $e_v := [i - 1, j]$
- 253 4. $e_h := [i, j - 1]$
- 254 5. $c_v(e) := c_v(e_v)$
- 255 6. $c_h(e) := c_h(e_h)$

256 7. If $w_{c_v(e)+1}$ can be mapped to e
 257 8. If $c_v(e) + 1 = w + h + 1$
 258 9. Terminate - M contains P as a submatrix
 259 10. If $w_{c_v(e)+1}$ is a vertical entry
 260 11. $c_v(e) := c_v(e) + 1$
 261 12. Else
 262 13. $c_h(e) := \max\{c_h(e), c_v(e) + 1\}$
 263 14. If $w_{c_h(e)+1}$ can be mapped to e
 264 15. If $c_h(e) + 1 = w + h + 1$
 265 16. Terminate - M contains P as a submatrix
 266 17. If $w_{c_h(e)+1}$ is a vertical entry
 267 18. $c_v(e) := \max\{c_v(e), c_h(e) + 1\}$
 268 19. Else
 269 20. $c_h(e) := \max\{c_h(e), c_h(e) + 1\}$



270 3.2.2 Inner structure

271 To run properly the algorithm needs two structures. The first one is a description
 272 of the walk, which is just an array of the values of its entries as well as the
 273 information whether the entry is vertical or horizontal. The second structure is
 274 a matrix of the values c_v and c_h as described above.

275 3.2.3 Correctness

276 We need to show that the values of c_v and c_h are always correct for the recomputed
277 elements after at the end of the function. We proceed by induction.

278 For the first element it is definitely true since there can be only the first entry
279 of the pattern mapped and we check just that.

280 When we compute an element e of a computed diagonal d , by induction as-
281 sumption all the diagonals $d' < d$ are correctly computed. In particular, the
282 values are correct in the diagonal $d - 1$. To compute the correct values of e , we
283 use the values of two element on the diagonal $d - 1$: e_v , which is right above e and
284 e_h , which is the first element to the left of e . If e_v or e_h are outside the matrix
285 then from that direction we cannot expect to find anything more than just the
286 first entry of the pattern and that is what we check for.

287 Let v be the true length of the longest part of P in $M_{\leq e}$ continuing vertically
288 in the same column as e . Now if e itself is not an entry of that part of the pattern,
289 it is a different element e' in the same column. But then the value of $c_v(e')$ is
290 correctly computed by the inductive hypothesis and it is copied to all element
291 underneath. Especially e_v gets the value and the algorithm copies the value from
292 it to e . On the other hand if e is an entry of the part of the pattern we work with,
293 it is the last entry. The entry right before the last one needs to be mapped to
294 the same row or column; therefore, either e_v or e_h contain the part of the pattern
295 shorter by one and the algorithm extends it to a correct value.

296 3.2.4 Generalization

297 The same algorithm, just rotated by 90 degrees, can be also used for a pattern
298 where all one-entries are contained on a walk from top right corner to the bottom
299 left one. Indeed the program uses it and if given a walking pattern it determines
300 by itself which walk it is.

301 On the other hand a direct generalization for a general pattern does not work.
302 While we can index all entries of the pattern, when trying to map a certain w_k
303 to an element it is not sufficient to just check whether w_l is above and w'_l to the
304 left from the element.

$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & \boxed{1} \\ 1 & \textcircled{1} & \diamond 1 \end{pmatrix} \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & \boxed{1} \\ 0 & 1 & \textcircled{1} & \diamond 1 \end{pmatrix}$$

305 In the picture, let the matrix on the left side be the pattern P and let M
306 be the other matrix. The entry in the square can be mapped to the element in
307 the square and the same holds for entries in the circle but it is not a sufficient
308 condition for the entry in the triangle to be mapped to the element in the triangle.

309 4. Improvements to basic 310 algorithms

311 4.1 General pattern

312 4.1.1 Improving memory consumption

313 As the algorithm for testing avoidance of a general pattern was described in
314 [chapter2], it creates all possible partial mappings and checks whether at least
315 one can be extended to a full mapping. Note that to compute all the partial
316 mappings of some level l , it only uses mappings of level $l - 1$; therefore, it is
317 enough to only store partial mappings of two levels in memory.

318 In [chapter2] we also introduced the idea of (un)important lines for a partial
319 mapping of level l and equivalence based on not using unimportant lines at all -
320 as they are fully bounded by other already mapped lines. When a line becomes
321 unimportant it stays unimportant till the end of the run; as a result, we can forget
322 where we mapped those lines to save memory. This is not as big of a deal as the
323 previous observation was but note there are cases of patterns in which each line
324 becomes unimportant just two levels after it gets added.

325 4.1.2 Not mapping empty lines

326 An empty line is a row or a column that does not contain any one-entries. Such
327 a line can be mapped to any line and if the algorithm leaves space for it (which
328 it does), we do not need to map it at all.

329 4.1.3 Using the last changed position

330 As the MCMC process works, it always changes one element of the big matrix
331 and asks whether it still avoids the pattern. If it does not and we know that
332 before the change it did, we are sure the changed element $[r, c]$ is a part of the
333 pattern. It is hard to use this fact in the algorithm. It just maps one line after
334 another and we do not know at the beginning to which line the changed position
335 lines should be mapped.

336 What we can do is to enforce that neither the r -th line nor the c -th one
337 get skipped. We will only look at the restriction for rows. The restrictions for
338 columns are symmetrical. There are three situations we want to avoid:

- 339 • The first row of P is mapped under the r -th row. This prevents any other
340 row to be mapped to r -th one and we don't want that.
- 341 • The last row of P is mapped above the r -th row. This again prevents any
342 other row to be mapped to r -th one.
- 343 • Two adjacent rows $l, l + 1$ of P are mapped to $L < L'$ respectively and
344 $L < r < L'$ which leaves no other row to be mapped to r .

345 4.1.4 Line order

346 An important thing, if we want the algorithm to run fast, is to choose a good
347 line order. A line which is unimportant in level l in a line order may easily be
348 important till the nearly last level in a different order.

349 We choose line order to hopefully enforce two things:

- 350 • Make as many unimportant lines as possible. This really allows the equiv-
351 alence based improvements to kick in. The more lines are unimportant the
352 more mappings become equivalent and the faster it is to iterate through all
353 of them.
- 354 • Recognize hopeless partial mappings as soon as possible. A partial mapping
355 gets extended if the line does not break the rule that there is a one-entry
356 where it needs to be. If we map all the rows first, the rule will get broken
357 only after we start to map columns and we probably want to find out sooner.

358
359 In the program a user can either choose their own custom order or one of four
360 algorithms with different main purposes:

- 361 • AUTO - this one tries the other three line orders and chooses the one which
362 shows the best performance over some iterations on a matrix. While this
363 may sound like a good thing to use, it is only so if an initial matrix is chosen
364 and it takes a lot of time since a lot of iterations need to be made in order
365 to make a good sample. I would recommend not to use AUTO order at all
366 and instead to try all the line orders by hand with a number of iterations
367 depending on the pattern and a good initial matrix; for instance, generated
368 with a smaller number of iterations on the same pattern and with any line
369 order.
- 370 • DESC - the lines are ordered in descending order depending on the number
371 of one-entries. This follows the idea to start with the lines that are the
372 hardest to map. Note that this algorithm does poorly if there are a lot of
373 lines with the same number of one-entries (for example an identity matrix).
- 374 • MAX - it orders the lines so that the maximum number of important lines
375 throughout the levels is as small as possible. This focuses straightforwardly
376 to having many unimportant lines, which the program does not remember.
- 377 • SUM - it orders the lines so that the sum of the numbers of the important
378 lines is the smallest possible throughout all levels. The purpose is the same
379 as in the MAX order and quite often it is the case both approaches produce
380 the same order.
- 381 • TWO - it orders the lines so that the maximum number of important lines
382 in two consecutive levels throughout all the levels is as small as possible.
383 This again focuses to having many unimportant lines, which the program
384 does not remember. The constant two is chosen due to the fact general
385 pattern always stores two levels of partial mapping at a time.

386 4.1.5 Mapping approaches

387 The one thing the approaches we will introduce have in common is that they try
 388 to recognize those partial mappings that have no chance to be extended to a full
 389 mapping as early as possible.

390 While the algorithm introduced in [chapter2] finds out the partial mapping
 391 is invalid only at the time it maps two lines having a one-entry at their inter-
 392 section to two lines having a zero-entry at the intersection, different approaches
 393 try to reveal the fact we would end up in the situation earlier by checking more
 394 conditions.

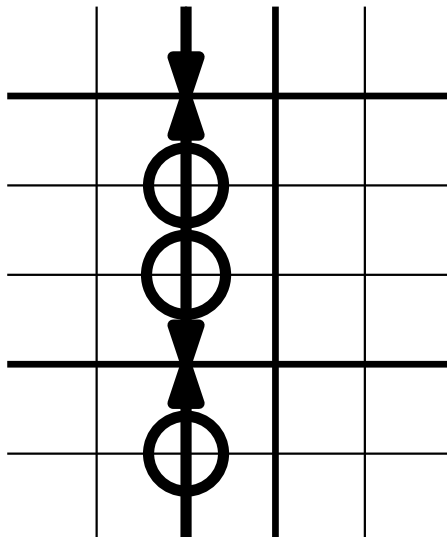
395 Imaging this as the pattern, where $0, \dots, 8$ are indices of the lines:

$$\begin{array}{c}
 \begin{array}{cccc}
 & 5 & 6 & 7 & 8 \\
 0 & \left(\begin{array}{c} 1 \\ 1 \\ 1 \\ 1 \end{array} \right) & \left(\begin{array}{c} 1 \\ 1 \\ 1 \\ 1 \end{array} \right) & \left(\begin{array}{c} 1 \\ 1 \\ 0 \\ 0 \end{array} \right) & \left(\begin{array}{c} 1 \\ 0 \\ 0 \\ 1 \end{array} \right) \\
 1 & & & & \\
 2 & & & & \\
 3 & & & & \\
 4 & & & &
 \end{array}
 \end{array}$$

396 We are in a situation when only lines 0, 3 and 7 are mapped and line 6 is
 397 currently being mapped. Here are some mapping approaches:

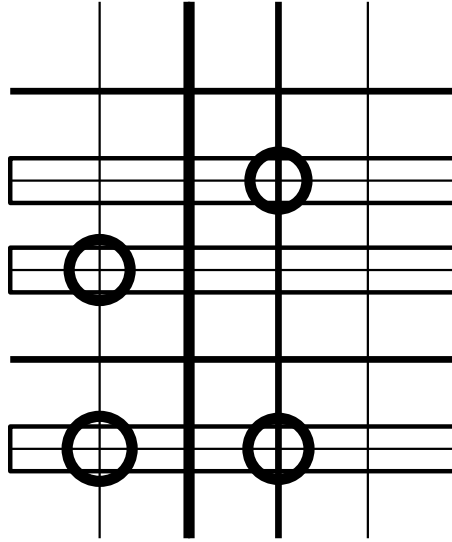
398 Enough one-entries

399 In the situation above, we not only want to check there are one-entries at the
 400 intersections of line 6 with lines 0 and 3, but we also check if there are enough
 401 one-entries in lines between where lines 0 and 3 are mapped so that there is a
 402 hope we can map lines 1 and 2 there and if there is a one-entry below the line
 403 where line 3 is mapped so we can map line 4 there later.



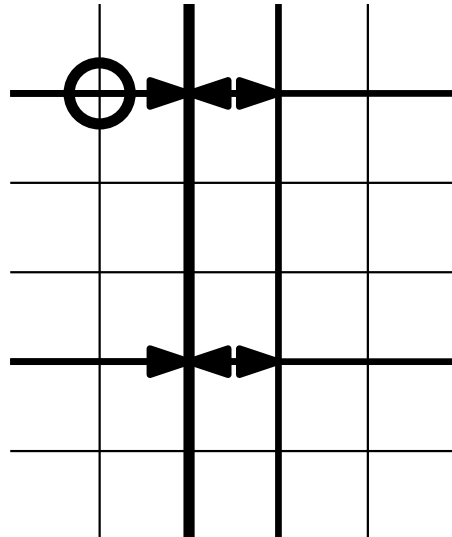
404 Recursive mapping

405 While we were only testing whether there are enough one-entries in between
406 already mapped lines in the previous approach, this time, we also check whether
407 those one-entries can be used for the lines that are intended to be mapped there.
408 For example, when we check there is a one-entry to be used for line 1 later, we
409 also check the line 1 can be mapped to that row, which in this situation means
410 to also check there is a one-entry at the intersection with the line to which the
411 line 7 is mapped.



412 Orthogonal bounds

413 When adding line 6, we check whether there are enough one-entries on the already
414 mapped lines orthogonal to line 6, in between line 6 and the closest mapped lines
415 next to line 6. The same idea as in “Enough one-entries”, but checking different
416 lines.



417 Usage

418 These restrictions on the added lines are not a fixed part of the program. A user
419 can decide which approaches they want to use in the configuration file. This is

420 due to the fact that there is no right path to choose.

421 In the testing that was done for a fixed pattern, we found out it is useful to
422 use all the mentioned restrictions when generating a matrix of size 100×100 , as
423 it turned out to be much faster than without the restrictions. On the other hand,
424 in the same test for a generated matrix of size 500×500 , it was much better not
425 to use any of those restrictions.

426 4.1.6 Using the whole structure in the next iteration

427 It may seem like a good idea to remember all the partial mappings, to propagate
428 them to the next iteration of the MCMC process and alter them depending upon
429 the change.

430 This really can be done. If the change is from zero-entry to one-entry, for each
431 partial mapping we already have we want to try to extend it by the line that just
432 changed and if we manage to do that we then try to extend it to a full mapping
433 in all possible ways if it is a new mapping or do nothing if it is equivalent with a
434 partial mapping of higher level. This can be easily done by means already used
435 in the standard algorithm and may lead to a better performing one.

436 However, if the element gets changed from one-entry to zero-entry we need
437 to go through the partial mappings and delete those that used the currently
438 changed one-entry. This gets a bit messy as we can no longer forget unimportant
439 lines and moreover for each partial mapping we need to remember how many
440 partial mappings of the previous level can be extended to that one, to delete that
441 mapping from the list if there are no longer any mappings extensible to that one.

442 This can all be done, but the whole thing comes with three huge inconven-
443 niences:

- 444 • Memory consumption - there can be a LOT of partial mappings and we
445 need to remember them all. Of course we can still use the equivalence but
446 we need to remember mappings of all levels.
- 447 • The change from one-entry to zero-entry is no longer for free. If this change
448 is done, we already know the pattern is not contained in M , but we still
449 need to do a lot of work to change the structure in order to use it in the
450 next iteration.
- 451 • Reverting - if the change is unsuccessful (the pattern is contained) we need
452 to revert the change which means to completely revert all changes we did to
453 the list of partial mappings. This can be either done by making a backup
454 copy of the whole structure and override the structure if needed, which again
455 is very costly as the structure is huge, or we can remember what partial
456 mappings are new and we go through all partial mappings and remove those
457 new ones. This again means to iterate through the big structure one more
458 time.

459 After realizing these issues it no longer looks useful to me to implement this
460 version of the algorithm.

4.2 MCMC parallelism

To speed up the computations, it is often possible to use parallelism. In this section, we show how to make the MCMC generator parallel, while still allowing both types of the pattern.

While the serial MCMC generator in each iteration changes one element in the generated matrix and checks whether it still avoids forbidden patterns, the parallel one makes several iterations at once, one on each copy of the generated matrix. This means that while iteration x is being computed by a thread, iteration $x + 1$ can at the same time be computed by a different thread. The only issue is that iteration $x + 1$ does not know what is going to be the state of the generated matrix at the time it should start. It expects iteration x to fail - not change the generated matrix at all, counting on the fact at some point it is unlikely a change does not create a mapping of the pattern, and starts with the same matrix as iteration x . If iteration x succeeds though, then the computed iteration $x + 1$ is invalid and the iteration is going to be recomputed again, starting with the altered matrix.

When the parallel version of MCMC generator is chosen and it is assigned n threads, it creates $n - 1$ private copies of the generated matrix and assigns one thread, called worker, to each of them. The last thread, which we call the main thread and which has exclusive access to the master copy of the generated matrix, makes one change of a bit in each private copy of the matrix and makes the corresponding worker check the avoidance.

The job of a worker is only to check if its copy of the matrix still avoids the pattern when one bit is changed. On the other hand all synchronization is left to the main thread. As mentioned before, one iteration of the MCMC process can be recomputed several times. We still want the generator to satisfy the conditions we have for the Markov chain (more in [chapter1]) in order to generate a random matrix. To achieve that, if a computed iteration x succeeds (and changes the generated matrix), all the other computed iterations that would follow after the iteration x become invalid and they all have to be recomputed. The process ends when all iterations get computed.

From now on, we won't be talking about iterations but about tasks. A task is basically one iteration of the MCMC process. The usefulness of this notation comes with an ID - a number, unique per task, assigned to each task, starting with 1 and always increasing. For a pair of consecutive iterations x and $x + 1$ it will always be the case that if task a is the last task to compute iteration x (which means the iteration does not get recomputed ever again after) and task b is the last task to compute iteration $x + 1$, then the ID of a is lower than the ID of b . Also there is no point in which two different tasks would be computing the same iteration at the same time. If tasks with IDs $a < b$ computed the same iteration, it must have been the case an earlier iteration succeeded when task with ID a was computed and after it got removed, task with ID b was assigned to recompute.

At any point in time, we only consider those tasks, that are being computed or those that wait to be processed (not those that have been processed), which means the lowest ID of tasks we consider increases in time.

When a task ends and it has the lowest ID (we can always wait for the task

508 with the lowest ID) we do:

509 • if it fails:

510 – Do nothing - there is no change to propagate to the master copy of the
511 generated matrix and all the tasks with higher ID expected this task
512 to fail, which it did.

513 – This increases the lowest ID by exactly one, as the task we speak of
514 got processed.

515 • if it succeeds:

516 – The main thread propagates the change tested by the task to the
517 master copy of the generated matrix.

518 – All the rest of the task get removed as they all had a higher ID -
519 computed iterations that follow after the one just computed and they
520 expected the task to fail, which it did not.

521 – This increases the lowest ID by more then one, because there are tasks
522 that got removed and one that got processed.

523 4.2.1 Example of the MCMC process for n threads

524 At first, iterations 1 to $n - 1$ are assigned one to each worker as tasks with ID 1 to
525 $n - 1$ with the same order as the order of iterations. If iteration 1 is not successful
526 (which all the other iterations count on), everything is alright. However, if the
527 iteration (its task) is successful, all the results of other tasks (and some of them
528 might have been already finished) are cleared and those iterations get recomputed
529 in tasks n to $2n - 3$ and the worker that computed task with ID 1 is assigned a
530 new task with ID $2n - 2$ - to compute iteration n . The result of the task gets
531 propagated to the master copy of the generated matrix only if all the tasks n to
532 $2n - 3$ fail, else is gets recomputed. This is what happens till the end.

533 4.2.2 Speculative computing

534 It may easily happen that a task not having the lowest ID ends first. In that
535 case we could just wait until it has the lowest ID and process it later. This is
536 not a very efficient approach. Instead we process the task immediately, but we
537 don't propagate the changes to the master copy of the generated matrix until all
538 tasks with lower ID fail and we do not stop the workers processing a task with
539 lower ID. If the task succeeds we remove all the changes computed by tasks with
540 higher ID and override their private copy of the generated matrix. Also it might
541 happen a task with even lower ID succeeds as well. This leads to more and more
542 overriding. Luckily this is the only precarious situation we may encounter and it
543 is not that hard to deal with it, even without copying the possibly huge generated
544 matrix all over the place.

545 The way we deal with these inconveniences is described in [chapter5] and
546 should be clear from the code itself.

547 4.2.3 Reverting and synchronizing in the main thread

548 The speculative computing discussed above is not the only improvement we can
549 make. It turns out to be costly to wake a thread so it computes a trivial function,
550 sets a few atomic variables and falls asleep again. This happens a lot in the
551 MCMC process. Every time a task succeeds it makes other workers revert the
552 changes they computed and synchronize the successful change, which are both
553 trivial functions.

554 To workaround this problem we make a theoretically bad decision which has
555 very nice practical results. All the reverts and synchronizations are computed by
556 the main thread instead of by an appropriate worker. There is no problem with
557 concurrency because the worker is always asleep when a task is to be assigned and
558 using the fact those tasks are really trivial, it does not make the rest of threads
559 wait for the main thread for too long while it computes changes.

560 4.3 Walking pattern

561 While the brute force implementation of an avoid algorithm for a general pattern
562 was improved heavily, the algorithm for a walking pattern is very fast in its nature
563 and cannot be much better. Or can it be?

564 4.3.1 Using the last changed position

565 As the MCMC process works it always changes one element of the big matrix and
566 asks whether it still avoids the pattern. If it does not and we know that before the
567 change it did, we are sure the changed element is a part of the pattern. Knowing
568 that and using the same inductive proof as we did in the proof of correctness of
569 the avoid algorithm (see [chapter2]) it is sufficient to only recompute the part of
570 the inner structure under the changed element and check if the last entry of the
571 pattern can be found there.

572 Not only that. We also know, using the fact the structure was completely
573 correct before the change, that if the values of both c_v and c_h of an element did
574 not change, the element won't cause the element underneath it to change and we
575 no longer have to recompute the other parts of the structure.

576 To use both these facts we replace the cycle through the diagonals by a simple
577 queue, starting at the position of the last changed element and putting more
578 positions in if the values of c_v or c_h are different than they were before. The
579 function ends either when the pattern was discovered or when the queue becomes
580 empty.

581 4.3.2 Lazy avoid

582 Lazy avoid is a variant of avoid function used when the MCMC parallelism is
583 chosen. While all the other types of patterns have a trivial implementation of
584 revert function, when using the walking pattern the inner structure needs to be
585 modified even when reverting. The MCMC parallelism turned out to work much
586 better if the revert calls are handled by the main thread (more in [chapter4]) and
587 it requires the function to run as fast as possible so the other threads are not

588 blocked by the call for too long. That is a reason why functions lazy revert and
589 lazy avoid were created.

590 The avoid function expects the inner structure of the walking pattern (see
591 [chapter3]) to be in a valid state and that requires some effort. To make lazy
592 revert the fastest possible, we postpone the work until the next call of lazy avoid,
593 meaning that lazy avoid then needs to do more things at once. It is no longer
594 sufficient to only compute the submatrix under the position changed last as we
595 did above, but it needs to also compute changes in the positions changed in those
596 lazy revert calls that are postponed.

597 We discuss several approaches, starting with the easiest one and ending with
598 the one that is fast and used in the final implementation.

599 **Recompute the whole structure every time**

600 The easiest way how to implement lazy avoid would be to always recompute the
601 whole inner structure. In that case we do not worry which positions are correct
602 and which are not, because every time we find the pattern, we recomputed all
603 the entries that form it, so we know it really is there. On the other hand, if we
604 manage to recompute the whole structure without finding the last entry of the
605 pattern, it just is not there.

606 The issue is efficiency. If the whole structure was correct and there was a
607 change of the last entry of the matrix it is sufficient to only recompute that one
608 entry. Instead we recompute a possibly very big structure. This results in a very
609 bad performance negating the advantage of parallel computation.

610 **Recompute only a part of the structure diagonal by diagonal**

611 A simple improvement would be to remember the changes done in previous calls
612 of lazy revert and together with the change done in lazy avoid call only recompute
613 the part of the structure that has possibly altered.

614 This gets a bit tricky when lazy avoid call actually discovers the pattern
615 because we cannot be sure the rest of the structure is in a correct order. It is
616 still possible to remember some horizontal, vertical and diagonal bounds and use
617 them to restrict the recomputed part of the matrix. The improvement is not that
618 significant though and we can do better.

619 **Queue of positions to recompute**

620 A different approach is closer to the one used in a standard avoid function. In-
621 stead of going through diagonal one after another, we have a queue of entries-to-
622 recompute. It is no longer sufficient to have a standard queue since in different
623 calls of lazy revert/avoid we can possibly change an entry of different priority (the
624 higher the more important) so we need to have some kind of a priority queue.
625 That is exactly what I tried.

626 Using `std::priority_queue` the function had no more problems with recomput-
627 ing the entries that were not influenced by the changes and used all the benefits
628 mentioned in the previous section. But the container does not come for free and
629 in the end I found out the price I paid for the operations on the priority queue
630 made the whole implementation comparably slow as in the previous attempt.

631 **Two leveled queue of positions to recompute**

632 The final solution comes with the same idea, but a different storage type. As
633 the priority depends upon a diagonal (two entries on the same diagonal can be
634 recomputed in any order) we only remember a priority queue of diagonals and an
635 array of diagonals saying whether a diagonal is already a member of the priority
636 queue. As far as the entries are concerned for every diagonal we have a `std::vector`
637 of entries-to-recompute as well as an array saying whether an entry is already a
638 member of the vector. So finally it is the case that the storage used is not only
639 good theoretically but as the numbers say, also practically. [reference to a table
640 of measurements or something]

641 5. Technical documentation

642 In this chapter, we cover those parts of the algorithm that may be hard to un-
643 derstand just from the code. This only means functions that are technically hard
644 - functions with unexpected dependencies, side effects and so on. Algorithmic
645 difficult tasks are explained in [chapter4].

646 5.1 General pattern

647 The general pattern class contains a lot of function. Most of them are easy to
648 follow and they all should be commented enough in the code. The only part
649 which deserves more attention is the constructor.

650 5.1.1 Construction

651 In the constructor of a general pattern, there are a few function that are easy in
652 nature but as they somehow use each other it is hard not to lose track of their
653 dependencies and results. In order to make this part of the code, which is a very
654 important part indeed, more understandable, we go through the constructor and
655 explain all that is happening in the order it is happening in.

656 Storing the pattern

657 The first thing, which is done right after initialization of variable, is storing the
658 pattern. Instead of storing the pattern in a `Matrix<bool>`, I decided to use to
659 store lines into a number, where in the binary coding a one-entry in the position
660 i means there is a one-entry in the line at the intersection with i -th orthogonal
661 line. This comes handy when computing lines orders. At the same time we also
662 find those lines that are empty (more in [chapter4]) and remember them, because
663 we do not have to map them at all.

664 Choosing the line order

665 After that we need to choose the right line order (again more in [chapter4]).
666 To compute MAX or SUM order we just use a brute force algorithm that checks
667 sequences of line adding and for each it computes how many lines are unimportant.
668 Then it just chooses the order which is the best in chosen metric.

669 What to remember

670 In the next step, we find what do we need to remember in each level of partial
671 mappings with respect to chosen order. As mention earlier, for MAX or SUM
672 order it is already computed when finding the order, but for other variants of
673 orders it is not, so we just compute it every time. What to remember is based
674 on the equivalence introduced in [chapter2] and the decision not to remember
675 unimportant lines (which we explained in [chapter4]).

676 **Parallel bound indices**

677 Now comes the hardest to follow part - precomputing the indices for searching for
678 parallel bounds. The idea behind is simple. When we are adding a new line and
679 we already have a partial mapping, it restricts to where we can add the line. For
680 example, if there are three rows in the pattern and the rows 1 and 3 are mapped,
681 the second one need to be mapped in between those two. The question is, where
682 are those two lines mapped to? First we add in a chosen order and second we do
683 not remember all lines, as some are unimportant. What do we want is to have
684 a instant access to the index of the line, which bounds added line, in the partial
685 mapping so we do not need to compute the index over and over again. That is
686 exactly what gets computed when the function “find_parallel_bound_indices” is
687 called. The series of other function calls follows just because we compute the
688 indices for all added lines in the order in which they are going to be added.

689 **Extending order**

690 The last function, “find_extending_order” just specifies how the next partial map-
691 ping will look like a from where from the previous mapping the values will be
692 copied. Again, unimportant lines play their role here and it may easily be the
693 case from a partial mapping storing k lines, after mapping one more line, we end
694 up with a partial mapping only storing $k - 1$ lines, because two lines become
695 unimportant by adding the line.

696 **5.2 Parallel computing**

697 **5.2.1 MCMC parallelism**

698 While the idea behind MCMC parallelism is described in [chapter4.3] and the
699 code is heavily commented, the work done by the main thread may still be hard
700 to understand.

701 Let I be the ID the process is currently waiting for, that is, the lowest ID of
702 a task that is being tested by a worker. In a structure called “queue” (which is
703 `std::vector<std::deque>`) each worker has a queue of tasks related to it. In the
704 queue, there are tasks that are either being computed or have been computed.
705 The history of tasks is needed to allow reverting changes that should have not
706 happen when the main thread encounters a different successful task with lower
707 ID. There is no need to have a complete history of all tasks computed. There
708 are only those tasks, that have higher ID than I or have lower ID, but those are
709 going to be removed from the “queue” as soon as possible. The name “queue” is
710 not random, it describes the order in which the tasks are being stored - the tasks
711 with lower ID have been inserted earlier and therefore they are at the bottom.

712 Now that we know the most important structure let's see how the main thread
713 works with that and what are the situations.

- 714 • `pop_front`: The main thread deletes the first tasks (the one with the lowest
715 ID) if one of two things happen:
 - 716 – The ID of the task being deleted is equal to I . That means the change
717 computed by the task is being propagated to the generated matrix and

718 there is no need to remember the task anymore. This also increases I ,
719 not necessarily by one.

720 – The ID of the task being deleted is less than I . This situation is due
721 to synchronization. The worker was supposed to synchronize a task
722 computed by a different worker that did not have the lowest ID at the
723 time. Therefore the task needs to be in the list of tasks so we can
724 revert it if needed. If there is no need to revert it and the lowest ID
725 gets greater or equal to the ID of the task, we can just delete it from
726 the “queue”.

727 • `pop_back`: There is only one reason to delete tasks from the end of the
728 “queue” and that is reverting. Imagine there is a task with id J at the
729 end of the “queue”. Now a different worker computes a task with lower
730 ID and finds out the change is successful. This means the task J won’t
731 propagate to the generated matrix and there is no use for it. If it is still
732 being computed, we cannot do much about it, so we just tell the worker to
733 stop computing and deal with it later. If the task is finished, we need to
734 revert it, but only in case the task was successful, because if it was not, it
735 had already been reverted by the worker. So we revert the task if needed
736 and we can just delete it from “queue” as it will never be used.

737 • `emplace_back`: The main thread only inserts new tasks to the end of the
738 “queue” and there are two reasons to insert:

739 – Worker is assigned a completely new task to check the avoidance. In
740 this situation the task is given a new, globally highest ID and we add
741 the task at the end of the list.

742 – The second reason to insert into “queue” are, again, synchronizations.
743 The situation is the same as it was in the case, when we `pop_back` -
744 after we revert all the tasks in the list, we need to synchronize changes
745 that forced reverting and if their ID is not lower or equal to I , we need
746 to add them to the list so they can be reverted if needed.

747 5.3 Library interface

748 6. User documentation

749 In the last chapter of the thesis we show how to make the program generate
750 random matrices or to test whether a certain matrix avoids a given forbidden
751 pattern. In the first section we find out how to set a configuration file. After that
752 we go through necessary input and output files and see how they are formatted.

753 6.1 Configuration file

754 In order to modify what the program computes, we use a configuration file. There
755 is only command line options to be used and that is a path to the configuration
756 file. If no path is inserted the configuration file is expected to be located in the
757 same directory as the executable file is and its name is “config.txt”.

758 The file is a standard text file which can be modified by any text editor and
759 is structured into four sections:

- 760 • input
- 761 • pattern
- 762 • output
- 763 • statistics

764 The order of the sections is not fixed and there can be additional empty lines for
765 better readability. In each section there is a list of values that can be set either
766 to arbitrary value or to a specific one. There is at most one command of format
767 “option=value” per line and there might be additional white spaces surrounding
768 the “=” sign.

769 If an option is set more than once, the latter value is always used. If, on the
770 other hand, an option is not set at all, the default value is used. If there is a
771 line encountered that sets a wrong option, for instance when the user mistypes a
772 valid option, the line is skipped and the user gets a warning in the standard error
773 output.

774 Let us provide a list of all options for each section together with their default
775 values.

776 6.1.1 Input

777 In the first section of the configuration file we set the generating process.

- 778 • size: The size of the generated matrix. Results in $M \in \{0, 1\}^{size \times size}$.
Possible value: $s \in \mathbb{N}$
779 Default value: 100
- 780 • iterations: The number of iterations of the MCMC process.
Possible value: $i \in \mathbb{N}$
781 Default value: 10,000

- 782 • `random_seed`: The random seed for the MCMC process.
 Possible value: $s \in \mathbb{N}$
 “random” - chooses a random seed
 783 Default value: “random”
- 784 • `init_matrix`: A $size \times size$ matrix the MCMC process starts with.
 Possible value: *input file path*
 “zero” - a matrix containing no one-entries
 785 Default value: “zero”
- 786 • `parallel_mode`: choose whether to compute in parallel or serial.
 Possible value: “serial”
 “mcmc” - more iterations of the MCMC generator are tested in parallel
 787 “map” - more partial mappings are being extended in parallel
 Default value: “serial”
- 788 • `threads_count`: The number of threads that are going to be used if a parallel
 789 mode is chosen.
 Possible value: $t \in \mathbb{N}$
 -1 - chosen according to the number of cores
 790 Default value: 1

791 6.1.2 Pattern

792 It is time to set the options that matter the most - matrix patterns. As you
 793 generate a matrix which avoids more than just one pattern, the section [pattern]
 794 can be used multiple times, specifying one pattern for each occurrence.

- 795 • `pattern_file`: A path to a input matrix file - the pattern.
 Possible value: matrix file path
 796 Default value: “input.txt”
- 797 • `pattern_type`: The type of the pattern. Determines the method used for
 798 testing avoidance.
 Possible value: “general”
 “walking” - see [chapter3]
 799 “slow” - brute force algorithm for a general pattern
 Default value: “general”

800 The next options are only useful if the general pattern type is chosen. It
 801 specifies how the mappings are stored as well as what the map function tests.

- 802 • `map_one_entries`: If set to “yes”, the map function not only checks that
 803 already mapped lines do cross the currently added one in a one-entry if
 804 needed but also tests whether there is enough one-entries in between those
 805 already mapped lines. The result is that the map function takes a little
 806 longer but it recognizes hopeless mappings earlier.
 Possible value: “yes”
 “no”
 807 Default value: “yes”

- 808 • `map_recursion`: If set to “yes” and the `map_one_entries` is also set to “yes”,
809 the `map` function not only tests there are enough one-entries in between
810 already mapped lines on the currently added one, but for those non-mapped
811 lines crossing the current one in a one-entry, it also tests if the partial
812 mapping has one-entries everywhere they are supposed to be. Although
813 this comes with a big time consumption it can filter a big portion of partial
814 mapping that would later become unfeasible.
 Possible value: “yes”
815 “no”
 Default value: “yes”

- 816 • `map_orthogonal_bounds`: If set to “yes”, the `map` function also tests the
817 orthogonal bounds of added line - see [chapter2].
 Possible value: “yes”
818 “no”
 Default value: “no”

- 819 • `map_container`: A container in which the partial mappings are stored.
 Possible value: “set” - `std::set` (red-black tree)
 “hash” - `std::unordered_set` (hash table)
820 “vector” - `std::vector` (dynamic array)
 Default value: “hash”

- 821 • `line_order`: Choose the order in which the lines are being added to the
822 partial mapping. See [chapter2]
 Possible value: “max”
 “two”
 “sum”
823 “desc”
 “auto”
 “order file path”
 Default value: “max”

824 6.1.3 Output

825 There is no reason to generate a matrix noone will ever see. In this section we
826 make the matrix show. As the matrix can be output to console, a text file or a
827 bmp file an option in the section can be set more than once and all the settings
828 make difference.

- 829 • `matrix_output`: The generated matrix can be output as a bmp file in which
830 one-entries are black pixels and zero-entries white. To do that, the file path
831 has to have a pattern “path.bmp”. If a different path is given the file is
832 stored as a matrix text file. It can also be output into a console if “console”
833 is set. In that case it has the text format.
 Possible value: “console”
 matrix bmp file path
834 matrix text file path
 “no”
 Default value: “no”

- 835 • `performance_stats`: If the serial computation or the map parallelism is cho-
836 sen, the program can output a statistics like the percentage of avoid call
837 success, how long did one call take on average and what was the average
838 size of structures. Note that this is not very useful as a matrix statistics,
839 see the next section for that. If more patterns are set, the statistics may get
840 misleading as it also counts the cases when the first pattern is contained in
841 the matrix and the other patterns are not tested at all.

842 Possible value: `"console"`
 performance file path
 `"no"`
842 Default value: `"no"`

- 843 • `performance_csv_stats`: The same information as above but formatted to a
844 csv file so the data can be more easily worked with.

845 Possible value: `"console"`
 csv file path
845 Possible value: `"no"`
845 Default value: `"no"`

- 846 • `time_to_console`: Prints how long the computation took into a console.

847 Possible value: `"yes"`
847 Possible value: `"no"`
847 Default value: `"no"`

- 848 • `patterns_to_console`: Prints all the used patterns into the console.

849 Possible value: `"yes"`
849 Possible value: `"no"`
849 Default value: `"no"`

850 6.1.4 Statistics

851 The last section handles the options important for scientists. While generating
852 a random matrix is a great result, on its way the program can also create some
853 statistics, namely make a histogram of occurrences of one-entries in a generated
854 matrix as the MCMC iterates as well as store the matrix with the highest amount
855 of one-entries. As the process usually does not start with a random matrix, the
856 user can decide to only compute the statistics after a certain number of iterations
857 has been done and to only check a small portion of iterations, every 10th for
858 instance, as a single iteration may not make any difference and counting the
859 histogram takes time.

- 860 • `histogram_frequency`: Sets how often the histogram gets refreshed.

861 Possible value: $f \in \mathbb{N}$
 0 - the histogram is not computed at all
861 Default value: 0

- 862 • `histogram_initial`: Sets the initial iteration of the MCMC process when the
863 histogram gets refreshed.

864 Possible value: $i \in \mathbb{N}$
864 Default value: 1,000

- 865 • histogram_final: Sets the last iteration of the MCMC process when the
866 histogram gets refreshed.
Possible value: $f \in \mathbb{N}$
867 -1 - the histogram is computed till the end
Default value: -1
- 868 • histogram_file: Sets where to output the histogram computed during the
869 MCMC process.
Possible value: matrix bmp file path
matrix text file path
870 "console"
"no"
Default value: "no"
- 871 • max_ones_matrix_file: Sets where to output the matrix that had the most
872 one-entries among all matrices iterated through during the MCMC process.
Possible value: matrix bmp file path
matrix text file path
873 "console"
"no"
Default value: "no"

874 6.2 File input

875 There are only two types of input files expected by the program. Either you want
876 to read a matrix file, which can be a pattern or an initial matrix, or an order file
877 that determines an order in which the lines are going to be mapped if the general
878 pattern is chosen.

879 6.2.1 Matrix file

880 A matrix file is a standard text file having the format as follows:

- 881 • 2 natural numbers specifying the number of rows and columns in this order.
- 882 • a sequence of zeros and ones of length rows×columns specifying the matrix
883 from the top left corner one row after another.

884 **Example:** 2 3
 1 0 1
 1 1 0

885 6.2.2 Order file

886 If you want to choose the order in which the lines are going to be mapped when
887 a general pattern is chosen, it is your responsibility to check that all lines that
888 need to be mapped are mapped. It is for example possible to only map three lines
889 even if the pattern consists of six lines just because there is for example no need

890 to map empty lines at all. Therefore the program does not check the validity of
891 the order and just uses it.

892 Now that the user has been warned, the format of the custom order file is
893 simple. It consist of the indices of the lines of the pattern numbered starting
894 with 0 and starting from the top row and ending with the right column.

895 One possible order for the matrix given as an example in [6.2.1] is this file:

896 2 1 0 3 4

897 First mapping the left column, the second and first row after that and finishing
898 the mapping with the middle column and the right one.

899 6.3 File output

900 Let us now find out what the output files look like.

901 6.3.1 Matrix text file

902 The matrix text file has the same format as the input one. It consists of:

- 903 • 2 natural numbers specifying the number of rows and columns in this order.
- 904 • a sequence of zeros and ones of length rows×columns specifying the matrix
905 from the top left corner one row after another.

906 The matrix is binary except for the one produced as a histogram, which can have
907 higher natural numbers and contains the number of samples as the last number.
908 If you then divide all the entries by the last number, you get a percentage of the
909 entry being a one-entry.

910 6.3.2 Matrix bmp file

911 For an $n \times n$ matrix the standard bmp file contains $n \times n$ pixel of black color
912 meaning a one-entry and a white color for a zero-entry. If the histogram is output
913 as a bmp file, the pixels are greyscaled and the darker a pixel is the more often
914 the entry was a one-entry during the MCMC process.

⁹¹⁵ Conclusion

916 Bibliography

917 Neal Madras. *Lectures on monte carlo methods*, volume 16. Springer Science &
918 Business, 2002.

919 Neal Madras and Hailong Liu. Random pattern-avoiding permutations. *Algo-*
920 *rithmic Probability and Combinatorics, AMS, Providence, RI*, pages 173–194,
921 2010.

922 List of Figures

923	1	Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$	
924		satisfies all the conditions. On the other hand, matrix M_2 avoids	
925		P as there is no such mapping.	4
926	2.1	An example showing unimportant line and equivalent mappings. .	9
927	3.1	An example of a walk and the order of its entries.	10

928 List of Tables

929 List of Abbreviations

930 Attachments