

Charles University in Prague
Faculty of Mathematics and Physics

BACHELOR THESIS



Stanislav Kučera

Generating random pattern-avoiding matrices

Computer Science Institute of Charles University

Supervisor of the bachelor thesis: RNDr. Vít Jelínek, Ph.D.

Study programme: Computer Science

Study branch: General Computer Science

Prague 2016

I declare that I carried out this bachelor thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University in Prague has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In date

signature of the author

Title: Generating random pattern-avoiding matrices

Author: Stanislav Kučera

Department: Computer Science Institute of Charles University

Supervisor: RNDr. Vít Jelínek, Ph.D., department

Abstract: Binary matrices not containing a smaller matrix as a submatrix became an interesting topic recently. In my thesis, I introduce two new algorithms to test whether a big square binary matrix contains a smaller binary matrix together with a process using randomness, which approximates a uniformly random matrix not containing a given matrix. The reason to create such algorithms is to allow researchers test their conjectures on random matrices. Thus, my thesis also contains an effective cross-platform implementation of all mentioned algorithms.

Keywords: binary matrix pattern-avoiding Markov chain Monte Carlo

I would like to thank my supervisor, Vít Jelínek, for his patience and willingness to spend so much time helping me improve the thesis. I also thank to Paul Macklin, author of EasyBMP (<http://easybmp.sourceforge.net/>) – a library for working with bmp files, which I use in my program.

Contents

Preface	3
Introduction	5
1 Markov chain Monte Carlo	6
1.1 Markov chains	6
1.2 Markov chain for pattern-avoiding binary matrices	6
2 An algorithm for testing pattern-avoidance of a general pattern	8
2.1 Sketch of a brute force algorithm	8
2.2 Equivalent mappings	8
3 An algorithm for testing pattern-avoidance of a special pattern	10
3.1 Walking pattern	10
3.2 Dynamic program	10
3.2.1 Inner structures	11
3.2.2 The algorithm	11
3.2.3 Correctness	12
3.2.4 Generalization	13
4 Improvements to basic algorithms	15
4.1 General pattern	15
4.1.1 Improving memory consumption	15
4.1.2 Not mapping empty lines	15
4.1.3 Using the last changed position	15
4.1.4 Line order	16
4.1.5 Mapping approaches	17
4.1.6 Using the whole structure in the next iteration	19
4.2 MCMC parallelism	20
4.2.1 Example of the MCMC process for n threads	21
4.2.2 Speculative computing	22
4.2.3 Reverting and synchronizing in the main thread	22
4.3 Walking pattern	22
4.3.1 Using the last changed position	22
4.3.2 Lazy avoid	23
5 Technical documentation	25
5.1 Classes and API	25
5.1.1 Matrix	25
5.1.2 Pattern	25
5.1.3 Statistics	25
5.2 General_pattern	25
5.2.1 Construction	26
5.3 MCMC parallelism	27

6	User documentation	29
6.1	Installation	29
6.1.1	Windows	29
6.1.2	Unix, Linux, iOS	29
6.2	Configuration file	29
6.2.1	Input	30
6.2.2	Pattern	30
6.2.3	Output	32
6.2.4	Statistics	32
6.3	File input	33
6.3.1	Matrix file	34
6.3.2	Order file	34
6.4	File output	34
6.4.1	Matrix text file	34
6.4.2	Matrix bmp file	35
	Conclusion	36
	Bibliography	37
	List of Figures	38
	List of Tables	39
	List of Abbreviations	40
	Attachments	41

1 Preface

2 Theses at the Faculty of Mathematics and Physics of Charles University in Prague
3 usually fit into one of three categories:

- 4 1. Theoretical thesis
- 5 2. Experimental thesis
- 6 3. Implementation thesis

7 My thesis does not fit entirely into only one category and it does not try to. The
8 project consists of several similarly important parts which are:

- 9 • Design of algorithms for generating a special binary matrix
- 10 • Making the algorithms run fast on inputs that are usual for researchers
- 11 • Implementing the algorithms to provide practical tool

12 One point would not make sense without others, but together the thesis may
13 become a very useful tool for scientists interested in matrices with forbidden
14 patterns as the thesis provides with a process of generating random pattern-
15 avoiding matrices.

Introduction

We let $M \in \{0, 1\}^{n \times m}$ denote a *binary matrix* of size n by m . The *height* of M , denoted by n , is the number of rows of M and m is its width (the number of columns). A *line* of a matrix is one of its rows or columns and we denote by $L(M)$ the ordered set of all lines of M . Its order is given by the standard indexing of rows and columns.

Definition 1. We say a binary matrix M contains a binary matrix P , which we call a “pattern”, as a submatrix, if there is a mapping $f : L(P) \rightarrow L(M)$, such that

- $l \in L(P)$ is a row of P iff $f(l) \in L(M)$ is a row of M
 - $\forall l, l' \in L(P) : l < l' \Rightarrow f(l) < f(l')$ (preserves the order)
 - $\forall l, l' \in L(P) : \text{if lines } l \text{ and } l' \text{ intersect and there is a one-entry at the intersection, then there is a one-entry at the intersection of } f(l) \text{ and } f(l').$
- otherwise, it avoids the pattern P .

$$P = \begin{matrix} & \begin{matrix} 2 & 3 \end{matrix} \\ \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \end{matrix} \quad M_1 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix} \quad M_2 = \begin{matrix} & \begin{matrix} 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \end{matrix} & \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix} \end{matrix}$$

Figure 1: Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$ satisfies all the conditions. On the other hand, matrix M_2 avoids P as there is no such mapping.

The interesting cases are square matrices of size n by n , where n is big (going to infinity) and the size of a pattern (not necessarily square matrix) is small (constant). Even for a constant size forbidden pattern it is hard to determine the number of matrices of size n that avoid it or to characterize, what properties they have. Sometimes we consider matrices avoiding more than just one forbidden pattern, in which case we denote the set of all forbidden matrices by \mathcal{P} . When a matrix avoids \mathcal{P} , it avoids every $P \in \mathcal{P}$.

Definition 2. We denote by $\mathcal{M}_n(\mathcal{P})$ a set of all binary matrices of size n by n avoiding \mathcal{P} as submatrices.

Definition 3. We always call M the square binary matrix for which we test the containing and P the pattern (if there is only one) that is being tested. Moreover, we denote by h the height (the number of rows) of P and by w its width.

The area of pattern avoidance has been heavily studied for permutations and it also becomes more popular for their generalization - binary matrices. In most of the areas in combinatorics it is useful to explore properties of random objects and a lot of attention is directed towards random matrices when considering pattern avoidance. The goal of the work is, for given $n \in \mathbb{N}$ and set of forbidden patterns \mathcal{P} , to generate a uniformly random $M \in_R \mathcal{M}_n(\mathcal{P})$.




Figure 2: Example of a generated matrix avoiding I_{10} (unit matrix). Black dots are one-entries and white are zero-entries. As you can see, matrices avoiding a pattern can have a nice structure.

48 Generating random matrix

49 One way to get $M \in_R \mathcal{M}_n(\mathcal{P})$ is to choose a matrix of required size completely
 50 at random, for such, test whether it avoids the pattern and simply repeat the
 51 process until we find one, which does. However, in the most interesting cases,
 52 only a small fraction of all matrices avoid the pattern and the process takes too
 53 long, to be practically useful.

54 For generating random permutations avoiding forbidden pattern, a different
 55 technique was introduced in Madras and Liu [2010]. It uses a randomized pro-
 56 cess called Markov chain Monte Carlo, which we will abbreviate by MCMC. It
 57 is an iterative process, which for a well chosen Markov chain (more in Chapter
 58 1) approximates a random object. The algorithm by Madras and Liu was devel-
 59 oped for permutations (permutation matrices) and it cannot be used for general
 60 matrices. In Section 1.2 we show how to adapt the algorithm, which will lead
 61 us to a MCMC algorithm that approximates $M \in_R \mathcal{M}_n(\mathcal{P})$. To produce a good
 62 approximation the process needs to do a lot of iterations and despite the fact
 63 it is unknown what is the mixing time (the number of iterations required) of a
 64 MCMC process, in practice, the method does better than the trivial algorithm.

65 Testing avoidance

66 In each step of our MCMC process we need to test whether a matrix avoids a
 67 pattern. We will show a very fast algorithm that only works for a special class of
 68 binary matrices (explained in Chapter 3) together with a slightly less performing
 69 algorithm for a general pattern, which, again, comes as a generalization of an
 70 algorithm for permutations from the article by Madras and Liu and is described
 71 in Chapter 2.

72 In Chapter 4 we improve both our algorithms and introduce a parallel version
 73 of MCMC process, which further increases the performance of matrix generating.

74 In Chapter 5 some technical details are explained to make reading the code
 75 easier for reader and to describe user interface. The last chapter (Chapter 6)
 76 contains user documentation.

1. Markov chain Monte Carlo

Our goal to generate $M \in_R \mathcal{M}(\mathcal{P})$ heavily depends on the theory of Markov chains. In this work we only define useful terms and state two important theorems. If you are interested in more details, see Madras [2002].

1.1 Markov chains

Definition 4. Let \mathcal{S} be a finite set of states and for every $i, j \in \mathcal{S}$ $p_{i,j}$ prescribed probability of a change of state from i to j . Also let X_0 be a random variable with values from \mathcal{S} . We call a sequence X_0, X_1, \dots , where $X_i \in \mathcal{S}$ for every i a Markov chain if

$$\Pr(X_{t+1} = j | X_t = i) = p_{i,j} \quad (i, j \in \mathcal{S})$$

Definition 5. A Markov chain is said to be symmetric if $p_{i,j} = p_{j,i}$ for every pair of states i and j .

Definition 6. A Markov chain is irreducible if the chain can eventually get from each state to every other state, that is, for every $i, j \in \mathcal{S}$ there exists a $k \geq 0$ (depending on i and j) such that $\Pr(X_k = j | X_0 = i) > 0$.

Definition 7. If an irreducible chain has $p_{i,i} > 0$ for some i , then it is aperiodic.

Let $p_{i,j}^{(k)} = \Pr(X_{t+k} = j | X_t = i)$ denote the k -step transition probabilities for $k = 0, 1, \dots$ and $i, j \in \mathcal{S}$. The transition probability matrix is $P = (p_{i,j})$.

Next we state two theorems allowing us to expect Markov chains to converge to a uniformly random state in \mathcal{S} even if the initial state X_0 is not random. Both theorem can be found in Madras [2002].

Theorem 1. Consider an aperiodic irreducible Markov chain with state space \mathcal{S} . For every $i, j \in \mathcal{S}$, the limit $\lim_{k \rightarrow \infty} p_{i,j}^{(k)}$ exists and is independent of i ; call it π_j . Furthermore, if \mathcal{S} is finite, then

$$\sum_{j \in \mathcal{S}} \pi_j = 1 \quad \wedge \quad \sum_{i \in \mathcal{S}} \pi_i p_{i,j}^{(1)} = \pi_j$$

for every $j \in \mathcal{S}$. That is, if we write π to denote the row vector whose entries are π_i , then $\pi P = \pi$.

Theorem 2. Suppose that an irreducible Markov chain on the finite state space \mathcal{S} is symmetric. Then the equilibrium distribution is uniform on \mathcal{S} .

1.2 Markov chain for pattern-avoiding binary matrices

To generate a binary matrix $M \in \{0, 1\}^{n \times n}$ avoiding patterns in \mathcal{P} , we create a Markov chain, whose states space is $\mathcal{M}_n(\mathcal{P})$. After sufficiently many iterations (m) of MCMC process we set $M := X_m \in \mathcal{M}_n(\mathcal{P})$. We always begin with an initial matrix X_0 and the process looks like this:

- 103 1. For $i := 1, 2, \dots, m$:
- 104 2. Set $X_i := X_{i-1}$.
- 105 3. Choose $r \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 106 4. Choose $c \in_R \{0, 1, \dots, n-1\}$ uniformly at random.
- 107 5. Flip the bit at $X_i[r, c]$.
- 108 6. If X_i contains \mathcal{P} , flip the bit back.

109 If the process starts with a matrix X_0 that avoids \mathcal{P} , then after every step it
 110 still avoids \mathcal{P} . Note that an iteration does not change the matrix if the condition
 111 6 is satisfied. We need to show the Markov chain we presented meets all the
 112 conditions of both theorems:

113 Symmetry

114 Imagine a sequence of bits flipping that changes the i -th matrix to j -th one. The
 115 reversed order of the same sequence changes the j -th matrix to the i -th one.

116 Irreducibility

117 As the steps go, it is easy to see we can with non-zero probability create any
 118 matrix $M_1 \in \mathcal{M}_n(\mathcal{P})$ from the zero matrix $0_n = 0^{n \times n}$ by choosing the one-entries
 119 of M_1 . When we can get from 0_n to M_2 by a sequence of flip changes, the reversed
 120 sequence is a sequence of steps from $M_2 \in \mathcal{M}_n(\mathcal{P})$ to 0_n . Thus, with non-zero
 121 probability we can always reach M_2 from M_1 ; therefore, the Markov chain is
 122 irreducible.

123 Aperiodicity

124 The Markov chain is irreducible so it suffices to show that there is an i for which
 125 $p_{i,i} > 0$. Clearly, there is a matrix for which there is at least one bit that cannot
 126 be flipped without creating a pattern (for example the one with the maximum
 127 number of one-entries) and this forces $p_{i,i} > 0$.

128 2. An algorithm for testing 129 pattern-avoidance of a general 130 pattern

131 In this chapter and Chapter 3 we show algorithms for testing whether a pattern
132 P is contained in a square binary matrix M .

133 We begin with a very basic algorithm, which we then improve a lot to get a
134 fast algorithm for testing avoidance of a general pattern.

135 2.1 Sketch of a brute force algorithm

136 Let $L = (l_1, l_2, \dots, l_{w+h-1})$ be a permutation of lines (rows and columns) of the
137 pattern P and $k \in [w + h - 1]$. *Partial mapping of level k* of lines of P is a
138 function f from $L' := \{l_1, l_2, \dots, l_k\} \subseteq L$ to lines of the big matrix M satisfying
139 three conditions:

- 140 • Both $l' \in L'$ and $f(l')$ are rows or they are both columns.
- 141 • If $l' \in L'$ and $l'' \in L'$ are both rows or columns and $l' < l''$, then $f(l') < f(l'')$.
142 This means partial mapping keeps the order of the lines.
- 143 • If $l' \in L'$ is a row of P and $l'' \in L'$ is a column of P and there is a one-entry
144 at the intersection of l' and l'' , then there is a one-entry at the intersection
145 of $f(l')$ and $f(l'')$.

146 The basic algorithm we use goes as follows. First it maps l_1 to all possible lines
147 of M , creating partial mappings of $\{l_1\} \subseteq L$. For $k = 2, \dots, w + h - 1$ it takes
148 each partial mapping from the previous iteration and extends it by adding line l_k
149 to the partial mapping in all possible ways. If we manage to map all the lines of
150 P , then M does not avoid it and if at some point there are no partial mappings
151 to extend it means M avoids P .

152 The algorithm can be improved in two ways. Firstly, we can try to recognize
153 unextendable partial mappings earlier than at the moment a line can no longer be
154 mapped, for example by counting whether there is enough one-entries in between
155 already mapped lines (more in Section 4.1.5). Secondly, which is going to be
156 fundamental for us, we can try not to remember more copies of different mappings
157 that can be extended in the same way.

158 2.2 Equivalent mappings

159 There is no need to remember two different partial mappings of the same level
160 if they can be both extended exactly the same way, because our function is only
161 supposed to check whether a pattern can be mapped to a big matrix not to find
162 all such mappings.

163 **Definition 8.** We call a line l of a pattern P important for chosen permutation
164 of lines of P , if one of the conditions is met:

- 165 • An adjacent line of the pattern has not been mapped yet.
- 166 • There is a one-entry on the line l at the intersection with line l' that has
- 167 not been mapped yet.
- 168 . Otherwise the line is unimportant for the permutation.

169 Whether a line is important or not only depends on the permutation, so if
 170 we have a line unimportant in a partial mapping of level k , it is unimportant in
 171 every partial mapping of level k .

172 At the beginning, when no line is mapped, all lines are important. After some
 173 lines get mapped, a line can become unimportant in the partial mapping as all
 174 lines that bound it are in the mapping as well. If a line is unimportant in a partial
 175 mapping of some level, it will stay unimportant in all extensions of the mapping
 176 we can find.

177 **Definition 9.** We say two partial mappings of the same level are equivalent if
 178 all important lines in the mapping of that level are mapped to the same lines of
 179 the big matrix in both mappings.

$$\begin{array}{c}
 \begin{array}{c}
 4 \quad 5 \quad 6 \quad 7 \\
 0 \left(\begin{array}{|c|c|c|c|} \hline 1 & 1 & 0 & 0 \\ \hline 1 & 0 & 1 & 1 \\ \hline 2 & 0 & 1 & 0 \\ \hline 3 & 0 & 1 & 1 \\ \hline \end{array} \right) \\
 \end{array}
 \quad
 \begin{array}{c}
 5 \quad 6 \quad 7 \quad 8 \quad 9 \\
 0 \left(\begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 0 & 1 & 0 & 1 \\ \hline 2 & 1 & 1 & 0 & 0 \\ \hline 3 & 0 & 1 & 1 & 0 \\ \hline 4 & 1 & 1 & 1 & 0 \\ \hline \end{array} \right) \\
 \end{array}
 \end{array}
 \quad
 M=$$

Figure 2.1: An example showing unimportant line and equivalent mappings.

180 For P and M , binary matrices in Figure 2.1, in partial mapping of level 4
 181 $f = \{(1, 1), (2, 2), (3, 4), (5, 6)\}$, line 2 is unimportant because both lines 1 and 3
 182 are mapped and so is line 5 - the only line to intersect line 2 in a one-entry. Line
 183 3 is important, because there is line 7 intersecting it in one-entry, which is not
 184 mapped.

185 In the same situation as above, consider a different partial mapping $f' =$
 186 $\{(1, 1), (2, 3), (3, 4), (5, 6)\}$, which is a mapping of the same level as f and only
 187 differs from f in mapping line 2. The line 2 is unimportant and by the definition
 188 of equivalent partial mappings, f and f' are equivalent. The idea behind this
 189 notion is simple. It is not important where we map line 2, because it does not
 190 restrict where we can map any other line that has not been mapped yet. This
 191 means that if a partial mapping f can be somehow extended, the equivalent
 192 partial mapping f' can be extended in the same way; therefore, it is sufficient to
 193 only extend one of them in order to find one full mapping. Note that it would
 194 be also sufficient to only extend one of the partial mappings if we were looking
 195 for all full mappings, but, in that case, we would need to keep the information
 196 about where the unimportant lines were mapped to.

3. An algorithm for testing pattern-avoidance of a special pattern

In the previous chapter, we have seen an algorithm for a general forbidden pattern. In this chapter, we introduce a special kind of a pattern, satisfying additional conditions, for which we can produce a much faster algorithm.

3.1 Walking pattern

Definition 10. A walk in a matrix P is a sequence of some of its entries, beginning in the top left corner and ending in the bottom right one. If an entry at the position $[i, j]$ is in the sequence, the next one is either $[i + 1, j]$ or $[i, j + 1]$. Let w denote the width of P and h denote its height, the length of an arbitrary walk is equal to $w + h - 1$ and we denote elements of the sequence by $w_1, w_2, \dots, w_{w+h-1}$.

Definition 11. We call a binary matrix P a walking pattern if there is a walk in P such that all the one-entries of P are contained on the walk.

$$M = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{pmatrix} \quad \begin{matrix} w_1 \\ w_2 & w_3 & w_4 \\ & w_5 & w_6 \\ & & w_7 \end{matrix}$$

Figure 3.1: An example of a walk W in matrix M and the order of entries in W .

In Figure 3.1 matrix M is a walking pattern as all the one-entries are included in a walk. We can also see that not all entries of a walk need to be one-entries.

It can be shown a walking pattern is exactly a matrix avoiding a forbidden pattern $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$.

3.2 Dynamic program

Next, we show an algorithm deciding whether a walking pattern P is contained in a big matrix M or not.

The pattern P is a walking pattern, so there is a walk containing all the one-entries of P . We choose one such walk arbitrarily. For each entry of the walk we remember whether its value in P is one or zero and whether the walk continues from the entry vertically, in which case we call it a *vertical entry* or horizontally, calling it a *horizontal entry*.

Definition 12. For an element e of M at the position $[i, j]$, the matrix $M_{\leq e}$ is a $(i + 1) \times (j + 1)$ submatrix of M consisting of rows with the index smaller than or equal to i and columns with the index smaller than or equal to j . The element e then lies in the bottom right corner. Similarly, $M_{\geq e}$ is a $(n - i) \times (n - j)$ submatrix of M consisting of rows with the index greater than or equal to i and columns with index greater than or equal to j . The element e is its first element.

To determine whether P is contained in M we find out for each element e of M what is the biggest index k such that there exists a mapping of $P_{\leq w_k}$ to $M_{\leq e}$. If there is an element for which we manage to find the whole pattern ($k = w + h - 1$), P is contained in M ; otherwise, it is avoided.

3.2.1 Inner structures

The algorithm uses two structures. For each w_k we remember whether it is a one-entry or zero-entry in P and whether it is a vertical entry or horizontal entry.

The second structure is a matrix of the same size as M . For each element e at the position $[i, j]$ we store two numbers. The number $c_v(e)$ is the biggest index k such that w_k is a vertical entry and there is a mapping of $P_{\leq w_k}$ to $M_{\leq e}$, in which w_k is being mapped to the j -th column. The number $c_h(e)$, symmetrically, is the biggest index k such that w_k is a horizontal entry and there is a mapping of $P_{\leq w_k}$ to $M_{\leq e}$, in which w_k is being mapped to the i -th row.

3.2.2 The algorithm

Definition 13. A diagonal of the matrix M is a subset of elements of M , such that all elements have the same sum of their coordinates.

For example, the zero diagonal only consists of the element $[0, 0]$, the first diagonal contains elements $[0, 1]$ and $[1, 0]$, and so on.




Figure 3.2: Diagonals of an matrix and the order in which the algorithm for walking pattern iterates through them.

The algorithm iterates through diagonals. For simplicity, in the pseudo-code below we do not deal with elements outside M (like $[-1, 0]$) explicitly. Instead, for those elements, we assume the values of c_v and c_h are always equal to zero and

250 $[-1, i]$ is a vertical entries and $[i, -1]$ is a horizontal for every $i \geq 0$. When we ask
 251 whether w_k can be mapped to e , where e is an element of M , we check whether
 252 w_k stands for a one-entry of P and if it does, we require e to be a one-entry too.

253 For an $n \times n$ matrix M the algorithm works as follows:

- 254 1. For $d = 0, \dots, 2n - 2$
- 255 2. For e element of d -th diagonal at the position $[i, j]$
- 256 3. $e_v := [i - 1, j]$
- 257 4. $e_h := [i, j - 1]$
- 258 5. $c_v(e) := c_v(e_v)$
- 259 6. $c_h(e) := c_h(e_h)$
- 260 7. If $w_{c_v(e_v)+1}$ can be mapped to e
- 261 8. If $c_v(e_v) + 1 = w + h - 1$
- 262 9. Terminate - M contains P as a submatrix
- 263 10. If $w_{c_v(e_v)+1}$ is a vertical entry
- 264 11. $c_v(e) := c_v(e_v) + 1$
- 265 12. Else
- 266 13. $c_h(e) := \max\{c_h(e), c_v(e_v) + 1\}$
- 267 14. If $w_{c_h(e_h)+1}$ can be mapped to e
- 268 15. If $c_h(e_h) + 1 = w + h - 1$
- 269 16. Terminate - M contains P as a submatrix
- 270 17. If $w_{c_h(e_h)+1}$ is a vertical entry
- 271 18. $c_v(e) := \max\{c_v(e), c_h(e_h) + 1\}$
- 272 19. Else
- 273 20. $c_h(e) := \max\{c_h(e), c_h(e_h) + 1\}$

274 3.2.3 Correctness

275 The first observation we make is that for every element e of M and any element
 276 e' above e in the same column $c_v(e') \leq c_v(e)$. This holds because whenever
 277 we manage to map $P_{\leq w_k}$ to $M_{\leq e'}$, then the same mapping maps $P_{\leq w_k}$ to $M_{\leq e}$.
 278 Similarly, it also holds for every e element of M and any element e' to the left of
 279 e in the same row that $c_h(e') \leq c_h(e)$.

280 The function can terminate before recomputing all elements and we have no
 281 guarantee about the state of elements, which have not been recomputed. If the
 282 function finds the pattern ending in entry e , it stops computing at that point, but

283 to prove correctness it is enough to prove the values are correct in $M_{\leq e}$, which
 284 has been fully recomputed. If, on the other hand, the function does not find the
 285 pattern, it recomputes the whole structure.

286 We need to show that the values of c_v and c_h are always correct for the
 287 recomputed elements at the end of the function. We proceed by induction on
 288 diagonals.

289 For the first diagonal it is definitely true since there can only be mapped w_1
 290 and we check that on lines 7 and 14.

291 When we are recomputing the values of $c_v(e)$ and $c_h(e)$ of an element e in the
 292 diagonal d , by induction hypothesis, all elements in diagonals $d' < d$ are correctly
 293 recomputed. Let cor denote the correct value of $c_v(e)$ as it is defined and com be
 294 the computed value. We need to show $cor = com$.

295 We can already see $cor \geq com$ because it holds after setting $c_v(e)$ on line 5
 296 and we only increase it, if we manage to find an extension of a mapping, in which
 297 case there really is a mapping; therefore, cor is greater or equal to the updated
 298 value.

299 To prove $cor \leq com$ we proceed by contradiction. Let us assume $cor > com$.
 300 It means there is a mapping of $P_{\leq w_{cor}}$ to $M_{\leq e}$ we have never found. Every such
 301 mapping has to map w_{cor} to e , because if it did not, the mapping would be
 302 possible even for diagonal $d - 1$, which is recomputed correctly and the value cor
 303 would be copied to com on line 5. Let us assume that w_{cor-1} is a vertical entry
 304 (else we proceed analogously). If $P_{\leq w_{cor}}$ can be mapped to $M_{\leq e}$ and w_{cor-1} is a
 305 vertical entry, then $P_{\leq w_{cor-1}}$ can be mapped to $M_{\leq e_v}$ and w_{cor-1} must be mapped
 306 to the same column as e . That means that $c_v(e_v) \geq cor - 1$. If $c_v(e_v) = cor - 1$
 307 and from knowing w_{cor} can be mapped to e , $com \geq c_v(e) \geq c_v(e_v) + 1 = cor$
 308 because of line 11. Otherwise $c_v(e_v) > cor - 1$, but then even from line 5 we get
 309 $com \geq cor$, resulting in contradiction.

310 To prove $c_h(e)$ has a correct value, we proceed symmetrically.

311 3.2.4 Generalization

312 The algorithm, with a few minor changes, can be also used for a pattern where all
 313 one-entries are contained on a walk from the top right corner to the bottom left
 314 one. The program supports both rotations of a walk and when walking pattern
 315 is chosen it automatically decides which variant to use.

316 On the other hand, a direct generalization for a general pattern does not work.
 317 While we can index all entries of the pattern, when trying to map a certain w_k
 318 to an element e of M , it is not sufficient to only check whether w_l is above and
 319 w'_l to the left from e .

320 In Figure 3.3, the entry of P in the square can be mapped to the element
 321 of M in the square and the same holds for entries in the circle but it is not a
 322 sufficient condition for the entry of P in the kite to be mapped to the element of
 323 M in the kite.

$$P = \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & \boxed{1} \\ 1 & \textcircled{1} & \diamond 1 \end{pmatrix} \quad M = \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & \boxed{1} \\ 0 & 1 & \textcircled{1} & \diamond 1 \end{pmatrix}$$

Figure 3.3: The algorithm testing avoidance for walking patterns cannot be easily generalized for all patterns.

324 4. Improvements to basic 325 algorithms

326 In this chapter we improve algorithms presented in previous chapters and intro-
327 duce a parallel method of testing pattern avoidance.

328 4.1 General pattern

329 We start by improving the brute force algorithm from Chapter 2.

330 4.1.1 Improving memory consumption

331 The algorithm creates all possible partial mappings and checks whether at least
332 one can be extended to a full mapping (mapping all lines of the pattern). To
333 compute all the partial mappings of some level l , it only uses mappings of level
334 $l-1$; therefore, it is enough to only store partial mappings of two levels in memory
335 at any time.

336 In Chapter 2 we also introduced the notion of (un)important lines and equiva-
337 lence based on not using unimportant lines at all (they are fully bounded by other
338 already mapped lines). When a line becomes unimportant, it stays unimportant
339 till the end of the test; as a result, we can forget where we mapped those lines to
340 save memory and only remember where we mapped important lines.

341 4.1.2 Not mapping empty lines

342 **Definition 14.** *An empty line is a row or a column that does not contain any*
343 *one-entries.*

344 An empty line can be mapped to any line and we do not need to map it at
345 all, as long as the algorithm does not map two lines surrounding an empty one
346 to two consecutive lines.

347 4.1.3 Using the last changed position

348 The MCMC process always changes one element of the big matrix and asks
349 whether it still avoids the pattern. If it does not and we know that before the
350 change it did, we are sure the changed element $[r, c]$ is a part of the pattern. It
351 is hard to use this fact in the algorithm. It just maps one line after another and
352 we do not know at the beginning to which line the changed position lines should
353 be mapped.

354 What we can do is to enforce that neither the r -th line nor the $n + c$ -th
355 one (c -th column) get skipped. We only look at the restriction for rows as the
356 restrictions for columns are symmetrical. There are three situations we want to
357 avoid:

- 358 • The first row of P is mapped under the r -th row. This prevents any other
359 row to be mapped to r -th one and we don't want that.

- The last row of P is mapped above the r -th row. This again prevents any other row to be mapped to r -th one.
- Two adjacent rows $l, l + 1$ of P are mapped to $L < L'$ respectively and $L < r < L'$ which leaves no other row to be mapped to r .

4.1.4 Line order

An important thing, if we want the algorithm to run fast, is to choose a good line order. A line which is unimportant in level l in a line order may easily be important till the nearly last level in a different order.

We choose line order to hopefully enforce two things:

- Make as many unimportant lines as possible. This really allows the equivalence based improvements to kick in. The more lines are unimportant the more mappings become equivalent and the faster it is to iterate through all of them.
- Recognize hopeless partial mappings as soon as possible. A partial mapping gets extended if the line does not break the rule that there is a one-entry where it needs to be. If we map all the rows first, the rule will get broken only after we start to map columns and we probably want to find out sooner.

In the program a user can either choose their own custom order or one of four algorithms with different main purposes:

- AUTO - this one tries the other three line orders and chooses the one which shows the best performance over some iterations on a matrix. While this may sound like a good thing to use, it is only so if an initial matrix is chosen and it takes a lot of time since a lot of iterations need to be made in order to make a good sample. I would recommend not to use AUTO order at all and instead to try all the line orders by hand with a number of iterations depending on the pattern and a good initial matrix; for instance, generated with a smaller number of iterations on the same pattern and with any line order.
- DESC - the lines are ordered in descending order depending on the number of one-entries. This follows the idea to start with the lines that are the hardest to map. Note that this algorithm does poorly if there are a lot of lines with the same number of one-entries (for example an identity matrix).
- MAX - it orders the lines so that the maximum number of important lines throughout the levels is as small as possible. This focuses straightforwardly to having many unimportant lines, which the program does not remember.
- SUM - it orders the lines so that the sum of the numbers of the important lines is the smallest possible throughout all levels. The purpose is the same as in the MAX order and quite often it is the case both approaches produce the same order.

400 • TWO - it orders the lines so that the maximum number of important lines
 401 in two consecutive levels throughout all the levels is as small as possible.
 402 This again focuses to having many unimportant lines, which the program
 403 does not remember. The constant two is chosen due to the fact general
 404 pattern always stores two levels of partial mapping at a time.

405 4.1.5 Mapping approaches

406 The one thing the approaches we will introduce have in common is that they try
 407 to recognize those partial mappings that have no chance to be extended to a full
 408 mapping as early as possible.

409 While the algorithm introduced in Chapter 2 finds out the partial mapping is
 410 invalid only at the time it maps two lines having a one-entry at their intersection
 411 to two lines having a zero-entry at the intersection, different approaches try to
 412 reveal the fact we would end up in the situation earlier by checking more condi-
 tions. Let P from Figure 4.1 be the forbidden pattern and imagine a situation,

$$P = \begin{matrix} & \begin{matrix} 5 & 6 & 7 & 8 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \end{matrix} & \begin{pmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix} \end{matrix}$$

Figure 4.1: Pattern P on which we demonstrate mapping approaches.

413 in which only lines 0, 3 and 7 are mapped and line 6 is currently being mapped.
 414 There are a few necessary conditions we can check:
 415

416 Enough one-entries

417 The first condition is that there is enough one-entries in between mapped lines,
 418 which is schematically shown in Figure 4.2. We check whether there is enough
 419 one-entries on lines in between those lines, where lines 0 and 3 are mapped, so
 420 that there is a hope we can map lines 1 and 2 there. Similarly, we check whether
 421 there is a one-entry below the line, where line 3 is mapped so we can map line 4
 422 there later.

423 Recursive mapping

424 While we were only testing whether there are enough one-entries in between al-
 425 ready mapped lines in the previous approach, as you can see in Figure ??(recursive),
 426 this time we also check whether those one-entries can be used for the lines that are
 427 intended to be mapped there. For example, when we check there is a one-entry
 428 to be used for line 1 later, we also check the line 1 can be mapped to the row




Figure 4.2: Checking whether there is enough one-entries. Bold lines are mapped or being mapped and in circles are the positions where we look for one-entries.




Figure 4.3: Checking whether crossed non-mapped lines can be mapped anywhere. Bold lines are mapped or being mapped, in rectangles are the line we check and in circles are the positions where we look for one-entries.

429 with one-entry, which in this situation means to also check there is a one-entry
 430 at the intersection with the line to which the line 7 is mapped.

431 **Orthogonal bounds**

432 As shown in Figure 4.4, when we are adding line 6, we check whether there is
 433 enough one-entries on the already mapped lines orthogonal to line 6 between line
 434 6 and the closest mapped lines next to line 6. The idea is same as in “Enough
 435 one-entries”, but we check different lines.




Figure 4.4: Checking whether there is enough one-entries on the orthogonal lines. Bold lines are mapped or being mapped and in circles are the positions where we look for one-entries.

436 Usage

437 These restrictions on the added lines are not a fixed part of the program. A user
 438 can decide which approaches they want to use in the configuration file.

When testing that was done for a fixed pattern, we found out it is useful to use all the mentioned restrictions when generating a matrix of size 100×100 , as it turned out to be much faster than without the restrictions. On the other hand, in the same test for a generated matrix of size 500×500 , it was much better not to use any of those restrictions.

++

TODO add chart and change text above

++

439 4.1.6 Using the whole structure in the next iteration

440 It may seem like a good idea to store all the partial mappings. In the next
 441 iteration, instead of finding all the partial mapping again, we only alter the
 442 mappings we remember. Let i be the number of the iteration we are in, and e be
 443 the element.

444 If the element e is changed from zero-entry to one-entry, for each partial
 445 mapping we have stored in previous iterations, we want to try to extend it only
 446 by the line that just changed. If we manage to extend a partial mapping, we
 447 then try to extend it to a full mapping in all possible ways (not only by using
 448 changed lines). When the new line in such a mapping becomes unimportant, we
 449 can stop looking for all possible extensions if the mapping is equivalent with a
 450 different one, which comes from previous iterations. This can be easily done by
 451 means already used in the standard algorithm.

452 However, if the element e gets changed from one-entry to zero-entry we need to
 453 go through the partial mappings and delete all those that use e . This complicates

the algorithm as we can no longer forget unimportant lines. Moreover, for each partial mapping we need to remember how many partial mappings of the previous level can be extended to that one, to delete that mapping from the list if there are no longer any mappings extensible to it.

This can all be done, but it comes with three huge inconveniences:

- Memory consumption - there can be a lot of partial mappings and we need to remember them all. We need to remember mappings of all levels and while we can still use the equivalence when extending a mapping, we need to also store all equivalent mappings for the purposes of deleting.
- The change from one-entry to zero-entry is no longer for free. If this change is done, we already know the pattern is not contained in M , but we still need to do a lot of work to change the structure in order to use it in the next iteration.
- Reverting - if the change is unsuccessful (the pattern is contained) we need to revert the change which means to completely revert all changes we did to the list of partial mappings. This can be either done by making a backup copy of the whole structure and override the structure if needed, which again is very costly as the structure is huge, or we can remember what partial mappings are new (or deleted) and we go through all partial mappings and remove (add) those. This means to iterate through the big structure one more time for every unsuccessful change.

After realizing these issues it no longer looks useful to me and this version of the algorithm is not a part of the implementation.

4.2 MCMC parallelism

To speed up computations, it is often possible to use parallelism. In this section, we show how to make the MCMC generator parallel, while still allowing both types of the pattern.

While the serial MCMC generator in each iteration changes one element in the generated matrix and checks whether it still avoids forbidden patterns, the parallel version makes several iterations at once, one on each copy of the generated matrix. This means that while iteration x is being computed by a thread, iteration $x + 1$ can at the same time be computed by a different thread. The issue is that the iteration $x + 1$ does not know what is going to be the state of the generated matrix at the time it should start. It expects iteration x to fail - not change the generated matrix at all, counting on the fact, it is unlikely a change does not create a mapping of the pattern, and starts with the same matrix as iteration x . If iteration x succeeds, then the computed iteration $x + 1$ is invalid and the iteration is going to be recomputed again, starting with the correct matrix.

When the parallel version of MCMC generator is chosen and it is assigned n threads, it creates $n - 1$ private copies of the generated matrix and assigns one thread, called worker, to each of them. The last thread, which we call the main thread and which has exclusive access to the master copy of the generated matrix, makes one change of a bit in each private copy of the matrix and makes the corresponding worker check the avoidance.

498 The job of a worker is only to check if its copy of the matrix still avoids the
 499 pattern when one bit is changed. On the other hand, all synchronization is left to
 500 the main thread. As mentioned before, one iteration of the MCMC process can
 501 be recomputed several times. We still want the generator to satisfy the conditions
 502 we have for the Markov chain (more in Section 1.2) in order to approximate a
 503 random matrix. To achieve that, if a computed iteration x succeeds (and changes
 504 the generated matrix), all the other computed iterations that would follow after
 505 the iteration x become invalid and they all have to be recomputed. The process
 506 ends when all iterations get computed.

507 For the purposes of clarity, from now on, we won't be talking about iterations
 508 but about tasks. A task is basically one iteration of the MCMC process. The
 509 usefulness of this notation comes with an ID - a number, unique per task, assigned
 510 to each task, starting with 1 and always increasing. For a pair of consecutive
 511 iterations x and $x + 1$ it will always be the case that if task a is the last task
 512 to compute iteration x (which means the iteration does not get recomputed ever
 513 again after) and task b is the last task to compute iteration $x + 1$, then the ID
 514 of a is lower then the ID of b . Also there is no point, in which two different
 515 tasks would be computing the same iteration at the same time. If tasks with IDs
 516 $a < b$ computed the same iteration, it must have been the case an earlier iteration
 517 succeeded when task with ID a was computed and after it got removed, task with
 518 ID b was assigned to recompute.

519 At any point in time, we only consider those tasks, that are being computed
 520 or those that wait to be processed (not those that have been processed), which
 521 means the lowest ID of tasks we consider increases in time.

522 When a task ends and it has the lowest ID (we can always wait for the task
 523 with the lowest ID) we do:

- 524 • if it fails:
 - 525 – Do nothing - there is no change to propagate to the master copy of the
 - 526 generated matrix and all the tasks with higher ID expected this task
 - 527 to fail, which it did.
 - 528 – This increases the lowest ID by exactly one, as the task we speak of
 - 529 got processed.
- 530 • if it succeeds:
 - 531 – The main thread propagates the change tested by the task to the
 - 532 master copy of the generated matrix.
 - 533 – All the rest of the task get removed as they all had a higher ID -
 - 534 computed iterations that follow after the one just computed and they
 - 535 expected the task to fail, which it did not.
 - 536 – This increases the lowest ID by more then one, because there are tasks
 - 537 that got removed and one that got processed.

538 4.2.1 Example of the MCMC process for n threads

539 At first, iterations 1 to $n - 1$ are assigned one to each worker as tasks with ID 1 to
 540 $n - 1$ with the same order as the order of iterations. If iteration 1 is not successful

(which all the other iterations count on), everything is alright. However, if the iteration (its task) is successful, all the results of other tasks (and some of them might have been already finished) are cleared and those iterations get recomputed in tasks n to $2n - 3$ and the worker that computed task with ID 1 is assigned a new task with ID $2n - 2$ - to compute iteration n . The result of the task gets propagated to the master copy of the generated matrix only if all the tasks n to $2n - 3$ fail, else is gets recomputed. This is what happens till the end.

4.2.2 Speculative computing

It may easily happen that a task not having the lowest ID ends first. In that case, we could just wait until it has the lowest ID and process it later. This is not a very efficient approach. Instead we process the task immediately, but we don't propagate the changes to the master copy of the generated matrix until all tasks with lower ID fail and we do not stop the workers processing tasks with lower ID. When a task succeeds we remove all the changes computed by tasks with higher ID and override their private copy of the generated matrix. Also it might happen a task with even lower ID succeeds as well. This leads to more and more overriding. Luckily this is the only precarious situation we may encounter and it can be dealt with, even without copying the possibly huge generated matrix.

The way we deal with these inconveniences is described in Chapter 5 and should be clear from the code itself.

4.2.3 Reverting and synchronizing in the main thread

The speculative computing discussed above is not the only improvement we can make. It turns out to be costly to wake a thread to compute a trivial function, to set a few atomic variables and to fall asleep again. This happens a lot in the MCMC process. Every time a task succeeds it makes other workers revert the changes they computed and synchronize the successful change, which are both trivial functions.

To workaround this problem we make a theoretically bad decision, which comes with very nice practical results. All the reverts and synchronizations are computed by the main thread instead of by an appropriate worker. There is no problem with concurrency because the worker is always asleep when a task is to be assigned and using the fact those tasks are really trivial, it does not make the rest of threads wait for the main thread for too long while it computes changes.

4.3 Walking pattern

While the brute force implementation of an avoid algorithm for a general pattern was improved heavily, the algorithm for a walking pattern (see Chapter 3) is very fast in its nature and cannot be improved. Or can it be?

4.3.1 Using the last changed position

The MCMC process always changes one element e of the big matrix and asks whether it still avoids the pattern. If it does not and we know that before the

581 change it did, we are sure e is a part of the pattern (a one-entry of the pattern
582 is mapped to it). Knowing that and using the same inductive proof as we did in
583 the proof of correctness of the avoid algorithm (see Chapter 2) it is sufficient to
584 only recompute the part of the inner structure under e and check if the last entry
585 of the pattern can be found there.

586 Not only that. We also know, using the fact the structure was completely
587 correct before the change, that if the values of both c_v and c_h of an element did
588 not change, the element won't cause the element underneath it to change and we
589 no longer have to recompute other parts of the structure.

590 To use both these facts we replace the cycle through the diagonals by a sim-
591 ple queue, starting at the position of the last changed element and putting more
592 positions in if the values of c_v or c_h are different than they were before recomput-
593 ing. The function ends either when the pattern is discovered or when the queue
594 becomes empty.

595 4.3.2 Lazy avoid

596 Lazy avoid is a variant of avoid function used when the MCMC parallelism (more
597 in Section 4.2) is chosen. While all the other types of patterns have a trivial
598 implementation of revert function, when using the walking pattern, the inner
599 structure needs to be modified even when reverting. The MCMC parallelism
600 turned out to work much better if the revert calls are handled by the main thread
601 and it requires the function to run as fast as possible so the other threads are not
602 blocked by the call for too long. That is a reason why functions lazy revert and
603 lazy avoid were created.

604 The avoid function expects the inner structure of the walking pattern to be
605 in a valid state and that requires some effort. To make lazy revert the fastest
606 possible, we postpone the work until the next call of lazy avoid, meaning that
607 lazy avoid then needs to do more things at once. It is no longer sufficient to only
608 compute the submatrix under the position changed last as we did above, but it
609 needs to also compute changes in the positions changed in those lazy revert calls
610 that are postponed.

611 We discuss several approaches, starting with the simplest one and ending with
612 the one that is fast and used in the final implementation.

613 Recompute the whole structure every time

614 The easiest way to implement lazy avoid is to always recompute the whole inner
615 structure. In that case, we do not worry which positions are correct and which
616 are not, because every time we find the pattern, we recomputed all the entries
617 that form it, so we know it really is there.

618 The weakness is efficiency. If the whole structure was correct and there was a
619 change of the last entry of the matrix it is sufficient to only recompute that one
620 entry. Instead we recompute a possibly very big structure. This results in a very
621 bad performance negating the advantage of parallel computation.

622 **Recompute only a part of the structure diagonal by diagonal**

623 A simple improvement is to remember the changes done in previous calls of lazy
624 revert and together with the change done in lazy avoid call only recompute the
625 part of the structure that has possibly changed.

626 This gets more complicated when lazy avoid call discovers the pattern in $M_{\leq e}$,
627 because we cannot be sure the rest of the structure (everything under the diagonal,
628 where e is present) is in a correct order. It is still possible to remember some
629 horizontal, vertical and diagonal bounds and use them to restrict the recomputed
630 part of the matrix. However, the improvement is not that significant and we can
631 do better.

632 **Queue of positions to recompute**

633 A different approach is closer to the one used in a standard avoid function. Instead
634 of going through diagonals one after another, we have a queue of entries-to-
635 recompute. It is no longer sufficient to have a standard queue since in different
636 calls of lazy revert/avoid we can possibly change an entry of different priority
637 (the smaller diagonal the more important) so we need to have some kind of a
638 priority queue. That is exactly what I tried.

639 Using `std::priority_queue`, the function has no more problems with recomput-
640 ing the entries that were not influenced by the changes and uses all the benefits
641 mentioned in the previous section. But the container does not come for free and
642 in the end it turns out the price we pay for the operations on the priority queue
643 make the whole implementation comparably slow as in the previous attempts.

644 **Two leveled queue of positions to recompute**

The final solution comes with the same idea, but a different storage. As the prior-
ity depends upon a diagonal (two entries on the same diagonal can be recomputed
in any order) we only remember a priority queue of diagonals and an array of
diagonals saying whether a diagonal is already a member of the priority queue.
As far as the entries are concerned, for every diagonal we have a `std::vector` of
entries-to-recompute as well as an array saying whether an entry is already a
member of the vector. Finally, it is the case that the storage used is not only
good theoretically but as the numbers say, also practically.

++

[reference to a table of measurements or something]

++

645 5. Technical documentation

646 In this chapter, we cover those parts of the algorithm that may be hard to un-
647 derstand just from the code. This only means functions that are technically
648 hard, for example functions with unexpected dependencies, side effects and so
649 on. Algorithmic difficult tasks are explained in Chapter 4.

650 5.1 Classes and API

651 First we list important classes of the program and explain their purpose.

652 5.1.1 Matrix

653 A minimalistic template container for storing and accessing matrices.

654 5.1.2 Pattern

655 An abstract class defining the interface of patterns. Three classes inherit from
656 the class:

- 657 • *General_pattern* – more in Chapter 2
- 658 • *Walking_pattern* – more in Chapter 3
- 659 • *Slow_pattern* – a class using a brute force algorithm to test pattern avoid-
660 ance.

661 For the case multiple patterns are chosen at the same time, every such Pattern
662 is stored in a container called Patterns, which is a class creating an interface
663 between Pattern and MCMCgenerator.

664 5.1.3 Statistics

665 To acquire, store and output statistics of the generating process, we use classes
666 in the file Statistics.hpp. There are two kinds of classes:

- 667 • *Matrix statistics* – these statistics store information about the structure of
668 all matrices that have been generated throughout the MCMC process. An
669 example of that is a histogram of occurrences of one-entries at all positions
670 of the matrix.
- 671 • *Performance statistics* – this is what we use to count how many changes
672 were successful and how long did it take to test a change.

673 5.2 General_pattern

674 The general pattern class contains a lot of function. Most of them are easy to
675 follow and they all should be commented enough in the code. The only part
676 which deserves more attention is the constructor.

677 5.2.1 Construction

678 In the constructor of a general pattern, there are a few function that are easy in
679 nature but as they somehow use each other it is hard not to lose track of their
680 dependencies and results. In order to make this part of the code, which is very
681 important, more understandable, we go through the constructor and explain all
682 that is happening in the order it is happening in.

683 Storing the pattern

684 The first thing, which is done right after initialization of variables, is storing the
685 pattern. Instead of storing the pattern in a `Matrix<bool>`, I decided to store
686 lines into a number, where in the binary coding a one-entry in the position i
687 means there is a one-entry in the line at the intersection with i -th orthogonal
688 line. This comes handy when computing line orders. At the same time we also
689 find those lines that are empty (more in Chapter 4) and remember them, because
690 we do not have to map them at all.

691 Choosing the line order

692 After that, we need to choose the right line order (again more in Chapter 4).
693 To compute MAX, SUM or TWO order we use a brute force algorithm that
694 checks sequences of line adding and for each it computes how many lines are
695 unimportant. Then it just chooses the order which is the best in chosen metric.
696 To compute DESC order, we sort the lines according to the number of one-entries.

697 What to remember

698 In the next step, we find which lines are important in each level of partial map-
699 pings with respect to chosen order, because what to remember is based on the
700 equivalence introduced in Chapter 2 and the decision not to remember unimpor-
701 tant lines, which we explained in Chapter 4.

702 Parallel bound indices

703 Now comes the hardest to follow part – precomputing the indices for searching
704 for parallel bounds. The idea is simple. When we are adding a new line and we
705 already have a partial mapping, it restricts to where we can add the line. For
706 example, if there are three rows in the pattern and the rows 1 and 3 are mapped,
707 then line 2 needs to be mapped in between those two. The question is, where are
708 those two lines mapped to?

709 First, we add in a chosen order and second we do not remember all lines, as
710 some are unimportant. What do we want is to have an instant access to indices
711 of lines, which bounds added line in the partial mapping, so we do not need to
712 compute the index over and over again. That is exactly what gets computed when
713 the function “`find_parallel_bound_indices`” is called. The series of other function
714 calls follows just because we compute the indices for all added lines in the order
715 in which they are going to be added.

716 Extending order

717 The last function, “find_extending_order” specifies how we store an extended
718 partial mapping. Again, unimportant lines play their role here and it may easily
719 be the case from a partial mapping storing k lines, after mapping one more
720 line, we end up with a partial mapping only storing $k - 1$ lines, because two
721 lines become unimportant by adding the line. This means we not only copy the
722 previous mapping and add the new mapped line but also remove unimportant
723 lines. This function precomputes which values are going to be copied and which
724 are going to be skipped.

725 5.3 MCMC parallelism

726 While the idea behind MCMC parallelism is described in Section 4.2 and the
727 code is heavily commented, the work done by the main thread may still be hard
728 to understand.

729 Let I be the ID the process is currently waiting for, that is, the lowest ID of
730 a task that is being tested by a worker. In a structure called “queue” (which is
731 `std::vector<std::deque>`) each worker has a queue of tasks assigned to it. In the
732 queue, there are tasks that are either being computed or have been computed.
733 The history of tasks is needed to allow reverting changes that should have not
734 happen when the main thread encounters a different successful task with lower
735 ID. There is no need to have a complete history of all tasks computed. There
736 are only those tasks, that have higher ID than I or have lower ID, but those are
737 going to be removed from the “queue” as soon as possible. The name “queue” is
738 not random, it describes the order, in which the tasks are being stored – the tasks
739 with lower ID have been inserted earlier and therefore they are at the bottom.

740 Now that we know the most important structure let’s see how the main thread
741 works with it. This is a list of operations changing “queue” and the situations,
742 in which we perform them:

- 743 • `pop_front`: The main thread deletes the first task (the one with the lowest
744 ID) if one of two things happen:
 - 745 – The ID of the task being deleted is equal to I . That means the change
746 computed by the task is being propagated to the generated matrix and
747 there is no need to remember the task anymore. This also increases I ,
748 not necessarily by one.
 - 749 – The ID of the task being deleted is less than I . This situation happens
750 due to synchronization. The worker was supposed to synchronize a
751 task computed by a different worker that did not have the lowest ID
752 at the time. Therefore, the task needs to be in the list of tasks so we
753 can revert it later, if needed. If there is no need to revert it and the
754 lowest ID gets greater or equal to the ID of the task, we can just delete
755 it from the “queue”.
- 756 • `pop_back`: There is only one reason to delete tasks from the end of the
757 “queue” and that is reverting. Imagine there is a task with id J at the end
758 of the “queue”. A different worker computes a task with lower ID and finds

759 out the change is successful. This means the task J won't propagate to the
760 generated matrix and there is no use for it. If it is still being computed, we
761 cannot do much about it, so we tell the worker to stop computing and deal
762 with it later. If the task is finished, we need to revert it, but only in case
763 the task was successful, because if it was not, it had already been reverted
764 by the worker. So we revert the task if needed and we can just delete it
765 from "queue" as it will never be used.

- 766 • `emplace_back`: The main thread only inserts new tasks to the end of the
767 "queue" and there are two reasons to insert:
 - 768 – Worker is assigned a completely new task to check the avoidance. In
769 this situation, the task is given a new, globally highest ID and we add
770 the task at the end of the "queue".
 - 771 – The second reason to insert into "queue" are synchronizations. The
772 situation is the same as it was in the case, when we `pop_back` – after
773 we revert all the tasks in the list, we need to synchronize changes that
774 forced reverting and if their ID is not lower or equal to I , we need to
775 add them to the list so they can be reverted if needed.

776 6. User documentation

777 In the last chapter of the thesis, we first describe how to install the program and
778 then show how to make the program generate random matrices or to test whether
779 a certain matrix avoids a given forbidden pattern.

780 6.1 Installation

781 The program is written in C++ and should be compilable on any standard plat-
782 form. To use it, you either just use an executable file (Windows) or build the
783 program using C++ compiler.

784 6.1.1 Windows

785 Windows users can run the application using an executable file “matrix-win.exe”
786 either directly, in which case the default configuration file will be used, or using
787 command line with an optional parameter specifying the configuration file.

788 6.1.2 Unix, Linux, iOS

789 Users on other platforms than Windows can build the solution using command
790 line easily by running “./build.sh”, which uses G++ compiler. Compiler can be
791 switched by rewriting g++ to some other variant (for example clang) in build.sh
792 file. This leads to creating a an executable file “matrix.exe” which can be run
793 with an optional parameter specifying the configuration file.

794 6.2 Configuration file

795 In order to modify what the program computes, we use a configuration file. The
796 configuration file can be chosen when running the program in command line and
797 relative path to it is the first (and only) option. If no path is inserted then the
798 configuration file is expected to be located in the same directory as the executable
799 file and its name is “config.txt”.

800 The file is a standard text file which can be modified by any text editor and
801 is structured into four sections:

- 802 • input
- 803 • pattern
- 804 • output
- 805 • statistics

806 The order of the sections is not fixed and there can be additional empty lines for
807 better readability. In each section, there is a list of values that can be set either
808 to arbitrary value or to a specific one. There is at most one command of format

809 “option=value” per line and there might be additional white spaces surrounding
810 the “=” sign.

811 If an option is set more than once, the latter value is always used. If, on the
812 other hand, an option is not set at all, the default value is used. If there is a
813 line encountered that sets a wrong option, for instance when the user mistypes a
814 valid option, the line is skipped and the user gets a warning in the standard error
815 output.

816 Let us provide a list of all options for each section together with their default
817 values.

818 6.2.1 Input

819 In the first section of the configuration file, we set the generating process.

- 820 • size: The size of the generated matrix. Results in $M \in \{0, 1\}^{size \times size}$.

Possible value: $s \in \mathbb{N}$
821 Default value: 100

- 822 • iterations: The number of iterations of the MCMC process.

Possible value: $i \in \mathbb{N}$
823 -1 - tests avoidance of the initial pattern
Default value: 10,000

- 824 • random_seed: The random seed for the MCMC process.

Possible value: $s \in \mathbb{N}$
825 “random” - chooses a random seed
Default value: “random”

- 826 • init_matrix: A $size \times size$ matrix the MCMC process starts with.

Possible value: *matrix file path*
827 “zero” - a matrix containing no one-entries
Default value: “zero”

- 828 • parallel_mode: A choice to compute in parallel or serial.

Possible value: “serial”
829 “mcmc” - more in Section 4.2
Default value: “serial”

- 830 • threads_count: The number of threads if a parallel mode is chosen.

Possible value: $t \in \mathbb{N}$
831 -1 - chosen according to the number of cores
Default value: 1

832 6.2.2 Pattern

833 In this section we set the options that matter the most – matrix patterns. As
834 we are allowed to generate a matrix which avoids more than just one pattern,
835 the section [pattern] can be used multiple times, specifying one pattern for each
836 occurrence.

837 • `pattern_file`: A relative path to a input matrix file - the pattern.
 Possible value: *matrix file path*
 838 Default value: `"pattern/input.txt"`

839 • `pattern_type`: The type of the pattern. Determines the method used for
 840 testing avoidance.
 Possible value: `"general"`
 `"walking"` - see Chapter 3
 841 `"slow"` - brute force algorithm for a general pattern
 Default value: `"general"`

842 The next options are only useful if the general pattern type is chosen. It
 843 specifies how the mappings are stored as well as what the map function tests.
 844 First we can decide what mapping approaches to use. More about them in
 845 Section 4.1.5.

846 • `map_one_entries`: If set to `"yes"`, the map function tests whether there is
 847 enough one-entries in between already mapped lines.
 Possible value: `"yes"`
 848 `"no"`
 Default value: `"yes"`

849 • `map_recursion`: If set to `"yes"` and the `map_one_entries` is also set to `"yes"`,
 850 the map function tests mapping recursively.
 Possible value: `"yes"`
 851 `"no"`
 Default value: `"yes"`

852 • `map_orthogonal_bounds`: If set to `"yes"`, the map function also tests the
 853 orthogonal bounds of added line.
 Possible value: `"yes"`
 854 `"no"`
 Default value: `"no"`

855 • `map_container`: A container in which the partial mappings are stored.
 Possible value: `"set"` - `std::set` (red-black tree)
 `"hash"` - `std::unordered_set` (hash table)
 856 `"vector"` - `std::vector` (dynamic array)
 Default value: `"hash"`

857 • `line_order`: Choose the order in which the lines are being added to the
 858 partial mapping. See Section 4.1.4.
 Possible value: `"max"`
 `"two"`
 `"sum"`
 859 `"desc"`
 `"auto"`
 `"order file path"`
 Default value: `"max"`

860 6.2.3 Output

861 In this section we specify, where to output the generated matrix or statistics files.
862 As the matrix can be output to console, a text file or a bmp file, each option in
863 the section can be set more than once and every line will make a new output.

- 864 • `matrix_output`: The generated matrix can be output as a bmp file in which
865 one-entries are black pixels and zero-entries white. To do that, the file path
866 has to have a pattern “*.bmp”. If a different path is given the file is stored
867 as a matrix text file. It can also be output into a console if “console” is set.
868 In that case it has the text format.

Possible value: “console”
matrix bmp file path
matrix text file path
“no”

Default value: “no”

- 870 • `performance_stats`: If the serial computation is chosen, the program can
871 output a statistics like the percentage of avoid call success, how long did
872 one call take on average and what was the average size of structures. If more
873 patterns are chosen at the same time, the statistics may get misleading as
874 they also count the cases when the first pattern is contained in the matrix
875 and the other patterns are not tested at all.

Possible value: “console”
performance file path
“no”

Default value: “no”

- 877 • `performance_csv_stats`: The same information as above but formatted to a
878 csv file so the data can be more easily worked with.

Possible value: “console”
csv file path
“no”

Default value: “no”

- 880 • `time_to_console`: Prints how long the computation took into a console.

Possible value: “yes”
“no”

Default value: “no”

- 882 • `patterns_to_console`: Prints all the used patterns into the console.

Possible value: “yes”
“no”

Default value: “no”

884 6.2.4 Statistics

885 The last section handles the options important for scientists. While generating
886 a random matrix is a great result, on its way the program can also create some
887 statistics, namely make a histogram of occurrences of one-entries in a generated

888 matrix as the MCMC iterates as well as store the matrix with the highest amount
 889 of one-entries. As the process usually does not start with a random matrix, the
 890 user can decide to only compute the statistics after a certain number of iterations
 891 has been done and to only check a small portion of iterations, every 10th for
 892 instance, as a single iteration may not make any difference and counting the
 893 histogram takes time.

- 894 • histogram_frequency: Sets how often the histogram gets refreshed.
 Possible value: $f \in \mathbb{N}$
 0 - the histogram is not computed at all
 Default value: 0
- 896 • histogram_initial: Sets the initial iteration of the MCMC process when the
 897 histogram gets refreshed.
 Possible value: $i \in \mathbb{N}$
 898 Default value: 1,000
- 899 • histogram_final: Sets the last iteration of the MCMC process when the
 900 histogram gets refreshed.
 Possible value: $f \in \mathbb{N}$
 -1 - the histogram is computed till the end
 901 Default value: -1
- 902 • histogram_file: Sets where to output the histogram computed during the
 903 MCMC process.
 Possible value: matrix bmp file path
 matrix text file path
 "console"
 "no"
 904 Default value: "no"
- 905 • max_ones_matrix_file: Sets where to output the matrix that had the most
 906 one-entries among all matrices iterated through during the MCMC process.
 Possible value: matrix bmp file path
 matrix text file path
 "console"
 "no"
 907 Default value: "no"

908 6.3 File input

909 There are only two types of input files expected by the program. Either you want
 910 to read a matrix file, which can be a pattern or an initial matrix, or an order file
 911 that determines an order in which the lines are going to be mapped if the general
 912 pattern is chosen.

913 6.3.1 Matrix file

914 A matrix file is a standard text file having the format as follows:

- 915 • 2 natural numbers specifying the number of rows and columns in this order.
- 916 • a sequence of zeros and ones of length rows×columns specifying the matrix
- 917 from the top left corner one row after another.

918 **Example:** 2 3
 1 0 1
 1 1 0

919 6.3.2 Order file

920 If you want to choose the order in which the lines are going to be mapped when
921 a general pattern is chosen, it is your responsibility to check that all lines that
922 need to be mapped are mapped. It is for example possible to only map three lines
923 even if the pattern consists of six lines just because there is for example no need
924 to map empty lines at all. Therefore the program does not check the validity of
925 the order and just uses it.

926 Now that the user has been warned, the format of the custom order file is
927 simple. It consist of the indices of the lines of the pattern numbered starting
928 with 0 and starting from the top row and ending with the right column.

929 One possible order for the matrix given as an example in [6.2.1] is this file:

930 2 1 0 3 4

931 First mapping the left column, the second and first row after that and finishing
932 the mapping with the middle column and the right one.

933 6.4 File output

934 Let us now find out what the output files look like.

935 6.4.1 Matrix text file

936 The matrix text file has the same format as the input one. It consists of:

- 937 • 2 natural numbers specifying the number of rows and columns in this order.
- 938 • a sequence of zeros and ones of length rows×columns specifying the matrix
- 939 from the top left corner one row after another.

940 The matrix is binary except for the one produced as a histogram, which can have
941 higher natural numbers and contains the number of samples as the last number.

942 If you then divide all the entries by the last number, you get a percentage of the
943 entry being a one-entry.

944 **6.4.2 Matrix bmp file**

945 For an $n \times n$ matrix the standard bmp file contains $n \times n$ pixel of black color
946 meaning a one-entry and a white color for a zero-entry. If the histogram is output
947 as a bmp file, the pixels are greyscaled and the darker a pixel is the more often
948 the entry was a one-entry during the MCMC process.

949 Conclusion

950 Bibliography

951 Neal Madras. *Lectures on monte carlo methods*, volume 16. Springer Science &
952 Business, 2002.

953 Neal Madras and Hailong Liu. Random pattern-avoiding permutations. *Algo-*
954 *rithmic Probability and Combinatorics, AMS, Providence, RI*, pages 178–180,
955 2010.

List of Figures

957	1	Matrix M_1 contains the pattern P , because a mapping $\{(0, 0), (1, 2), (2, 3), (3, 4)\}$	
958		satisfies all the conditions. On the other hand, matrix M_2 avoids	
959		P as there is no such mapping.	4
960	2	Example of a generated matrix avoiding I_{10} (unit matrix). Black	
961		dots are one-entries and white are zero-entries. As you can see,	
962		matrices avoiding a pattern can have a nice structure.	5
963	2.1	An example showing unimportant line and equivalent mappings. .	9
964	3.1	An example of a walk W in matrix M and the order of entries in W .	10
965	3.2	Diagonals of an matrix and the order in which the algorithm for	
966		walking pattern iterates through them.	11
967	3.3	The algorithm testing avoidance for walking patterns cannot be	
968		easily generalized for all patterns.	14
969	4.1	Pattern P on which we demonstrate mapping approaches.	17
970	4.2	Checking whether there is enough one-entries. Bold lines are mapped	
971		or being mapped and in circles are the positions where we look for	
972		one-entries.	18
973	4.3	Checking whether crossed non-mapped lines can be mapped any-	
974		where. Bold lines are mapped or being mapped, in rectangles are	
975		the line we check and in circles are the positions where we look for	
976		one-entries.	18
977	4.4	Checking whether there is enough one-entries on the orthogonal	
978		lines. Bold lines are mapped or being mapped and in circles are	
979		the positions where we look for one-entries.	19

980 List of Tables

⁹⁸¹ List of Abbreviations

982 Attachments